

✓ Lab 5 — Recovery / Continue From Last Checkpoint

This notebook is a **recovery notebook**, not a fresh Lab 5 run.

It reads the existing Lab 5 data/checkpoints from Google Drive and determines the latest completed stage. The supplied notebook shows that the last successfully executed experiment was **Cell 34 — Best Optimizer**; Cell 35 (Hyperparameter Tuning) was not completed. Therefore the normal starting point is Cell 35.

The notebook is designed to:

- reuse the existing Oxford-IIIT Pet dataset on Drive;
- load existing optimizer/regularization/initialization results where available;
- skip completed hyperparameter configurations;
- save each hyperparameter result immediately to Drive;
- continue through transfer learning, fine-tuning, cross-validation, final training and evaluation;
- recover from another runtime interruption without starting completed configurations again;
- save plots as both PNG and EPS at 600 DPI.

```
# =====
# CELL R1 - IMPORTS + GOOGLE DRIVE
# =====

import os
import gc
import json
import time
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    precision_score,
    recall_score,
    f1_score
)

from google.colab import drive

# -----
# Mount Google Drive
# -----

drive.mount("/content/drive", force_remount=False)

# -----
# Drive paths
# -----

LAB5_DRIVE_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

# Main recovery folders
CHECKPOINT_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "checkpoints"
)

RESULTS_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "results"
)

MODELS_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "models"
)
```

```

OUTPUT_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "outputs"
)

# -----
# Dataset location
# -----
# Actual structure:
#
# My Drive/
#   └─ Lab5/
#       ├── checkpoints/
#       ├── models/
#       ├── outputs/
#       ├── results/
#       └─ lab5/
#           ├── checkpoints/
#           ├── dataset/
#           └─ results/
#
# Therefore the dataset is inside Lab5/lab5/.

DATA_ROOT = os.path.join(
    LAB5_DRIVE_ROOT,
    "dataset"
)

IMAGES_DIR = os.path.join(
    DATA_ROOT,
    "images"
)

ANNOTATIONS_DIR = os.path.join(
    DATA_ROOT,
    "annotations"
)

# -----
# Create required directories
# -----

for d in [
    CHECKPOINT_DIR,
    RESULTS_DIR,
    MODELS_DIR,
    OUTPUT_DIR
]:
    os.makedirs(d, exist_ok=True)

# -----
# Verify dataset
# -----

print("Dataset root:", DATA_ROOT)
print("Images directory:", IMAGES_DIR)
print("Annotations directory:", ANNOTATIONS_DIR)

if not os.path.exists(DATA_ROOT):
    raise FileNotFoundError(
        f"Dataset not found at:\n{DATA_ROOT}"
    )

if not os.path.exists(IMAGES_DIR):
    raise FileNotFoundError(
        f"Images directory not found at:\n{IMAGES_DIR}"
    )

if not os.path.exists(ANNOTATIONS_DIR):
    raise FileNotFoundError(
        f"Annotations directory not found at:\n{ANNOTATIONS_DIR}"
    )

print("Dataset folders verified.")

# -----
# Experiment settings
# -----

```

```

SEED = 42

IMG_SIZE = (224, 224)

NUM_CLASSES = 37

BATCH_SIZE = 32

EPOCHS = 5

# -----
# Reproducibility
# -----

random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)

# -----
# Environment information
# -----

print("\n" + "=" * 60)
print("LAB 5 RECOVERY ENVIRONMENT")
print("=" * 60)

print("Drive root:", LAB5_DRIVE_ROOT)
print("Dataset:", DATA_ROOT)
print("Checkpoint directory:", CHECKPOINT_DIR)
print("Results directory:", RESULTS_DIR)
print("Models directory:", MODELS_DIR)
print("Output directory:", OUTPUT_DIR)

print("TensorFlow:", tf.__version__)
print(
    "Available GPUs:",
    tf.config.list_physical_devices("GPU")
)

print("=" * 60)
print("Lab 5 recovery environment ready.")

```

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
Dataset root: /content/drive/MyDrive/Lab5/Lab5/dataset
Images directory: /content/drive/MyDrive/Lab5/Lab5/dataset/images
Annotations directory: /content/drive/MyDrive/Lab5/Lab5/dataset/annotations
Dataset folders verified.

```

```

=====
LAB 5 RECOVERY ENVIRONMENT
=====
Drive root: /content/drive/MyDrive/Lab5/Lab5
Dataset: /content/drive/MyDrive/Lab5/Lab5/dataset
Checkpoint directory: /content/drive/MyDrive/Lab5/Lab5/checkpoints
Results directory: /content/drive/MyDrive/Lab5/Lab5/results
Models directory: /content/drive/MyDrive/Lab5/Lab5/models
Output directory: /content/drive/MyDrive/Lab5/Lab5/outputs
TensorFlow: 2.20.0
Available GPUs: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]
=====
Lab 5 recovery environment ready.

```

```

# =====
# CELL R2 – FIGURE SAVING
# =====

plt.rcParams["font.family"] = "DejaVu Sans"
plt.rcParams["axes.titlesize"] = 14
plt.rcParams["axes.labelsize"] = 12
plt.rcParams["xtick.labelsize"] = 10
plt.rcParams["ytick.labelsize"] = 10
plt.rcParams["legend.fontsize"] = 10

def save_figure(name):
    eps_path = os.path.join(OUTPUT_DIR, name + ".eps")
    png_path = os.path.join(OUTPUT_DIR, name + ".png")

    plt.savefig(
        eps_path,
        format="eps",

```

```

        dpi=600,
        bbox_inches="tight"
    )
    plt.savefig(
        png_path,
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

    print("Saved:", eps_path)
    print("Saved:", png_path)

```

```

# =====
# CELL R3 – RECOVERY HELPERS
# =====

def experiment_dir(name):
    path = os.path.join(CHECKPOINT_DIR, name)
    os.makedirs(path, exist_ok=True)
    return path

def status_path(name):
    return os.path.join(experiment_dir(name), "status.json")

def write_status(name, status, **extra):
    payload = {
        "experiment": name,
        "status": status,
        "updated_at": time.strftime("%Y-%m-%d %H:%M:%S"),
        **extra
    }
    with open(status_path(name), "w") as f:
        json.dump(payload, f, indent=2)

def read_status(name):
    path = status_path(name)
    if not os.path.exists(path):
        return None
    try:
        with open(path, "r") as f:
            return json.load(f)
    except Exception:
        return None

def history_path(name):
    return os.path.join(experiment_dir(name), "history.json")

def save_history(name, history, metadata=None):
    payload = {
        "history": {
            k: [float(x) for x in v]
            for k, v in history.history.items()
        },
        "metadata": metadata or {},
        "saved_at": time.strftime("%Y-%m-%d %H:%M:%S")
    }
    with open(history_path(name), "w") as f:
        json.dump(payload, f, indent=2)

def load_history(name):
    path = history_path(name)
    if not os.path.exists(path):
        return None
    try:
        with open(path, "r") as f:
            payload = json.load(f)
        return payload.get("history", payload)
    except Exception:
        return None

def save_csv(df, filename):
    path = os.path.join(RESULTS_DIR, filename)
    df.to_csv(path, index=False)

```

```

    return path

def load_csv(filename):
    path = os.path.join(RESULTS_DIR, filename)
    if not os.path.exists(path):
        return None
    try:
        return pd.read_csv(path)
    except Exception:
        return None

def make_callbacks(name):
    d = experiment_dir(name)

    return [
        keras.callbacks.ModelCheckpoint(
            os.path.join(d, "best_model.weights.h5"),
            monitor="val_accuracy",
            mode="max",
            save_best_only=True,
            save_weights_only=True,
            verbose=1
        ),
        keras.callbacks.ModelCheckpoint(
            os.path.join(d, "last_model.weights.h5"),
            monitor="val_accuracy",
            mode="max",
            save_best_only=False,
            save_weights_only=True,
            verbose=0
        ),
        keras.callbacks.BackupAndRestore(
            backup_dir=os.path.join(d, "training_backup"),
            delete_checkpoint=False
        )
    ]

print("Recovery helpers ready.")

```

Recovery helpers ready.

```

# =====
# CELL R4 - RECREATE DATASET FROM DRIVE
# =====

# Actual Drive structure:
#
# My Drive/
# └─ Lab5/
#     │ checkpoints/
#     │ models/
#     │ outputs/
#     │ results/
#     └─ lab5/
#         │ checkpoints/
#         │ dataset/
#         │ │ images/
#         │ │ └─ annotations/
#         └─ results/

DATA_ROOT = os.path.join(
    LAB5_DRIVE_ROOT,
    "dataset"
)

IMAGES_DIR = os.path.join(
    DATA_ROOT,
    "images"
)

ANNOTATIONS_DIR = os.path.join(
    DATA_ROOT,
    "annotations"
)

# -----
# Verify dataset location

```

```

# -----

print("Dataset root:", DATA_ROOT)
print("Images directory:", IMAGES_DIR)
print("Annotations directory:", ANNOTATIONS_DIR)

if not os.path.exists(IMAGES_DIR):
    raise FileNotFoundError(
        f"Images directory not found:\n{IMAGES_DIR}"
    )

if not os.path.exists(ANNOTATIONS_DIR):
    raise FileNotFoundError(
        f"Annotations directory not found:\n{ANNOTATIONS_DIR}"
    )

# =====
# READ SPLIT FILES
# =====

def read_split_file(filename):

    names = []

    with open(filename, "r") as f:

        for line in f:

            line = line.strip()

            if not line or line.startswith("#"):
                continue

            names.append(line.split()[0])

    return set(names)

LIST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "list.txt"
)

TRAINVAL_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "trainval.txt"
)

TEST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "test.txt"
)

# =====
# READ ALL IMAGES
# =====

samples = []

with open(LIST_FILE, "r") as f:

    for line in f:

        line = line.strip()

        if not line or line.startswith("#"):
            continue

        parts = line.split()

        image_name = parts[0]

        image_path = os.path.join(
            IMAGES_DIR,
            image_name + ".jpg"
        )

        if os.path.exists(image_path):
            samples.append(image_path)

```

```

# =====
# CREATE CLASS LABELS
# =====

class_names = sorted(
    set(
        os.path.basename(p)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
        for p in samples
    )
)

class_to_index = {
    name: i
    for i, name in enumerate(class_names)
}

image_paths = []
labels = []

for path in samples:

    breed = (
        os.path.basename(path)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
    )

    if breed in class_to_index:

        image_paths.append(path)

        labels.append(
            class_to_index[breed]
        )

image_paths = np.array(image_paths)

labels = np.array(
    labels,
    dtype=np.int32
)

# =====
# TRAIN / TEST SPLIT
# =====

trainval_names = read_split_file(
    TRAINVAL_FILE
)

test_names = read_split_file(
    TEST_FILE
)

trainval_paths = []
trainval_labels = []

test_paths = []
test_labels = []

for path, label in zip(
    image_paths,
    labels
):

    name = (
        os.path.basename(path)
        .replace(".jpg", "")
    )

    if name in trainval_names:

        trainval_paths.append(path)

```

```

        trainval_labels.append(label)

    elif name in test_names:

        test_paths.append(path)
        test_labels.append(label)

trainval_paths = np.array(
    trainval_paths
)

trainval_labels = np.array(
    trainval_labels,
    dtype=np.int32
)

test_paths = np.array(
    test_paths
)

test_labels = np.array(
    test_labels,
    dtype=np.int32
)

# =====
# DETERMINISTIC TRAIN / VALIDATION SPLIT
# =====

rng = np.random.default_rng(SEED)

indices = np.arange(
    len(trainval_paths)
)

rng.shuffle(indices)

val_size = int(
    0.20 * len(indices)
)

val_indices = indices[:val_size]

train_indices = indices[val_size:]

train_paths = trainval_paths[
    train_indices
]

train_labels = trainval_labels[
    train_indices
]

val_paths = trainval_paths[
    val_indices
]

val_labels = trainval_labels[
    val_indices
]

# =====
# IMAGE LOADING
# =====

def load_image(path, label):

    image = tf.io.read_file(path)

    image = tf.image.decode_jpeg(
        image,
        channels=3
    )

    image = tf.image.resize(
        image,
        IMG_SIZE
    )

```

```

        image = tf.cast(
            image,
            tf.float32
        )

        image = preprocess_input(
            image
        )

        return image, label

# =====
# DATASET CREATION
# =====

def make_dataset(
    paths,
    labels,
    batch_size=32,
    shuffle=False
):

    ds = tf.data.Dataset.from_tensor_slices(
        (paths, labels)
    )

    if shuffle:

        ds = ds.shuffle(
            len(paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        batch_size
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

train_ds = make_dataset(
    train_paths,
    train_labels,
    BATCH_SIZE,
    shuffle=True
)

val_ds = make_dataset(
    val_paths,
    val_labels,
    BATCH_SIZE,
    shuffle=False
)

test_ds = make_dataset(
    test_paths,
    test_labels,
    BATCH_SIZE,
    shuffle=False
)

# =====
# SUMMARY
# =====

print("\n" + "=" * 60)
print("DATASET RECONSTRUCTED FROM DRIVE")
print("=" * 60)

```

```
print("Dataset:", DATA_ROOT)
print("Classes:", len(class_names))
print("Train:", len(train_paths))
print("Validation:", len(val_paths))
print("Test:", len(test_paths))
```

```
print("=" * 60)
```

```
Dataset root: /content/drive/MyDrive/Lab5/Lab5/dataset
Images directory: /content/drive/MyDrive/Lab5/Lab5/dataset/images
Annotations directory: /content/drive/MyDrive/Lab5/Lab5/dataset/annotations
```

```
=====
DATASET RECONSTRUCTED FROM DRIVE
=====
Dataset: /content/drive/MyDrive/Lab5/Lab5/dataset
Classes: 37
Train: 2944
Validation: 736
Test: 3669
=====
```

```
# =====
# CELL R5 - LOAD ALL EXISTING RESULTS
# =====
```

```
initialization_df = load_csv("initialization_results.csv")
regularization_df = load_csv("regularization_results.csv")
optimizer_df = load_csv("optimizer_results.csv")
hyperparameter_df = load_csv("hyperparameter_results.csv")
transfer_df = load_csv("transfer_learning_results.csv")
fine_tune_df = load_csv("correct_fine_tuning_results.csv")
cv_df = load_csv("cv_results.csv")
final_df = load_csv("final_model_results.csv")
```

```
print("Existing saved results:")
for label, df in [
    ("Initialization", initialization_df),
    ("Regularization", regularization_df),
    ("Optimizer", optimizer_df),
    ("Hyperparameters", hyperparameter_df),
    ("Transfer learning", transfer_df),
    ("Fine-tuning", fine_tune_df),
    ("Cross-validation", cv_df),
    ("Final model", final_df),
]:
    if df is None:
        print(f" {label}: not found")
    else:
        print(f" {label}: {len(df)} row(s)")
```

```
Existing saved results:
Initialization: 4 row(s)
Regularization: 4 row(s)
Optimizer: 4 row(s)
Hyperparameters: not found
Transfer learning: not found
Fine-tuning: not found
Cross-validation: not found
Final model: not found
```

```
# =====
# CELL R6 - RECOVERY UTILITIES
# =====
```

```
import os
import json
import gc
import time
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
```

```
# -----
# Confirm paths
# -----
```

```

print("LAB5 ROOT      :", LAB5_DRIVE_ROOT)
print("CHECKPOINT DIR:", CHECKPOINT_DIR)
print("RESULTS DIR   :", RESULTS_DIR)
print("OUTPUT DIR    :", OUTPUT_DIR)

# -----
# Safe metric extraction
# -----

def get_history_values(history):

    h = history.history

    # Explicit metric name is used in all new models.
    train_acc = h.get("accuracy", h.get("acc"))
    val_acc = h.get("val_accuracy", h.get("val_acc"))

    loss = h.get("loss")
    val_loss = h.get("val_loss")

    if train_acc is None:
        raise RuntimeError(
            f"Training accuracy missing. Available keys: {list(h.keys())}"
        )

    if val_acc is None:
        raise RuntimeError(
            f"Validation accuracy missing. Available keys: {list(h.keys())}"
        )

    if loss is None or val_loss is None:
        raise RuntimeError(
            f"Loss values missing. Available keys: {list(h.keys())}"
        )

    return (
        np.asarray(train_acc),
        np.asarray(val_acc),
        np.asarray(loss),
        np.asarray(val_loss)
    )

# -----
# Save Keras history
# -----

def save_history_json(history, path):

    with open(path, "w") as f:
        json.dump(
            {
                k: [float(x) for x in v]
                for k, v in history.history.items()
            },
            f,
            indent=2
        )

# -----
# Save experiment state
# -----

def save_json(data, path):

    with open(path, "w") as f:
        json.dump(data, f, indent=2)

# -----
# Safe figure saver
# -----

def save_figure(name):

    eps_path = os.path.join(
        OUTPUT_DIR,
        name + ".eps"
    )

```

```

    png_path = os.path.join(
        OUTPUT_DIR,
        name + ".png"
    )

    plt.savefig(
        eps_path,
        format="eps",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        png_path,
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

    print("Saved:", eps_path)
    print("Saved:", png_path)

# -----
# Common compilation
# -----

def compile_model(model, learning_rate):

    model.compile(
        optimizer=keras.optimizers.RMSprop(
            learning_rate=learning_rate
        ),
        loss="sparse_categorical_crossentropy",
        metrics=[
            keras.metrics.SparseCategoricalAccuracy(
                name="accuracy"
            )
        ]
    )

print("Recovery utilities ready.")

```

```

LAB5_ROOT      : /content/drive/MyDrive/Lab5/Lab5
CHECKPOINT_DIR : /content/drive/MyDrive/Lab5/Lab5/checkpoints
RESULTS_DIR    : /content/drive/MyDrive/Lab5/Lab5/results
OUTPUT_DIR     : /content/drive/MyDrive/Lab5/Lab5/outputs
Recovery utilities ready.

```

```

# =====
# CELL R7 - HYPERPARAMETER RECOVERY SETUP
# =====

HYPERPARAMETER_DIR = os.path.join(
    CHECKPOINT_DIR,
    "hyperparameters"
)

os.makedirs(
    HYPERPARAMETER_DIR,
    exist_ok=True
)

learning_rates = [
    0.001,
    0.0001
]

batch_sizes = [
    16,
    32,
    64
]

dropout_rates = [
    0.0,
    0.25,
    0.5
]

```

```

hyperparameter_histories = {}

hyperparameter_results = []

# -----
# Recover LR checkpoints by validation evaluation
# -----

def build_hyperparameter_model(dropout_rate=0.0):

    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    base_model.trainable = False

    inputs = keras.Input(
        shape=(224, 224, 3)
    )

    x = base_model(
        inputs,
        training=False
    )

    x = layers.GlobalAveragePooling2D()(x)

    if dropout_rate > 0:
        x = layers.Dropout(
            dropout_rate
        )(x)

    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax"
    )(x)

    return keras.Model(
        inputs,
        outputs
    )

print("Hyperparameter recovery setup ready.")

```

Hyperparameter recovery setup ready.

```

# =====
# CELL R8 - RECOVER LEARNING-RATE CHECKPOINTS
# =====

lr_checkpoint_map = {
    0.001: os.path.join(
        CHECKPOINT_DIR,
        "hyperparameter_lr_0_001",
        "best_model.weights.h5"
    ),
    0.0001: os.path.join(
        CHECKPOINT_DIR,
        "hyperparameter_lr_0_0001",
        "best_model.weights.h5"
    )
}

for lr in learning_rates:

    checkpoint = lr_checkpoint_map[lr]

    if not os.path.exists(checkpoint):

        print(
            f"Checkpoint missing for learning rate {lr}"
        )

        continue

```

```

print("\n" + "=" * 60)
print("Recovering learning rate:", lr)
print("=" * 60)

model = build_hyperparameter_model(
    dropout_rate=0.0
)

compile_model(
    model,
    lr
)

model.load_weights(
    checkpoint
)

results = model.evaluate(
    val_ds,
    verbose=1,
    return_dict=True
)

val_accuracy = (
    results["accuracy"] * 100
)

val_loss = results["loss"]

hyperparameter_results.append({
    "Hyperparameter":
        "Learning Rate",

    "Value":
        lr,

    "Best Validation Accuracy (%)":
        val_accuracy,

    "Final Training Accuracy (%)":
        np.nan,

    "Final Validation Accuracy (%)":
        val_accuracy,

    "Final Training Loss":
        np.nan,

    "Final Validation Loss":
        val_loss,

    "Training Time (s)":
        np.nan
})

print(
    f"Recovered validation accuracy: "
    f"{val_accuracy:.2f}%"
)

del model

tf.keras.backend.clear_session()
gc.collect()

print("\nRecovered learning-rate results.")

```

```

=====
Recovering learning rate: 0.001
=====
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\_v2/mobilenet\_v2\_weights\_tf\_dim\_9406464/9406464 0s 0us/step
/usr/local/lib/python3.13/dist-packages/keras/src/saving/saving_lib.py:797: UserWarning: Skipping variable loading for optimizer. load_own_variables(weights_store.get(inner_path))

```

```
23/23 ----- 116s 4s/step - accuracy: 0.9647 - loss: 0.1171
Recovered validation accuracy: 96.47%
```

```
=====
Recovering learning rate: 0.0001
=====
```

```
23/23 ----- 8s 140ms/step - accuracy: 0.7853 - loss: 1.2185
Recovered validation accuracy: 78.53%
```

```
Recovered learning-rate results.
```

```
# =====
# CELL R8.5 - HYPERPARAMETER MODEL BUILDER
# =====
```

```
def build_hp_model(dropout_rate=0.0):

    # Fresh pretrained MobileNetV2
    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    # Freeze pretrained backbone
    base_model.trainable = False

    # Input
    inputs = keras.Input(
        shape=(224, 224, 3)
    )

    # Feature extraction
    x = base_model(
        inputs,
        training=False
    )

    # Global pooling
    x = layers.GlobalAveragePooling2D()(x)

    # Optional dropout
    if dropout_rate > 0:
        x = layers.Dropout(
            dropout_rate
        )(x)

    # Classification head
    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax"
    )(x)

    # Build model
    model = keras.Model(
        inputs=inputs,
        outputs=outputs
    )

    return model
```

```
print("build_hp_model() is ready.")
```

```
build_hp_model() is ready.
```

```
# =====
# CELL R9 - RESUMABLE BATCH-SIZE EXPERIMENT
# =====
```

```
print("=" * 70)
print("BATCH SIZE EXPERIMENT")
print("=" * 70)
```

```
batch_sizes = [16, 32, 64]
```

```
# Make sure the results dataframe exists
if hyperparameter_df is None:
    hyperparameter_df = pd.DataFrame()
```

```
# -----
```

```

# Check whether a configuration is already saved
# -----

def hp_done(hyperparameter, value):

    if hyperparameter_df.empty:
        return False

    if "Hyperparameter" not in hyperparameter_df.columns:
        return False

    rows = hyperparameter_df[
        hyperparameter_df["Hyperparameter"].astype(str)
        == str(hyperparameter)
    ]

    if rows.empty:
        return False

    if "Value" not in rows.columns:
        return False

    return np.any(
        np.isclose(
            rows["Value"].astype(float).to_numpy(),
            float(value)
        )
    )

# -----
# Run each batch size
# -----

for batch_size in batch_sizes:

    key = f"Batch Size = {batch_size}"

    print("\n" + "=" * 60)
    print(f"Batch Size: {batch_size}")
    print("=" * 60)

    # Skip if already completed
    if hp_done("Batch Size", batch_size):

        print(f"SKIP - already saved: {key}")
        continue

    # -----
    # Create datasets using the variables that actually exist
    # -----

    train_batch_ds = make_dataset(
        train_paths,
        train_labels,
        batch_size=batch_size,
        shuffle=True
    )

    val_batch_ds = make_dataset(
        val_paths,
        val_labels,
        batch_size=batch_size,
        shuffle=False
    )

    # -----
    # Create fresh model
    # -----

    model = build_hp_model(
        dropout_rate=0.0
    )

    # -----
    # Compile
    # -----

    model.compile(
        optimizer=keras.optimizers.RMSprop(
            learning_rate=0.001
        ),

```

```

        loss="sparse_categorical_crossentropy",
        metrics=[
            keras.metrics.SparseCategoricalAccuracy(
                name="accuracy"
            )
        ]
    )

# -----
# Checkpoint paths
# -----

exp_name = f"hyperparameter_bs_{batch_size}"

exp_checkpoint_dir = os.path.join(
    CHECKPOINT_DIR,
    exp_name
)

os.makedirs(
    exp_checkpoint_dir,
    exist_ok=True
)

best_model_path = os.path.join(
    exp_checkpoint_dir,
    "best_model.keras"
)

history_path = os.path.join(
    exp_checkpoint_dir,
    "history.json"
)

status_path = os.path.join(
    exp_checkpoint_dir,
    "status.json"
)

# -----
# Callbacks
# -----

callbacks = [

    keras.callbacks.ModelCheckpoint(
        best_model_path,
        monitor="val_accuracy",
        mode="max",
        save_best_only=True,
        save_weights_only=False,
        verbose=1
    ),

    keras.callbacks.CSVLogger(
        os.path.join(
            exp_checkpoint_dir,
            "training_log.csv"
        ),
        append=True
    )
]

# -----
# Mark as running
# -----

with open(status_path, "w") as f:
    json.dump(
        {
            "status": "running",
            "hyperparameter": "Batch Size",
            "value": batch_size
        },
        f,
        indent=4
    )

# -----
# Train
# -----

```

```

start_time = time.time()

history = model.fit(
    train_batch_ds,
    validation_data=val_batch_ds,
    epochs=EPOCHS,
    verbose=1,
    callbacks=callbacks
)

training_time = time.time() - start_time

# -----
# Extract metrics safely
# -----

h = history.history

print("Available metrics:", list(h.keys()))

train_accuracy = h.get("accuracy", [])
val_accuracy = h.get("val_accuracy", [])
train_loss = h.get("loss", [])
val_loss = h.get("val_loss", [])

if len(val_accuracy) == 0:
    raise RuntimeError(
        f"No validation accuracy was recorded for "
        f"batch size {batch_size}. "
        f"Available metrics: {list(h.keys())}"
    )

# -----
# Save history
# -----

with open(history_path, "w") as f:
    json.dump(
        {
            k: [float(x) for x in v]
            for k, v in h.items()
        },
        f,
        indent=4
    )

# -----
# Store result
# -----

result = {

    "Hyperparameter": "Batch Size",

    "Value": batch_size,

    "Best Validation Accuracy (%)":
        max(val_accuracy) * 100,

    "Final Training Accuracy (%)":
        train_accuracy[-1] * 100
        if len(train_accuracy) else np.nan,

    "Final Validation Accuracy (%)":
        val_accuracy[-1] * 100,

    "Final Training Loss":
        train_loss[-1]
        if len(train_loss) else np.nan,

    "Final Validation Loss":
        val_loss[-1]
        if len(val_loss) else np.nan,

    "Training Time (s)":
        training_time
}

hyperparameter_df = pd.concat(
    [
        hyperparameter_df,
        pd.DataFrame([result])
    ]

```

```

        ],
        ignore_index=True
    )

    # -----
    # Save results immediately
    # -----

    save_csv(
        hyperparameter_df,
        "hyperparameter_results.csv"
    )

    # -----
    # Mark completed
    # -----

    with open(status_path, "w") as f:
        json.dump(
            {
                "status": "completed",
                "hyperparameter": "Batch Size",
                "value": batch_size
            },
            f,
            indent=4
        )

    print(
        f"Completed Batch Size {batch_size} | "
        f"Best Val Accuracy: "
        f"{max(val_accuracy) * 100:.2f}%"
    )

    # -----
    # Clear memory
    # -----

    del model
    del train_batch_ds
    del val_batch_ds
    del history

    tf.keras.backend.clear_session()
    gc.collect()

    print("\n" + "=" * 70)
    print("BATCH SIZE EXPERIMENT COMPLETE")
    print("=" * 70)

    display(
        hyperparameter_df[
            hyperparameter_df["Hyperparameter"]
            == "Batch Size"
        ]
    )

```

```
=====
BATCH SIZE EXPERIMENT
=====
```

```
=====
Batch Size: 16
=====
```

```
Epoch 1/5
184/184 ----- 0s 2s/step - accuracy: 0.5845 - loss: 1.6790
Epoch 1: val_accuracy improved from None to 0.90489, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperpar

Epoch 1: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperparameter_bs_16/best_model.keras
184/184 ----- 451s 2s/step - accuracy: 0.7653 - loss: 0.9191 - val_accuracy: 0.9049 - val_loss: 0.3368
Epoch 2/5
184/184 ----- 0s 57ms/step - accuracy: 0.9333 - loss: 0.2413
Epoch 2: val_accuracy improved from 0.90489 to 0.90625, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyper

Epoch 2: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperparameter_bs_16/best_model.keras
184/184 ----- 71s 88ms/step - accuracy: 0.9344 - loss: 0.2295 - val_accuracy: 0.9062 - val_loss: 0.3066
Epoch 3/5
184/184 ----- 0s 54ms/step - accuracy: 0.9555 - loss: 0.1647
Epoch 3: val_accuracy did not improve from 0.90625
184/184 ----- 14s 75ms/step - accuracy: 0.9579 - loss: 0.1511 - val_accuracy: 0.8995 - val_loss: 0.2850
Epoch 4/5
184/184 ----- 0s 51ms/step - accuracy: 0.9744 - loss: 0.0944
Epoch 4: val_accuracy did not improve from 0.90625
184/184 ----- 13s 71ms/step - accuracy: 0.9688 - loss: 0.1034 - val_accuracy: 0.9035 - val_loss: 0.2775
Epoch 5/5
183/184 ----- 0s 54ms/step - accuracy: 0.9841 - loss: 0.0681
Epoch 5: val_accuracy improved from 0.90625 to 0.91304, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyper

Epoch 5: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperparameter_bs_16/best_model.keras
184/184 ----- 14s 75ms/step - accuracy: 0.9783 - loss: 0.0744 - val_accuracy: 0.9130 - val_loss: 0.2693
Available metrics: ['accuracy', 'loss', 'val_accuracy', 'val_loss']
Completed Batch Size 16 | Best Val Accuracy: 91.30%
```

```
=====
Batch Size: 32
=====
```

```
Epoch 1/5
92/92 ----- 0s 114ms/step - accuracy: 0.5473 - loss: 1.8587
Epoch 1: val_accuracy improved from None to 0.88723, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperpar

Epoch 1: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperparameter_bs_32/best_model.keras
92/92 ----- 25s 197ms/step - accuracy: 0.7357 - loss: 1.1055 - val_accuracy: 0.8872 - val_loss: 0.4325
Epoch 2/5
92/92 ----- 0s 117ms/step - accuracy: 0.9340 - loss: 0.3015
Epoch 2: val_accuracy improved from 0.88723 to 0.89538, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyper

Epoch 2: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/hyperparameter_bs_32/best_model.keras
92/92 ----- 32s 145ms/step - accuracy: 0.9287 - loss: 0.2828 - val_accuracy: 0.8954 - val_loss: 0.3488
Epoch 3/5
92/92 ----- 0s 115ms/step - accuracy: 0.9514 - loss: 0.1912
Epoch 3: val_accuracy did not improve from 0.89538
92/92 ----- 13s 139ms/step - accuracy: 0.9558 - loss: 0.1735 - val_accuracy: 0.8899 - val_loss: 0.3186
Epoch 4/5
92/92 ----- 0s 106ms/step - accuracy: 0.9728 - loss: 0.1113
```

```
# =====
# CELL R10 - DROPOUT EXPERIMENT
# =====
```

```
for dropout_rate in dropout_rates:
```

```
    marker_name = (
        str(dropout_rate)
        .replace(".", "_")
    )
```

```
    marker = os.path.join(
        HYPERPARAMETER_DIR,
        f"dropout_{marker_name}.done"
    )
```

```
    if os.path.exists(marker):
```

```
        print(
            f"Skipping completed dropout "
            f"{dropout_rate}"
        )
```

```
        continue
```

```
    print("\n" + "=" * 60)
    print("Dropout:", dropout_rate)
```

```
print("=" * 60)

model = build_hyperparameter_model(
    dropout_rate=dropout_rate
)

compile_model(
    model,
    0.001
)

start_time = time.time()

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=EPOCHS,
    verbose=1
)

training_time = (
    time.time() - start_time
)

train_acc, val_acc, loss, val_loss = (
    get_history_values(history)
)

hyperparameter_histories[
    f"Dropout = {dropout_rate}"
] = history

row = {

    "Hyperparameter":
        "Dropout",

    "Value":
        dropout_rate,

    "Best Validation Accuracy (%)":
        float(np.max(val_acc) * 100),

    "Final Training Accuracy (%)":
        float(train_acc[-1] * 100),

    "Final Validation Accuracy (%)":
        float(val_acc[-1] * 100),

    "Final Training Loss":
        float(loss[-1]),

    "Final Validation Loss":
        float(val_loss[-1]),

    "Training Time (s)":
        training_time
}

append_hyperparameter_result(
    row
)

checkpoint_path = os.path.join(
    HYPERPARAMETER_DIR,
    f"dropout_{marker_name}.keras"
)

model.save(
    checkpoint_path
)

history_path = os.path.join(
    HYPERPARAMETER_DIR,
```

```
        f"dropout_{marker_name}_history.json"
    )

    save_history_json(
        history,
        history_path
    )

    open(marker, "w").close()

    del model

    tf.keras.backend.clear_session()
    gc.collect()

hyperparameter_df = pd.DataFrame(
    hyperparameter_results
)

hyperparameter_df.to_csv(
    hyperparameter_csv,
    index=False
)

display(
    hyperparameter_df
)
```

```
=====
Dropout: 0.0
=====
Epoch 1/5
92/92 ----- 29s 229ms/step - accuracy: 0.7408 - loss: 1.1270 - val_accuracy: 0.9022 - val_loss: 0.4312
Epoch 2/5
92/92 ----- 29s 145ms/step - accuracy: 0.9300 - loss: 0.2800 - val_accuracy: 0.9062 - val_loss: 0.3335
Epoch 3/5
92/92 ----- 14s 149ms/step - accuracy: 0.9579 - loss: 0.1714 - val_accuracy: 0.9035 - val_loss: 0.3073
Epoch 4/5
92/92 ----- 14s 148ms/step - accuracy: 0.9667 - loss: 0.1167 - val_accuracy: 0.9049 - val_loss: 0.2890
Epoch 5/5
92/92 ----- 13s 143ms/step - accuracy: 0.9789 - loss: 0.0829 - val_accuracy: 0.9090 - val_loss: 0.2761
Checkpointed: /content/drive/MyDrive/Lab5/Lab5/results/hyperparameter_results.csv
```

```
=====
Dropout: 0.25
=====
Epoch 1/5
92/92 ----- 25s 187ms/step - accuracy: 0.6841 - loss: 1.2974 - val_accuracy: 0.8899 - val_loss: 0.4857
Epoch 2/5
92/92 ----- 13s 142ms/step - accuracy: 0.8988 - loss: 0.3621 - val_accuracy: 0.8995 - val_loss: 0.3412
Epoch 3/5
92/92 ----- 20s 143ms/step - accuracy: 0.9348 - loss: 0.2352 - val_accuracy: 0.9117 - val_loss: 0.2876
Epoch 4/5
92/92 ----- 13s 144ms/step - accuracy: 0.9474 - loss: 0.1730 - val_accuracy: 0.8995 - val_loss: 0.2948
Epoch 5/5
92/92 ----- 12s 131ms/step - accuracy: 0.9592 - loss: 0.1309 - val_accuracy: 0.9008 - val_loss: 0.3009
Checkpointed: /content/drive/MyDrive/Lab5/Lab5/results/hyperparameter_results.csv
```

```
=====
Dropout: 0.5
=====
Epoch 1/5
92/92 ----- 26s 187ms/step - accuracy: 0.5839 - loss: 1.5846 - val_accuracy: 0.8533 - val_loss: 0.5678
Epoch 2/5
92/92 ----- 13s 144ms/step - accuracy: 0.8451 - loss: 0.5061 - val_accuracy: 0.8954 - val_loss: 0.3639
Epoch 3/5
92/92 ----- 19s 133ms/step - accuracy: 0.8916 - loss: 0.3423 - val_accuracy: 0.9022 - val_loss: 0.3197
Epoch 4/5
92/92 ----- 12s 132ms/step - accuracy: 0.9178 - loss: 0.2682 - val_accuracy: 0.9008 - val_loss: 0.3010
Epoch 5/5
92/92 ----- 13s 143ms/step - accuracy: 0.9283 - loss: 0.2164 - val_accuracy: 0.9008 - val_loss: 0.2974
Checkpointed: /content/drive/MyDrive/Lab5/Lab5/results/hyperparameter_results.csv
```

	Hyperparameter	Value	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Learning Rate	0.0010	96.467394	NaN	96.467394	NaN	0.117140	NaN
1	Learning Rate	0.0001	78.532606	NaN	78.532606	NaN	1.218464	NaN
2	Dropout	0.0000	90.896738	97.894019	90.896738	0.082908	0.276094	105.535548
3	Dropout	0.2500	91.168481	95.923913	90.081519	0.130857	0.300911	92.754024
4	Dropout	0.5000	90.217394	92.832881	90.081519	0.216385	0.297352	84.395862

```
# =====
# CELL R11 - PLOTS 8-10
# =====

# -----
# Plot 8 - Learning Rate
# -----

lr_df = hyperparameter_df[
    hyperparameter_df["Hyperparameter"]
    == "Learning Rate"
].copy()

plt.figure(figsize=(7, 5))

plt.plot(
    lr_df["Value"],
    lr_df["Best Validation Accuracy (%)"],
    marker="o",
    linewidth=1.8
)

plt.xscale("log")

plt.xlabel("Learning Rate")
plt.ylabel("Validation Accuracy (%)")
plt.title(
    "Learning Rate vs. Validation Accuracy"
```

```
)

plt.grid(alpha=0.2)

save_figure(
    "plot_8_learning_rate"
)

plt.show()
plt.close()

# -----
# Plot 9 – Batch Size
# -----

batch_df = hyperparameter_df[
    hyperparameter_df["Hyperparameter"]
    == "Batch Size"
].copy()

plt.figure(figsize=(7, 5))

plt.plot(
    batch_df["Value"],
    batch_df["Best Validation Accuracy (%)"],
    marker="o",
    linewidth=1.8
)

plt.xlabel("Batch Size")
plt.ylabel("Validation Accuracy (%)")
plt.title(
    "Batch Size vs. Validation Accuracy"
)

plt.grid(alpha=0.2)

save_figure(
    "plot_9_batch_size"
)

plt.show()
plt.close()

# -----
# Plot 10 – Dropout
# -----

dropout_df = hyperparameter_df[
    hyperparameter_df["Hyperparameter"]
    == "Dropout"
].copy()

plt.figure(figsize=(7, 5))

plt.plot(
    dropout_df["Value"],
    dropout_df["Best Validation Accuracy (%)"],
    marker="o",
    linewidth=1.8
)


plt.xlabel("Dropout Rate")
plt.ylabel("Validation Accuracy (%)")
plt.title(
    "Dropout Rate vs. Validation Accuracy"
)

plt.grid(alpha=0.2)


save_figure(
    "plot_10_dropout"
)

plt.show()
plt.close()
```

WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_8_learning_rate.eps
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_8_learning_rate.png



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_9_batch_size.eps
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_9_batch_size.png



```
# =====
# CELL R12 - BEST HYPERPARAMETERS
# =====

valid_hyperparameters = hyperparameter_df.dropna(
    subset=["Best Validation Accuracy (%)"]
).copy()

print("Available hyperparameter results:")
print(
    valid_hyperparameters[
        ["Hyperparameter", "Value", "Best Validation Accuracy (%)"]
    ].to_string(index=False)
)

print("\n" + "=" * 60)

# -----
# Helper
# -----

def get_best_hyperparameter(name):

    rows = valid_hyperparameters[
        valid_hyperparameters["Hyperparameter"].astype(str) == name
    ].copy()

    if rows.empty:
        return None
```

```

    return rows.loc[
        rows["Best Validation Accuracy (%)"].idxmax(),
        "Value"
    ]

best_lr = get_best_hyperparameter("Learning Rate")
best_batch_size = get_best_hyperparameter("Batch Size")
best_dropout = get_best_hyperparameter("Dropout")

# -----
# Display
# -----

print("BEST HYPERPARAMETERS")
print("=" * 50)

if best_lr is not None:
    print("Learning rate:", best_lr)
else:
    print("Learning rate: NOT AVAILABLE")

if best_batch_size is not None:
    print("Batch size:", best_batch_size)
else:
    print("Batch size: NOT AVAILABLE – batch-size experiment not completed")

if best_dropout is not None:
    print("Dropout:", best_dropout)
else:
    print("Dropout: NOT AVAILABLE – dropout experiment not completed")

```

```

Available hyperparameter results:
Hyperparameter Value Best Validation Accuracy (%)
Learning Rate 0.0010 96.467394
Learning Rate 0.0001 78.532606
Dropout 0.0000 90.896738
Dropout 0.2500 91.168481
Dropout 0.5000 90.217394

=====
BEST HYPERPARAMETERS
=====
Learning rate: 0.001
Batch size: NOT AVAILABLE – batch-size experiment not completed
Dropout: 0.25

```

```

# =====
# CELL R13 – FEATURE EXTRACTION VS FINE-TUNING
# =====

TRANSFER_DIR = os.path.join(
    CHECKPOINT_DIR,
    "transfer_learning"
)

os.makedirs(
    TRANSFER_DIR,
    exist_ok=True
)

transfer_histories = {}
transfer_results = []

def build_transfer_model(
    fine_tune=False,
    unfreeze_layers=20,
    dropout_rate=0.0
):

    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    if fine_tune:

        base_model.trainable = True

```

```

        for layer in base_model.layers[:-unfreeze_layers]:
            layer.trainable = False

    else:

        base_model.trainable = False

    inputs = keras.Input(
        shape=(224, 224, 3)
    )

    x = base_model(
        inputs,
        training=False
    )

    x = layers.GlobalAveragePooling2D()(x)

    if dropout_rate > 0:

        x = layers.Dropout(
            dropout_rate
        )(x)

    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax"
    )(x)

    return keras.Model(
        inputs,
        outputs
    )

transfer_configs = {

    "Feature Extraction": {
        "fine_tune": False,
        "lr": float(best_lr),
        "unfreeze": 0
    },

    "Fine-Tuning": {
        "fine_tune": True,
        "lr": 0.0001,
        "unfreeze": 20
    }
}

for name, config in transfer_configs.items():

    safe_name = (
        name.lower()
        .replace(" ", "_")
        .replace("-", "_")
    )

    marker = os.path.join(
        TRANSFER_DIR,
        safe_name + ".done"
    )

    if os.path.exists(marker):

        print(
            "Skipping completed:",
            name
        )

        continue

    print("\n" + "=" * 60)
    print(name)
    print("=" * 60)

```

```

model = build_transfer_model(
    fine_tune=config["fine_tune"],
    unfreeze_layers=config["unfreeze"],
    dropout_rate=float(best_dropout)
)

compile_model(
    model,
    config["lr"]
)

trainable_params = int(
    np.sum([
        np.prod(v.shape)
        for v in model.trainable_variables
    ])
)

total_params = int(
    model.count_params()
)

start_time = time.time()

history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=EPOCHS,
    verbose=1
)

training_time = (
    time.time() - start_time
)

train_acc, val_acc, loss, val_loss = (
    get_history_values(history)
)

transfer_histories[name] = history

transfer_results.append({
    "Transfer Learning Method":
        name,

    "Trainable Parameters":
        trainable_params,

    "Total Parameters":
        total_params,

    "Best Validation Accuracy (%)":
        float(np.max(val_acc) * 100),

    "Final Training Accuracy (%)":
        float(train_acc[-1] * 100),

    "Final Validation Accuracy (%)":
        float(val_acc[-1] * 100),

    "Final Training Loss":
        float(loss[-1]),

    "Final Validation Loss":
        float(val_loss[-1]),

    "Training Time (s)":
        training_time
})

model.save(
    os.path.join(

```

```

        TRANSFER_DIR,
        safe_name + ".keras"
    )
)

save_history_json(
    history,
    os.path.join(
        TRANSFER_DIR,
        safe_name + "_history.json"
    )
)

open(marker, "w").close()

del model

tf.keras.backend.clear_session()
gc.collect()

transfer_results_df = pd.DataFrame(
    transfer_results
)

transfer_results_df.to_csv(
    os.path.join(
        RESULTS_DIR,
        "transfer_learning_results.csv"
    ),
    index=False
)

display(
    transfer_results_df
)

```

```

=====
Feature Extraction
=====

```

```

Epoch 1/5
92/92 ————— 27s 185ms/step - accuracy: 0.6712 - loss: 1.2984 - val_accuracy: 0.8832 - val_loss: 0.4672
Epoch 2/5
92/92 ————— 13s 138ms/step - accuracy: 0.9056 - loss: 0.3496 - val_accuracy: 0.8927 - val_loss: 0.3115
Epoch 3/5
92/92 ————— 14s 153ms/step - accuracy: 0.9321 - loss: 0.2290 - val_accuracy: 0.8981 - val_loss: 0.3042
Epoch 4/5
92/92 ————— 20s 149ms/step - accuracy: 0.9477 - loss: 0.1687 - val_accuracy: 0.9076 - val_loss: 0.2694
Epoch 5/5
92/92 ————— 19s 131ms/step - accuracy: 0.9596 - loss: 0.1324 - val_accuracy: 0.9076 - val_loss: 0.2591

```

```

=====
Fine-Tuning
=====

```

```

Epoch 1/5
92/92 ————— 35s 241ms/step - accuracy: 0.6111 - loss: 1.5945 - val_accuracy: 0.8302 - val_loss: 0.5916
Epoch 2/5
92/92 ————— 13s 144ms/step - accuracy: 0.9022 - loss: 0.3978 - val_accuracy: 0.8587 - val_loss: 0.4197
Epoch 3/5
92/92 ————— 13s 143ms/step - accuracy: 0.9460 - loss: 0.2120 - val_accuracy: 0.8587 - val_loss: 0.3824
Epoch 4/5
92/92 ————— 12s 133ms/step - accuracy: 0.9752 - loss: 0.1148 - val_accuracy: 0.8804 - val_loss: 0.3537
Epoch 5/5
92/92 ————— 13s 138ms/step - accuracy: 0.9888 - loss: 0.0696 - val_accuracy: 0.8927 - val_loss: 0.3141

```

	Transfer Learning Method	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Feature Extraction	47397	2305381	90.760869	95.957881	90.760869	0.132448	0.259060	100.642656

```

# =====
# CELL R14 - PLOT 11
# FEATURE EXTRACTION VS FINE-TUNING
# =====

```

```

plt.figure(figsize=(7, 5))

for name, history in transfer_histories.items():

    _, val_acc, _, _ = get_history_values(
        history
    )

    plt.plot(
        np.arange(1, len(val_acc) + 1),
        val_acc * 100,
        marker="o",
        linewidth=1.8,
        label=name
    )

plt.xlabel("Epoch")
plt.ylabel("Validation Accuracy (%)")

plt.title(
    "Feature Extraction vs. Fine-Tuning"
)


plt.legend()
plt.grid(alpha=0.2)

save_figure(
    "plot_11_feature_extraction_vs_fine_tuning"
)

plt.show()
plt.close()

```

WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists w
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_11_feature_extraction_vs_fine_tuning.eps
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_11_feature_extraction_vs_fine_tuning.png



```

# =====
# CELL R15 - PLOT 12
# TRAINING AND VALIDATION LOSS
# =====

plt.figure(figsize=(7, 5))

for name, history in transfer_histories.items():

    _, _, loss, val_loss = get_history_values(
        history
    )

    epochs_axis = np.arange(
        1,
        len(loss) + 1
    )

```

```

plt.plot(
    epochs_axis,
    loss,
    linewidth=1.8,
    label=name + " - Train"
)

plt.plot(
    epochs_axis,
    val_loss,
    linestyle="--",
    linewidth=1.8,
    label=name + " - Validation"
)

plt.xlabel("Epoch")
plt.ylabel("Loss")

plt.title(
    "Training and Validation Loss"
)


plt.legend()
plt.grid(alpha=0.2)

save_figure(
    "plot_12_transfer_learning_loss"
)

plt.show()
plt.close()

```

WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists will be drawn as opaque.
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_12_transfer_learning_loss.eps
 Saved: /content/drive/MyDrive/Lab5/Lab5/outputs/plot_12_transfer_learning_loss.png



```

# =====
# CELL R16 - FINE-TUNING LEARNING-RATE STUDY
# =====

FT_STUDY_DIR = os.path.join(
    CHECKPOINT_DIR,
    "fine_tuning_lr"
)

os.makedirs(
    FT_STUDY_DIR,
    exist_ok=True
)

fine_tune_learning_rates = [
    0.001,
    0.0001
]

```

```
]

fine_tune_lr_results = []

for lr in fine_tune_learning_rates:

    safe_lr = str(lr).replace(".", "_")

    marker = os.path.join(
        FT_STUDY_DIR,
        f"lr_{safe_lr}.done"
    )

    if os.path.exists(marker):

        print(
            "Skipping completed LR:",
            lr
        )

        continue

    print("\n" + "=" * 60)
    print("Fine-tuning learning rate:", lr)
    print("=" * 60)

    model = build_transfer_model(
        fine_tune=True,
        unfreeze_layers=20,
        dropout_rate=float(best_dropout)
    )

    compile_model(
        model,
        lr
    )

    start_time = time.time()

    history = model.fit(
        train_ds,
        validation_data=val_ds,
        epochs=EPOCHS,
        verbose=1
    )

    training_time = (
        time.time() - start_time
    )

    train_acc, val_acc, loss, val_loss = (
        get_history_values(history)
    )

    fine_tune_lr_results.append({

        "Learning Rate":
            lr,

        "Best Validation Accuracy (%)":
            float(np.max(val_acc) * 100),

        "Final Validation Accuracy (%)":
            float(val_acc[-1] * 100),

        "Final Training Accuracy (%)":
            float(train_acc[-1] * 100),

        "Final Training Loss":
            float(loss[-1]),

        "Final Validation Loss":
            float(val_loss[-1]),

        "Training Time (s)":
```

```

        training_time
    })

    model.save(
        os.path.join(
            FT_STUDY_DIR,
            f"lr_{safe_lr}.keras"
        )
    )

    save_history_json(
        history,
        os.path.join(
            FT_STUDY_DIR,
            f"lr_{safe_lr}_history.json"
        )
    )

    open(marker, "w").close()

    del model

    tf.keras.backend.clear_session()
    gc.collect()

fine_tune_lr_df = pd.DataFrame(
    fine_tune_lr_results
)

fine_tune_lr_df.to_csv(
    os.path.join(
        RESULTS_DIR,
        "fine_tuning_learning_rate_results.csv"
    ),
    index=False
)

display(
    fine_tune_lr_df
)

```

```

=====
Fine-tuning learning rate: 0.001
=====
Epoch 1/5
92/92 ————— 29s 180ms/step - accuracy: 0.7418 - loss: 0.8867 - val_accuracy: 0.5720 - val_loss: 2.9210
Epoch 2/5
92/92 ————— 13s 141ms/step - accuracy: 0.9066 - loss: 0.3039 - val_accuracy: 0.5761 - val_loss: 3.2720
Epoch 3/5
92/92 ————— 13s 144ms/step - accuracy: 0.9409 - loss: 0.1836 - val_accuracy: 0.5897 - val_loss: 4.2398
Epoch 4/5
92/92 ————— 13s 137ms/step - accuracy: 0.9528 - loss: 0.1390 - val_accuracy: 0.3899 - val_loss: 7.9368
Epoch 5/5
92/92 ————— 12s 130ms/step - accuracy: 0.9704 - loss: 0.0984 - val_accuracy: 0.3071 - val_loss: 10.4191

=====
Fine-tuning learning rate: 0.0001
=====
Epoch 1/5
92/92 ————— 33s 226ms/step - accuracy: 0.6135 - loss: 1.5340 - val_accuracy: 0.8438 - val_loss: 0.5252
Epoch 2/5
92/92 ————— 13s 142ms/step - accuracy: 0.8906 - loss: 0.3982 - val_accuracy: 0.8682 - val_loss: 0.4066
Epoch 3/5
92/92 ————— 13s 142ms/step - accuracy: 0.9487 - loss: 0.2071 - val_accuracy: 0.8723 - val_loss: 0.3448
Epoch 4/5
92/92 ————— 21s 143ms/step - accuracy: 0.9718 - loss: 0.1181 - val_accuracy: 0.8981 - val_loss: 0.3065
Epoch 5/5
92/92 ————— 22s 159ms/step - accuracy: 0.9891 - loss: 0.0668 - val_accuracy: 0.8954 - val_loss: 0.3034

   Learning Rate  Best Validation Accuracy (%)  Final Validation Accuracy (%)  Final Training Accuracy (%)  Final Training Loss  Final Validation Loss  Final Training Time (s)
0         0.0010         58.967394         30.706522         97.044837         0.098368         10.419126         80.107404
1         0.0001         89.809781         89.538044         98.913044         0.066755         0.303381         101.260137

```

```

# =====
# CELL R17 - PARTIAL UNFREEZING STUDY
# =====

UNFREEZE_DIR = os.path.join(
    CHECKPOINT_DIR,
    "partial_unfreezing"
)

os.makedirs(
    UNFREEZE_DIR,
    exist_ok=True
)

unfreeze_values = [
    10,
    20,
    40
]

unfreeze_results = []

for n_layers in unfreeze_values:

    marker = os.path.join(
        UNFREEZE_DIR,
        f"layers_{n_layers}.done"
    )

    if os.path.exists(marker):

        print(
            "Skipping completed:",
            n_layers,
            "layers"
        )

        continue

    print("\n" + "=" * 60)
    print(
        "Unfreezing last",
        n_layers,
        "layers"
    )
    print("=" * 60)

    model = build_transfer_model(
        fine_tune=True,
        unfreeze_layers=n_layers,
        dropout_rate=float(best_dropout)
    )

    compile_model(
        model,
        0.0001
    )

    trainable_params = int(
        np.sum([
            np.prod(v.shape)
            for v in model.trainable_variables
        ])
    )

    start_time = time.time()

    history = model.fit(
        train_ds,
        validation_data=val_ds,
        epochs=EPOCHS,
        verbose=1
    )

```

```
training_time = (
    time.time() - start_time
)

train_acc, val_acc, loss, val_loss = (
    get_history_values(history)
)

unfreeze_results.append({
    "Unfrozen Layers":
        n_layers,

    "Trainable Parameters":
        trainable_params,

    "Best Validation Accuracy (%)":
        float(np.max(val_acc) * 100),

    "Final Validation Accuracy (%)":
        float(val_acc[-1] * 100),

    "Training Time (s)":
        training_time
})

model.save(
    os.path.join(
        UNFREEZE_DIR,
        f"layers_{n_layers}.keras"
    )
)

save_history_json(
    history,
    os.path.join(
        UNFREEZE_DIR,
        f"layers_{n_layers}_history.json"
    )
)

open(marker, "w").close()

del model

tf.keras.backend.clear_session()
gc.collect()

unfreeze_df = pd.DataFrame(
    unfreeze_results
)

unfreeze_df.to_csv(
    os.path.join(
        RESULTS_DIR,
        "partial_unfreezing_results.csv"
    ),
    index=False
)

display(
    unfreeze_df
)
```

```

=====
Unfreezing last 10 layers
=====
Epoch 1/5
92/92 ----- 28s 180ms/step - accuracy: 0.5859 - loss: 1.7342 - val_accuracy: 0.8383 - val_loss: 0.6674
Epoch 2/5
92/92 ----- 13s 142ms/step - accuracy: 0.8855 - loss: 0.4859 - val_accuracy: 0.8628 - val_loss: 0.4408
Epoch 3/5
92/92 ----- 21s 144ms/step - accuracy: 0.9310 - loss: 0.2791 - val_accuracy: 0.8764 - val_loss: 0.3725
Epoch 4/5
92/92 ----- 20s 140ms/step - accuracy: 0.9490 - loss: 0.1926 - val_accuracy: 0.8736 - val_loss: 0.3517
Epoch 5/5
92/92 ----- 13s 138ms/step - accuracy: 0.9711 - loss: 0.1358 - val_accuracy: 0.8736 - val_loss: 0.3461

=====
Unfreezing last 20 layers
=====
Epoch 1/5
92/92 ----- 29s 183ms/step - accuracy: 0.6192 - loss: 1.5928 - val_accuracy: 0.8438 - val_loss: 0.5843
Epoch 2/5
92/92 ----- 13s 144ms/step - accuracy: 0.8981 - loss: 0.4033 - val_accuracy: 0.8791 - val_loss: 0.3886
Epoch 3/5
92/92 ----- 16s 174ms/step - accuracy: 0.9497 - loss: 0.2083 - val_accuracy: 0.8859 - val_loss: 0.3655
Epoch 4/5
92/92 ----- 21s 184ms/step - accuracy: 0.9738 - loss: 0.1219 - val_accuracy: 0.8913 - val_loss: 0.3313
Epoch 5/5
92/92 ----- 16s 134ms/step - accuracy: 0.9895 - loss: 0.0722 - val_accuracy: 0.8995 - val_loss: 0.3068

=====
Unfreezing last 40 layers
=====
Epoch 1/5
92/92 ----- 34s 188ms/step - accuracy: 0.5965 - loss: 1.6289 - val_accuracy: 0.8234 - val_loss: 0.5834
Epoch 2/5
92/92 ----- 13s 145ms/step - accuracy: 0.9046 - loss: 0.3741 - val_accuracy: 0.8247 - val_loss: 0.4878
Epoch 3/5
92/92 ----- 14s 156ms/step - accuracy: 0.9579 - loss: 0.1786 - val_accuracy: 0.8832 - val_loss: 0.3941
Epoch 4/5
92/92 ----- 20s 146ms/step - accuracy: 0.9834 - loss: 0.0872 - val_accuracy: 0.8777 - val_loss: 0.3915
Epoch 5/5
92/92 ----- 13s 142ms/step - accuracy: 0.9939 - loss: 0.0467 - val_accuracy: 0.8804 - val_loss: 0.3373

  Unfrozen Layers  Trainable Parameters  Best Validation Accuracy (%)  Final Validation Accuracy (%)  Training Time (s)
0                10              779877              87.635869              87.364131              94.407162
1                20             1253477              89.945650              89.945650             104.119935
2                40             1728933              88.315219              88.043481              94.235107

```

```

# =====
# EMERGENCY CHECKPOINT - RECOVERY NOTEBOOK
# Place immediately after CELL R17
# =====

import os
import json
import pickle
import shutil
import glob
from datetime import datetime

# -----
# Correct Drive location
# -----

LAB5_DRIVE_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

RECOVERY_CHECKPOINT_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "checkpoints",
    "emergency_recovery"
)

RECOVERY_STATE_DIR = os.path.join(
    RECOVERY_CHECKPOINT_DIR,
    "python_state"
)

RECOVERY_RESULTS_DIR = os.path.join(
    RECOVERY_CHECKPOINT_DIR,
    "results_snapshot"
)

```

```

)

os.makedirs(RECOVERY_CHECKPOINT_DIR, exist_ok=True)
os.makedirs(RECOVERY_STATE_DIR, exist_ok=True)
os.makedirs(RECOVERY_RESULTS_DIR, exist_ok=True)

timestamp = datetime.now().strftime("%Y%m%d_%H%M%S")

print("=" * 70)
print("EMERGENCY CHECKPOINT — RECOVERY NOTEBOOK")
print("=" * 70)
print("Timestamp:", timestamp)
print("Drive:", LAB5_DRIVE_ROOT)

# -----
# 1. Save all DataFrames currently in memory
# -----

dataframe_names = [
    "initialization_df",
    "regularization_df",
    "optimizer_df",
    "hyperparameter_df",
    "transfer_df",
    "fine_tune_df",
    "cv_df",
    "final_df",
    "init_results_df",
    "regularization_results_df",
    "optimizer_results_df",
    "hyperparameter_results_df",
    "transfer_results_df",
    "fine_tune_results_df",
    "cv_results_df",
    "final_results"
]

saved_dataframes = []

for name in dataframe_names:

    if name in globals():

        obj = globals()[name]

        if isinstance(obj, pd.DataFrame):

            path = os.path.join(
                RECOVERY_RESULTS_DIR,
                f"{name}.csv"
            )

            obj.to_csv(
                path,
                index=False
            )

            saved_dataframes.append({
                "name": name,
                "rows": len(obj),
                "columns": list(obj.columns),
                "path": path
            })

            print(
                f"Saved DataFrame: {name} "
                f"({len(obj)} rows)"
            )

# -----
# 2. Save all training histories
# -----

history_objects = [
    "init_histories",
    "regularization_histories",
    "optimizer_histories",
    "hyperparameter_histories",
    "transfer_histories",
    "fine_tune_histories"
]

```

```

]

saved_histories = []

for container_name in history_objects:

    if container_name not in globals():
        continue

    container = globals()[container_name]

    if not isinstance(container, dict):
        continue

    history_data = {}

    for key, history in container.items():

        # Keras History object
        if hasattr(history, "history"):

            history_data[str(key)] = {
                metric: [
                    float(v)
                    for v in values
                ]
                for metric, values
                in history.history.items()
            }

        # Already converted dictionary
        elif isinstance(history, dict):

            history_data[str(key)] = history

    if history_data:

        path = os.path.join(
            RECOVERY_STATE_DIR,
            f"{container_name}.json"
        )

        with open(path, "w") as f:
            json.dump(
                history_data,
                f,
                indent=4
            )

        saved_histories.append({
            "name": container_name,
            "configurations": list(history_data.keys()),
            "path": path
        })

        print(
            f"Saved histories: {container_name}"
        )

# -----
# 3. Save individual `history` object if present
# -----

if "history" in globals():

    h = globals()["history"]

    if hasattr(h, "history"):

        path = os.path.join(
            RECOVERY_STATE_DIR,
            "latest_history.json"
        )

        with open(path, "w") as f:

            json.dump(
                {
                    key: [
                        float(v)
                        for v in values

```

```

        ]
        for key, values
        in h.history.items()
    },
    f,
    indent=4
)

print("Saved latest history.")

# -----
# 4. Save experiment variables that are easy to serialize
# -----

simple_variables = {}

for name in [
    "SEED",
    "IMG_SIZE",
    "NUM_CLASSES",
    "BATCH_SIZE",
    "EPOCHS",
    "best_lr",
    "best_batch_size",
    "best_dropout",
    "best_initialization",
    "best_optimizer"
]:

    if name in globals():

        value = globals()[name]

        try:
            json.dumps(value)

            simple_variables[name] = value

        except Exception:
            simple_variables[name] = str(value)

simple_variables_path = os.path.join(
    RECOVERY_STATE_DIR,
    "experiment_variables.json"
)

with open(
    simple_variables_path,
    "w"
) as f:

    json.dump(
        simple_variables,
        f,
        indent=4,
        default=str
    )

print("Saved experiment variables.")

# -----
# 5. Record important checkpoint files already on Drive
# -----

checkpoint_files = []

for root, dirs, files in os.walk(
    os.path.join(
        LAB5_DRIVE_ROOT,
        "checkpoints"
    )
):

    for filename in files:

        full_path = os.path.join(
            root,
            filename
        )

```

```

        checkpoint_files.append(
            os.path.relpath(
                full_path,
                LAB5_DRIVE_ROOT
            )
        )

# -----
# 6. Create recovery manifest
# -----

manifest = {

    "timestamp": timestamp,

    "notebook":
        "Lab5_Recovery_Continue_From_Checkpoint.ipynb",

    "checkpoint_type":
        "emergency_python_state_dump",

    "drive_root":
        LAB5_DRIVE_ROOT,

    "saved_dataframes":
        saved_dataframes,

    "saved_histories":
        saved_histories,

    "saved_variables":
        simple_variables,

    "existing_checkpoint_files":
        checkpoint_files,

    "runtime": {
        "python": __import__("sys").version,
        "tensorflow": tf.__version__,
        "gpu_devices": [
            str(x)
            for x in tf.config.list_physical_devices("GPU")
        ]
    }
}

manifest_path = os.path.join(
    RECOVERY_CHECKPOINT_DIR,
    f"recovery_manifest_{timestamp}.json"
)

with open(
    manifest_path,
    "w"
) as f:

    json.dump(
        manifest,
        f,
        indent=4,
        default=str
    )

# -----
# 7. Final status
# -----

print()
print("=" * 70)
print("EMERGENCY CHECKPOINT COMPLETE")
print("=" * 70)

print("DataFrames saved:",
      len(saved_dataframes))

print("History containers saved:",
      len(saved_histories))

print("Checkpoint files found:",

```

```

len(checkpoint_files))

print()
print("Manifest:")
print(manifest_path)

print()
print("Your Drive state is now backed up.")
print("=" * 70)

=====
EMERGENCY CHECKPOINT – RECOVERY NOTEBOOK
=====
Timestamp: 20260831_094050
Drive: /content/drive/MyDrive/Lab5/Lab5
Saved DataFrame: initialization_df (4 rows)
Saved DataFrame: regularization_df (4 rows)
Saved DataFrame: optimizer_df (4 rows)
Saved DataFrame: hyperparameter_df (5 rows)
Saved DataFrame: transfer_results_df (2 rows)
Saved histories: hyperparameter_histories
Saved histories: transfer_histories
Saved latest history.
Saved experiment variables.

=====
EMERGENCY CHECKPOINT COMPLETE
=====
DataFrames saved: 5
History containers saved: 2
Checkpoint files found: 105

Manifest:
/content/drive/MyDrive/Lab5/Lab5/checkpoints/emergency_recovery/recovery_manifest_20260831_094050.json

Your Drive state is now backed up.
=====

```

```

# =====
# CELL R18 – RESTORE ENVIRONMENT + SAVED RESULTS
# =====

import os
import gc
import json
import time
import random
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers, regularizers
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import (
    classification_report,
    confusion_matrix,
    precision_score,
    recall_score,
    f1_score
)

from google.colab import drive

drive.mount("/content/drive", force_remount=False)

# =====
# FIXED DRIVE PATH – NESTED LAB5
# =====

LAB5_DRIVE_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

CHECKPOINT_DIR = os.path.join(
    LAB5_DRIVE_ROOT, "checkpoints"
)

RESULTS_DIR = os.path.join(
    LAB5_DRIVE_ROOT, "results"
)

```

```

MODELS_DIR = os.path.join(
    LAB5_DRIVE_ROOT, "models"
)

OUTPUT_DIR = os.path.join(
    LAB5_DRIVE_ROOT, "outputs"
)

DATA_ROOT = os.path.join(
    LAB5_DRIVE_ROOT, "dataset"
)

print("=" * 70)
print("LAB 5 – CORRECT DRIVE ROOT")
print("=" * 70)

print("Drive root:", LAB5_DRIVE_ROOT)
print("Checkpoints:", CHECKPOINT_DIR)
print("Results:", RESULTS_DIR)
print("Dataset:", DATA_ROOT)

print("\nResults directory exists:",
      os.path.exists(RESULTS_DIR))

print("Checkpoints directory exists:",
      os.path.exists(CHECKPOINT_DIR))

print("Dataset directory exists:",
      os.path.exists(DATA_ROOT))

# -----
# Configuration
# -----

SEED = 42
IMG_SIZE = (224, 224)
NUM_CLASSES = 37
BATCH_SIZE = 32
EPOCHS = 5

random.seed(SEED)
np.random.seed(SEED)
tf.random.set_seed(SEED)

# -----
# Result loader
# -----

def load_csv(filename):
    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):
        return pd.read_csv(path)

    return None

# -----
# History helper
# -----

def get_history_dict(history):
    """
    Accept either a Keras History object or
    an already-saved dictionary.
    """

    if hasattr(history, "history"):
        return history.history

    if isinstance(history, dict):
        if "history" in history:
            return history["history"]

        return history

    return {}

```

```

# -----
# Load existing results
# -----

initialization_df = load_csv(
    "initialization_results.csv"
)

regularization_df = load_csv(
    "regularization_results.csv"
)

optimizer_df = load_csv(
    "optimizer_results.csv"
)

hyperparameter_df = load_csv(
    "hyperparameter_results.csv"
)

transfer_df = load_csv(
    "transfer_learning_results.csv"
)

fine_tune_df = load_csv(
    "correct_fine_tuning_results.csv"
)

cv_df = load_csv(
    "cv_results.csv"
)

final_df = load_csv(
    "final_model_results.csv"
)

bn_results_df = load_csv(
    "batch_normalization_results.csv"
)

# -----
# Display recovery status
# -----

print("=" * 70)
print("LAB 5 - RECOVERY ENVIRONMENT")
print("=" * 70)

print("Drive:", LAB5_DRIVE_ROOT)
print("TensorFlow:", tf.__version__)
print(
    "GPU:",
    tf.config.list_physical_devices("GPU")
)

print()
print("Saved results:")

for label, df in [
    ("Initialization", initialization_df),
    ("Regularization", regularization_df),
    ("Batch Normalization", bn_results_df),
    ("Optimizer", optimizer_df),
    ("Hyperparameters", hyperparameter_df),
    ("Transfer Learning", transfer_df),
    ("Fine-Tuning", fine_tune_df),
    ("Cross-Validation", cv_df),
    ("Final Model", final_df)
]:

    if df is None:
        print(f" {label}: not found")
    else:
        print(
            f" {label}: {len(df)} row(s)"
        )

print()
print("Ready to continue from R17.")

```

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)
=====
LAB 5 – CORRECT DRIVE ROOT
=====
Drive root: /content/drive/MyDrive/Lab5/Lab5
Checkpoints: /content/drive/MyDrive/Lab5/Lab5/checkpoints
Results: /content/drive/MyDrive/Lab5/Lab5/results
Dataset: /content/drive/MyDrive/Lab5/Lab5/dataset

Results directory exists: True
Checkpoints directory exists: True
Dataset directory exists: True
=====
LAB 5 – RECOVERY ENVIRONMENT
=====
Drive: /content/drive/MyDrive/Lab5/Lab5
TensorFlow: 2.20.0
GPU: []

Saved results:
  Initialization: 4 row(s)
  Regularization: 4 row(s)
  Batch Normalization: 1 row(s)
  Optimizer: 4 row(s)
  Hyperparameters: 5 row(s)
  Transfer Learning: 2 row(s)
  Fine-Tuning: not found
  Cross-Validation: not found
  Final Model: not found

Ready to continue from R17.

```

```

# =====
# CELL R19 – FINE-TUNING SETUP
# =====

import os
import gc
import time
import json
import numpy as np
import pandas as pd
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input

print("=" * 70)
print("LAB 5 – FINE-TUNING SETUP")
print("=" * 70)

# =====
# 1. RECREATE DATA PIPELINE
# =====

DATA_ROOT = os.path.join(
    LAB5_DRIVE_ROOT,
    "dataset"
)

IMAGES_DIR = os.path.join(
    DATA_ROOT,
    "images"
)

ANNOTATIONS_DIR = os.path.join(
    DATA_ROOT,
    "annotations"
)

LIST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "list.txt"
)

TRAINVAL_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "trainval.txt"
)

```

```

TEST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "test.txt"
)

if not os.path.exists(IMAGES_DIR):
    raise FileNotFoundError(
        f"Images directory not found:\n{IMAGES_DIR}"
    )

if not os.path.exists(ANNOTATIONS_DIR):
    raise FileNotFoundError(
        f"Annotations directory not found:\n{ANNOTATIONS_DIR}"
    )

def read_split_file(filename):

    names = []

    with open(filename, "r") as f:

        for line in f:

            line = line.strip()

            if not line or line.startswith("#"):
                continue

            names.append(line.split()[0])

    return set(names)

# -----
# Read all images
# -----

samples = []

with open(LIST_FILE, "r") as f:

    for line in f:

        line = line.strip()

        if not line or line.startswith("#"):
            continue

        image_name = line.split()[0]

        image_path = os.path.join(
            IMAGES_DIR,
            image_name + ".jpg"
        )

        if os.path.exists(image_path):
            samples.append(image_path)

# -----
# Class names
# -----

class_names = sorted(
    set(
        os.path.basename(path)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
        for path in samples
    )
)

class_to_index = {
    name: i
    for i, name in enumerate(class_names)
}

# -----

```

```

# Image paths + labels
# -----

image_paths = []
labels = []

for path in samples:

    breed = (
        os.path.basename(path)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
    )

    if breed in class_to_index:

        image_paths.append(path)
        labels.append(
            class_to_index[breed]
        )

image_paths = np.array(image_paths)

labels = np.array(
    labels,
    dtype=np.int32
)

# =====
# 2. RECREATE TRAIN / VALIDATION / TEST SPLIT
# =====

trainval_names = read_split_file(
    TRAINVAL_FILE
)

test_names = read_split_file(
    TEST_FILE
)

trainval_paths = []
trainval_labels = []

test_paths = []
test_labels = []

for path, label in zip(
    image_paths,
    labels
):

    name = (
        os.path.basename(path)
        .replace(".jpg", "")
    )

    if name in trainval_names:

        trainval_paths.append(path)
        trainval_labels.append(label)

    elif name in test_names:

        test_paths.append(path)
        test_labels.append(label)

trainval_paths = np.array(
    trainval_paths
)

trainval_labels = np.array(
    trainval_labels,
    dtype=np.int32
)

test_paths = np.array(
    test_paths
)

```

```

    )

    test_labels = np.array(
        test_labels,
        dtype=np.int32
    )

    # -----
    # Deterministic 80/20 train-validation split
    # -----

    rng = np.random.default_rng(
        SEED
    )

    indices = np.arange(
        len(trainval_paths)
    )

    rng.shuffle(indices)

    val_size = int(
        0.20 * len(indices)
    )

    val_indices = indices[:val_size]

    train_indices = indices[val_size:]

    train_paths = trainval_paths[
        train_indices
    ]

    train_labels = trainval_labels[
        train_indices
    ]

    val_paths = trainval_paths[
        val_indices
    ]

    val_labels = trainval_labels[
        val_indices
    ]

    print("Dataset reconstructed.")
    print("Classes:", len(class_names))
    print("Train:", len(train_paths))
    print("Validation:", len(val_paths))
    print("Test:", len(test_paths))

    # =====
    # 3. DATASET FUNCTIONS
    # =====

    def load_image(path, label):

        image = tf.io.read_file(path)

        image = tf.image.decode_jpeg(
            image,
            channels=3
        )

        image = tf.image.resize(
            image,
            IMG_SIZE
        )

        image = tf.cast(
            image,
            tf.float32
        )

        image = preprocess_input(
            image
        )

```

```

        return image, label

def make_dataset(
    paths,
    labels,
    batch_size=32,
    shuffle=False
):

    ds = tf.data.Dataset.from_tensor_slices(
        (paths, labels)
    )

    if shuffle:

        ds = ds.shuffle(
            len(paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        batch_size
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

# =====
# 4. LOAD BEST HYPERPARAMETERS
# =====

if (
    "hyperparameter_df" not in globals()
    or hyperparameter_df is None
):

    raise RuntimeError(
        "hyperparameter_df is not loaded. "
        "Run the results-loading cell first."
    )

hp = hyperparameter_df.copy()

hp[
    "Best Validation Accuracy (%)"
] = pd.to_numeric(
    hp["Best Validation Accuracy (%)"],
    errors="coerce"
)

hp = hp.dropna(
    subset=[
        "Best Validation Accuracy (%)"
    ]
)

def get_best_hp(
    parameter_name,
    default
):

    subset = hp[
        hp["Hyperparameter"]
        .astype(str)
        .str.strip()
        .str.lower()
        ==
        parameter_name.lower()
    ]

```

```

    if len(subset) == 0:

        print(
            f"No saved result for "
            f"{parameter_name}; using {default}"
        )

        return default

    return subset.loc[
        subset[
            "Best Validation Accuracy (%)"
        ].idxmax(),
        "Value"
    ]

best_lr = float(
    get_best_hp(
        "Learning Rate",
        0.001
    )
)

best_batch_size = int(
    get_best_hp(
        "Batch Size",
        32
    )
)

best_dropout = float(
    get_best_hp(
        "Dropout",
        0.0
    )
)

print("\nSelected hyperparameters:")
print("Learning rate :", best_lr)
print("Batch size    :", best_batch_size)
print("Dropout       :", best_dropout)

# =====
# 5. FINE-TUNING SETTINGS
# =====

FINE_TUNE_LR = 1e-5
FINE_TUNE_EPOCHS = 5

# CPU-friendly partial fine-tuning
UNFREEZE_LAST_N = 20

print("\nFine-tuning settings:")
print("Learning rate :", FINE_TUNE_LR)
print("Epochs       :", FINE_TUNE_EPOCHS)
print(
    "Unfreeze last :",
    UNFREEZE_LAST_N,
    "layers"
)

# =====
# 6. CREATE DATASETS
# =====

fine_train_ds = make_dataset(
    train_paths,
    train_labels,
    batch_size=best_batch_size,
    shuffle=True
)

fine_val_ds = make_dataset(
    val_paths,
    val_labels,
    batch_size=best_batch_size,

```

```

        shuffle=False
    )

# =====
# 7. BUILD MOBILE NET V2
# =====

base_model = MobileNetV2(
    input_shape=(224, 224, 3),
    include_top=False,
    weights="imagenet"
)

# Enable fine-tuning
base_model.trainable = True

# Freeze earlier layers
for layer in base_model.layers[
    :-UNFREEZE_LAST_N
]:

    layer.trainable = False

# Keep BatchNorm frozen
for layer in base_model.layers:

    if isinstance(
        layer,
        layers.BatchNormalization
    ):

        layer.trainable = False

trainable_layers = sum(
    layer.trainable
    for layer in base_model.layers
)

# =====
# 8. BUILD CLASSIFIER
# =====

inputs = keras.Input(
    shape=(224, 224, 3)
)

x = base_model(
    inputs,
    training=False
)

x = layers.GlobalAveragePooling2D()(x)

if best_dropout > 0:

    x = layers.Dropout(
        best_dropout
    )(x)

outputs = layers.Dense(
    NUM_CLASSES,
    activation="softmax"
)(x)

fine_tune_model = keras.Model(
    inputs,
    outputs
)

# =====
# 9. COMPILE
# =====

```

```

fine_tune_model.compile(
    optimizer=keras.optimizers.Adam(
        learning_rate=FINE_TUNE_LR
    ),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

print("\n" + "=" * 70)
print("FINE-TUNING MODEL READY")
print("=" * 70)

print(
    "Total parameters:",
    fine_tune_model.count_params()
)

print(
    "Trainable layers:",
    trainable_layers
)

print(
    "Trainable parameters:",
    sum(
        np.prod(v.shape)
        for v in fine_tune_model.trainable_weights
    )
)

```

```

=====
LAB 5 – FINE-TUNING SETUP
=====
Dataset reconstructed.
Classes: 37
Train: 2944
Validation: 736
Test: 3669
No saved result for Batch Size; using 32

Selected hyperparameters:
Learning rate : 0.001
Batch size    : 32
Dropout       : 0.25

Fine-tuning settings:
Learning rate : 1e-05
Epochs       : 5
Unfreeze last : 20 layers
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\_v2/mobilenet\_v2\_weights\_tf\_dim\_9406464/9406464 — 0s 0us/step

=====
FINE-TUNING MODEL READY
=====
Total parameters: 2305381
Trainable layers: 13
Trainable parameters: 1242277

```

```

# =====
# CELL R20 – FOLD DATASET HELPER
# =====

def create_fold_dataset_from_arrays(
    paths,
    labels,
    indices,
    batch_size,
    shuffle=False
):
    selected_paths = paths[
        indices
    ]

    selected_labels = labels[
        indices
    ]

    ds = tf.data.Dataset.from_tensor_slices(

```

```

        (
            selected_paths,
            selected_labels
        )
    )

if shuffle:

    ds = ds.shuffle(
        len(selected_paths),
        seed=SEED,
        reshuffle_each_iteration=True
    )

ds = ds.map(
    load_image,
    num_parallel_calls=tf.data.AUTOTUNE
)

ds = ds.batch(
    batch_size
)

ds = ds.prefetch(
    tf.data.AUTOTUNE
)

return ds

```

```

# =====
# CELL R21 - CV MODEL BUILDER
# =====

def build_cv_model(
    initialization="normal",
    fine_tune=False,
    dropout=0.0
):

    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    if fine_tune:

        base_model.trainable = True

        for layer in base_model.layers[:-20]:
            layer.trainable = False

    else:

        base_model.trainable = False

    inputs = keras.Input(
        shape=(224, 224, 3)
    )

    x = base_model(
        inputs,
        training=False
    )

    x = layers.GlobalAveragePooling2D()(x)

    if dropout > 0:

        x = layers.Dropout(
            dropout
        )(x)

    if initialization == "zero":

```

```

        outputs = layers.Dense(
            NUM_CLASSES,
            activation="softmax",
            kernel_initializer=keras.initializers.Zeros()
        )(x)

    else:

        outputs = layers.Dense(
            NUM_CLASSES,
            activation="softmax"
        )(x)

    return keras.Model(
        inputs,
        outputs
    )

```

```

# =====
# CELL R22 – RECREATE DATASETS FOR CONTINUATION
# =====

print("=" * 70)
print("RECREATING DATASETS")
print("=" * 70)

def load_image(path, label):

    image = tf.io.read_file(path)

    image = tf.image.decode_jpeg(
        image,
        channels=3
    )

    image = tf.image.resize(
        image,
        IMG_SIZE
    )

    image = tf.cast(
        image,
        tf.float32
    )

    image = preprocess_input(image)

    return image, label

def make_dataset(
    paths,
    labels,
    batch_size=32,
    shuffle=False
):

    ds = tf.data.Dataset.from_tensor_slices(
        (
            paths,
            labels
        )
    )

    if shuffle:

        ds = ds.shuffle(
            len(paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        batch_size
    )

```

```

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

# -----
# Standard datasets
# -----

train_ds = make_dataset(
    train_paths,
    train_labels,
    batch_size=BATCH_SIZE,
    shuffle=True
)

val_ds = make_dataset(
    val_paths,
    val_labels,
    batch_size=BATCH_SIZE,
    shuffle=False
)

test_ds = make_dataset(
    test_paths,
    test_labels,
    batch_size=BATCH_SIZE,
    shuffle=False
)

print("Datasets recreated successfully.")
print("Train:", len(train_paths))
print("Validation:", len(val_paths))
print("Test:", len(test_paths))
print("Batch size:", BATCH_SIZE)

```

```

=====
RECREATING DATASETS
=====
Datasets recreated successfully.
Train: 2944
Validation: 736
Test: 3669
Batch size: 32

```

```

# =====
# CELL R23 - FINE-TUNING EXPERIMENT
# =====

import os
import gc
import json
import time
import numpy as np
import pandas as pd
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2

# -----
# Paths
# -----

FINE_TUNE_DIR = os.path.join(
    CHECKPOINT_DIR,
    "fine_tuning_final"
)

os.makedirs(
    FINE_TUNE_DIR,
    exist_ok=True
)

FINE_TUNE_RESULTS_PATH = os.path.join(
    RESULTS_DIR,

```

```

    "correct_fine_tuning_results.csv"
)

FINE_TUNE_MODEL_PATH = os.path.join(
    FINE_TUNE_DIR,
    "fine_tuned_model.keras"
)

FINE_TUNE_WEIGHTS_PATH = os.path.join(
    FINE_TUNE_DIR,
    "best_model.weights.h5"
)

FINE_TUNE_HISTORY_PATH = os.path.join(
    FINE_TUNE_DIR,
    "history.json"
)

FINE_TUNE_BACKUP_DIR = os.path.join(
    FINE_TUNE_DIR,
    "training_backup"
)

# =====
# 1. LOAD COMPLETED EXPERIMENT IF AVAILABLE
# =====

if (
    os.path.exists(FINE_TUNE_MODEL_PATH)
    and os.path.exists(FINE_TUNE_HISTORY_PATH)
):

    print("=" * 70)
    print("LOADING COMPLETED FINE-TUNING EXPERIMENT")
    print("=" * 70)

    # Load complete model
    fine_model = keras.models.load_model(
        FINE_TUNE_MODEL_PATH
    )

    # Load history
    with open(
        FINE_TUNE_HISTORY_PATH,
        "r"
    ) as f:

        history_dict = json.load(f)

    fine_history = keras.callbacks.History()
    fine_history.history = history_dict

    # Load saved results if available
    if os.path.exists(FINE_TUNE_RESULTS_PATH):

        fine_tune_results = pd.read_csv(
            FINE_TUNE_RESULTS_PATH
        )

    else:

        h = fine_history.history

        fine_tune_results = pd.DataFrame([
            {
                "Method":
                    "Fine-Tuning",

                "Learning Rate":
                    0.0001,

                "Trainable Parameters":
                    int(np.sum([
                        np.prod(v.shape)
                        for v in fine_model.trainable_weights
                    ])),

                "Total Parameters":
                    int(fine_model.count_params()),

                "Best Validation Accuracy (%)":

```

```

        max(h["val_accuracy"]) * 100,

        "Final Training Accuracy (%)":
            h["accuracy"][-1] * 100,

        "Final Validation Accuracy (%)":
            h["val_accuracy"][-1] * 100,

        "Final Training Loss":
            h["loss"][-1],

        "Final Validation Loss":
            h["val_loss"][-1],

        "Training Time (s)":
            np.nan

    })

    fine_tune_results.to_csv(
        FINE_TUNE_RESULTS_PATH,
        index=False
    )

    print("Fine-tuned model loaded.")
    print(
        "Best validation accuracy:",
        f"{max(history_dict['val_accuracy']) * 100:.2f}%"
    )

    display(
        fine_tune_results
    )

# =====
# 2. OTHERWISE BUILD / RESUME FINE-TUNING
# =====

else:

    print("=" * 70)
    print("FINE-TUNING")
    print("=" * 70)

    # -----
    # Build MobileNetV2
    # -----

    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    base_model.trainable = True

    # Freeze all except last 20 layers
    for layer in base_model.layers[:-20]:

        layer.trainable = False

    # Keep BatchNorm frozen
    for layer in base_model.layers:

        if isinstance(
            layer,
            layers.BatchNormalization
        ):

            layer.trainable = False

    # -----
    # Build classifier
    # -----

    inputs = keras.Input(

```

```

        shape=(224, 224, 3)
    )

    x = base_model(
        inputs,
        training=False
    )

    x = layers.GlobalAveragePooling2D()(x)

    x = layers.Dropout(
        float(best_dropout)
    )(x)

    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax"
    )(x)

    fine_model = keras.Model(
        inputs,
        outputs
    )

    # -----
    # Compile
    # -----

    fine_model.compile(

        optimizer=keras.optimizers.RMSprop(
            learning_rate=0.0001
        ),

        loss="sparse_categorical_crossentropy",

        metrics=["accuracy"]
    )

    # -----
    # Resume weights if available
    # -----

    if os.path.exists(
        FINE_TUNE_WEIGHTS_PATH
    ):

        print(
            "Existing fine-tuning weights found."
        )

        fine_model.load_weights(
            FINE_TUNE_WEIGHTS_PATH
        )

        print(
            "Weights loaded successfully."
        )

    else:

        print(
            "No previous fine-tuning weights found."
        )

        print(
            "Starting fine-tuning from pretrained MobileNetV2."
        )

    # -----
    # Parameter counts
    # -----

    trainable_params = int(
        np.sum([
            np.prod(v.shape)
            for v in fine_model.trainable_weights
        ])
    )

```

```
total_params = int(
    fine_model.count_params()
)

print(
    "Trainable parameters:",
    trainable_params
)

print(
    "Total parameters:",
    total_params
)

# -----
# Callbacks
# -----

fine_callbacks = [

    keras.callbacks.ModelCheckpoint(

        FINE_TUNE_WEIGHTS_PATH,

        monitor="val_accuracy",

        mode="max",

        save_best_only=True,

        save_weights_only=True,

        verbose=1
    ),

    keras.callbacks.BackupAndRestore(

        backup_dir=FINE_TUNE_BACKUP_DIR,

        delete_checkpoint=False
    )
]

# -----
# Train
# -----

start_time = time.time()

fine_history = fine_model.fit(

    train_ds,

    validation_data=val_ds,

    epochs=EPOCHS,

    verbose=1,

    callbacks=fine_callbacks
)

fine_tune_time = (
    time.time() - start_time
)

# -----
# Validate history
# -----

h = fine_history.history

required_metrics = [
    "loss",
    "accuracy",
    "val_loss",
    "val_accuracy"
]
```

```

missing = [
    metric
    for metric in required_metrics
    if metric not in h
]

if missing:

    raise RuntimeError(
        "Fine-tuning did not record required "
        f"metrics: {missing}. "
        f"Available: {list(h.keys())}"
    )

# -----
# Save history
# -----

with open(
    FINE_TUNE_HISTORY_PATH,
    "w"
) as f:

    json.dump(

        {

            key: [
                float(x)
                for x in values
            ]

            for key, values in h.items()
        },

        f,

        indent=2
    )

# -----
# Save complete model
# -----

fine_model.save(
    FINE_TUNE_MODEL_PATH
)

# -----
# Results
# -----

fine_tune_results = pd.DataFrame([

    "Method":
        "Fine-Tuning",

    "Learning Rate":
        0.0001,

    "Trainable Parameters":
        trainable_params,

    "Total Parameters":
        total_params,

    "Best Validation Accuracy (%)":
        max(h["val_accuracy"]) * 100,

    "Final Training Accuracy (%)":
        h["accuracy"][-1] * 100,

    "Final Validation Accuracy (%)":
        h["val_accuracy"][-1] * 100,

    "Final Training Loss":
        h["loss"][-1],

    "Final Validation Loss":

```

```

        h["val_loss"][-1],

        "Training Time (s)":
            fine_tune_time

    })

    fine_tune_results.to_csv(
        FINE_TUNE_RESULTS_PATH,
        index=False
    )

    display(
        fine_tune_results
    )

    print()
    print("Fine-tuning completed and saved.")
    print(
        "Model:",
        FINE_TUNE_MODEL_PATH
    )

    print(
        "History:",
        FINE_TUNE_HISTORY_PATH
    )

    print(
        "Results:",
        FINE_TUNE_RESULTS_PATH
    )

# =====
# IMPORTANT
# =====
# Do NOT delete fine_model here.
#
# R24 needs fine_model and fine_history.
# They must remain in memory for the next cells.
# =====

print()
print("=" * 70)
print("R23 READY")
print("=" * 70)
print("fine_model exists:", "fine_model" in globals())
print("fine_history exists:", "fine_history" in globals())

```

```

=====
LOADING COMPLETED FINE-TUNING EXPERIMENT
=====
Fine-tuned model loaded.
Best validation accuracy: 89.27%

```

	Method	Learning Rate	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Fine-Tuning	0.0001	1242277	2305381	89.266306	96.705163	89.266306	0.093895	0.303505	1297.457243

```

=====
R23 READY
=====
fine_model exists: True

```

```

# =====
# CELL R24 - FINE-TUNING RESULTS
# =====

import os
import json
import pandas as pd
import numpy as np

```

```

# -----
# Check that fine-tuning completed
# -----

if "fine_history" not in globals():
    raise RuntimeError(
        "fine_history was not found. "
        "Run CELL R23 successfully before continuing."
    )

if "fine_model" not in globals():
    raise RuntimeError(
        "fine_model was not found. "
        "Run CELL R23 successfully before continuing."
    )

# -----
# Extract history safely
# -----

fine_hist = fine_history.history

print("=" * 70)
print("FINE-TUNING RESULTS")
print("=" * 70)

print("Available metrics:")
print(list(fine_hist.keys()))

# -----
# Handle validation metric
# -----

val_accuracy_key = None

for key in ["val_accuracy", "val_sparse_categorical_accuracy"]:
    if key in fine_hist:
        val_accuracy_key = key
        break

if val_accuracy_key is None:
    raise KeyError(
        "No validation accuracy metric was found. "
        f"Available metrics: {list(fine_hist.keys())}"
    )

train_accuracy_key = None

for key in ["accuracy", "sparse_categorical_accuracy"]:
    if key in fine_hist:
        train_accuracy_key = key
        break

# -----
# Calculate results
# -----

best_val_accuracy = (
    max(fine_hist[val_accuracy_key]) * 100
)

final_val_accuracy = (
    fine_hist[val_accuracy_key][-1] * 100
)

final_train_accuracy = None

if train_accuracy_key is not None:
    final_train_accuracy = (
        fine_hist[train_accuracy_key][-1] * 100
    )

final_train_loss = fine_hist["loss"][-1]

final_val_loss = fine_hist["val_loss"][-1]

```

```

# -----
# Create result table
# -----

fine_tuning_results = pd.DataFrame([
    {
        "Model": "Fine-Tuned MobileNetV2",
        "Best Validation Accuracy (%)":
            best_val_accuracy,
        "Final Training Accuracy (%)":
            final_train_accuracy,
        "Final Validation Accuracy (%)":
            final_val_accuracy,
        "Final Training Loss":
            final_train_loss,
        "Final Validation Loss":
            final_val_loss,
        "Number of Parameters":
            fine_model.count_params()
    }
])

display(fine_tuning_results)

# -----
# Save results
# -----

FINE_TUNE_RESULTS_CSV = os.path.join(
    RESULTS_DIR,
    "fine_tuning_results.csv"
)

fine_tuning_results.to_csv(
    FINE_TUNE_RESULTS_CSV,
    index=False
)

# -----
# Save history
# -----

FINE_TUNE_HISTORY_JSON = os.path.join(
    RESULTS_DIR,
    "fine_tuning_history.json"
)

with open(
    FINE_TUNE_HISTORY_JSON,
    "w"
) as f:
    json.dump(
        {
            key: [
                float(value)
                for value in values
            ]
            for key, values in fine_hist.items()
        },
        f,
        indent=2
    )

print()
print("Fine-tuning results saved.")
print("Results:", FINE_TUNE_RESULTS_CSV)
print("History:", FINE_TUNE_HISTORY_JSON)

```

```
=====
```

FINE-TUNING RESULTS

```
=====
```

Available metrics:

```
['accuracy', 'loss', 'val_accuracy', 'val_loss']
```

	Model	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Number of Parameters
0	Fine-Tuned MobileNetV2	89.266306	96.705163	89.266306	0.093895	0.303505	2305381

Fine-tuning results saved.

Results: /content/drive/MyDrive/Lab5/Lab5/results/fine_tuning_results.csv

History: /content/drive/MyDrive/Lab5/Lab5/results/fine_tuning_history.json

```
# =====
```

```
# CELL R25 - CROSS-VALIDATION SETUP
```

```
# =====
```

```
import os
import gc
import json
import time
import numpy as np
import pandas as pd
import tensorflow as tf
```

```
from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
```

```
# -----
# Cross-validation paths
# -----
```

```
CV_DIR = os.path.join(
    CHECKPOINT_DIR,
    "cross_validation"
)
```

```
os.makedirs(
    CV_DIR,
    exist_ok=True
)
```

```
CV_RESULTS_PATH = os.path.join(
    RESULTS_DIR,
    "cv_results.csv"
)
```

```
CV_SUMMARY_PATH = os.path.join(
    RESULTS_DIR,
    "cv_summary.csv"
)
```

```
# -----
# Load existing CV results if present
# -----
```

```
if os.path.exists(CV_RESULTS_PATH):
```

```
    cv_results_df = pd.read_csv(
        CV_RESULTS_PATH
    )
```

```
    cv_results = cv_results_df.to_dict(
        "records"
    )
```

```
    print(
        f"Loaded {len(cv_results)} existing CV result(s)."
    )
```

```
else:
```

```
    cv_results = []
```

```

    print(
        "No previous CV results found."
    )

# -----
# Create 5 stratified folds
# -----

from sklearn.model_selection import StratifiedKFold

skf = StratifiedKFold(
    n_splits=5,
    shuffle=True,
    random_state=SEED
)

folds = list(
    skf.split(
        train_paths,
        train_labels
    )
)

print()
print("=" * 70)
print("5-FOLD CROSS-VALIDATION SETUP")
print("=" * 70)

print(
    "Number of folds:",
    len(folds)
)

for i, (train_idx, val_idx) in enumerate(
    folds,
    start=1
):

    print(
        f"Fold {i}: "
        f"Train = {len(train_idx)}, "
        f"Validation = {len(val_idx)}"
    )

```

No previous CV results found.

```

=====
5-FOLD CROSS-VALIDATION SETUP
=====
Number of folds: 5
Fold 1: Train = 2355, Validation = 589
Fold 2: Train = 2355, Validation = 589
Fold 3: Train = 2355, Validation = 589
Fold 4: Train = 2355, Validation = 589
Fold 5: Train = 2356, Validation = 588

```

```

# =====
# CELL R26 - RESUMABLE 5-FOLD CROSS-VALIDATION
# =====

CV_CONFIGS = {

    "Fine-Tuned": {
        "initialization": "normal",
        "fine_tune": True,
        "dropout": float(best_dropout),
        "learning_rate": 0.0001,
        "batch_size": int(best_batch_size)
    },

    "Baseline": {
        "initialization": "normal",
        "fine_tune": False,
        "dropout": 0.0,
        "learning_rate": float(best_lr),
        "batch_size": int(best_batch_size)
    }
}

# -----

```

```

# Helper: create fold dataset
# -----

def create_cv_dataset(
    paths,
    labels,
    indices,
    batch_size,
    shuffle=False
):

    selected_paths = paths[indices]
    selected_labels = labels[indices]

    ds = tf.data.Dataset.from_tensor_slices(
        (
            selected_paths,
            selected_labels
        )
    )

    if shuffle:

        ds = ds.shuffle(
            len(selected_paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        batch_size
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

# -----
# Helper: CV model
# -----

def build_cv_model(
    initialization="normal",
    fine_tune=False,
    dropout=0.0
):

    base_model = MobileNetV2(
        input_shape=(224, 224, 3),
        include_top=False,
        weights="imagenet"
    )

    if fine_tune:

        base_model.trainable = True

        for layer in base_model.layers[:-20]:
            layer.trainable = False

        # Keep BatchNorm frozen
        for layer in base_model.layers:

            if isinstance(
                layer,
                layers.BatchNormalization
            ):
                layer.trainable = False

    else:

        base_model.trainable = False

```

```

inputs = keras.Input(
    shape=(224, 224, 3)
)

x = base_model(
    inputs,
    training=False
)

x = layers.GlobalAveragePooling2D()(x)

if dropout > 0:

    x = layers.Dropout(
        dropout
    )(x)

if initialization == "zero":

    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax",
        kernel_initializer=keras.initializers.Zeros()
    )(x)

else:

    outputs = layers.Dense(
        NUM_CLASSES,
        activation="softmax"
    )(x)

return keras.Model(
    inputs,
    outputs
)

# -----
# Helper: check completed fold
# -----

def cv_fold_completed(
    configuration,
    fold
):

    return any(
        str(row["Configuration"]) == str(configuration)
        and int(row["Fold"]) == int(fold)
        for row in cv_results
    )

# -----
# Start CV
# -----

print("=" * 70)
print("5-FOLD CROSS-VALIDATION")
print("=" * 70)

for config_name, config in CV_CONFIGS.items():

    print()
    print("#" * 70)
    print(config_name)
    print("#" * 70)

    for fold_number, (
        train_idx,
        val_idx
    ) in enumerate(
        folds,
        start=1
    ):

        # -----
        # Skip completed folds
        # -----

```

```

if cv_fold_completed(
    config_name,
    fold_number
):

    print(
        f"Skipping {config_name} "
        f"Fold {fold_number} - already saved."
    )

    continue

print()
print(
    f"{config_name} | "
    f"Fold {fold_number}/5"
)

# -----
# Fold datasets
# -----

fold_train_ds = create_cv_dataset(
    train_paths,
    train_labels,
    train_idx,
    batch_size=config["batch_size"],
    shuffle=True
)

fold_val_ds = create_cv_dataset(
    train_paths,
    train_labels,
    val_idx,
    batch_size=config["batch_size"],
    shuffle=False
)

# -----
# Build model
# -----

model = build_cv_model(
    initialization=config[
        "initialization"
    ],
    fine_tune=config[
        "fine_tune"
    ],
    dropout=config[
        "dropout"
    ]
)

# -----
# Compile
# -----

model.compile(

    optimizer=keras.optimizers.RMSprop(
        learning_rate=config[
            "learning_rate"
        ]
    ),

    loss="sparse_categorical_crossentropy",

    metrics=["accuracy"]
)

# -----
# Fold checkpoint directory
# -----

fold_dir = os.path.join(
    CV_DIR,

```

```

        config_name.lower().replace(
            " ",
            "_",
        ),
        f"fold_{fold_number}"
    )

    os.makedirs(
        fold_dir,
        exist_ok=True
    )

# -----
# Checkpoint callback
# -----

checkpoint_callback = (
    keras.callbacks.ModelCheckpoint(

        os.path.join(
            fold_dir,
            "best_model.weights.h5"
        ),

        monitor="val_accuracy",

        mode="max",

        save_best_only=True,

        save_weights_only=True,

        verbose=1
    )
)

# -----
# Train
# -----

start_time = time.time()

history = model.fit(

    fold_train_ds,

    validation_data=fold_val_ds,

    epochs=EPOCHS,

    verbose=1,

    callbacks=[
        checkpoint_callback
    ]
)

training_time = (
    time.time() - start_time
)

# -----
# Extract metrics
# -----

h = history.history

if "val_accuracy" not in h:

    raise RuntimeError(
        "val_accuracy was not recorded. "
        f"Available metrics: {list(h.keys())}"
    )

best_val = (
    max(h["val_accuracy"]) * 100
)

```

```
final_val = (
    h["val_accuracy"][-1] * 100
)

# -----
# Save history
# -----

history_path = os.path.join(
    fold_dir,
    "history.json"
)

with open(
    history_path,
    "w"
) as f:

    json.dump(
        {
            key: [
                float(x)
                for x in values
            ]
            for key, values in h.items()
        },
        f,
        indent=2
    )

# -----
# Save complete fold model
# -----

model.save(
    os.path.join(
        fold_dir,
        "model.keras"
    )
)

# -----
# Result
# -----

row = {

    "Configuration":
        config_name,

    "Fold":
        fold_number,

    "Best Validation Accuracy (%)":
        best_val,

    "Final Validation Accuracy (%)":
        final_val,

    "Training Time (s)":
        training_time
}

cv_results.append(
    row
)

# -----
# SAVE RESULT IMMEDIATELY
# -----

cv_results_df = pd.DataFrame(
    cv_results
)

cv_results_df.to_csv(
    CV_RESULTS_PATH,
```

```
        index=False
    )

    print()
    print(
        f"Saved {config_name} "
        f"Fold {fold_number}"
    )

    print(
        f"Best validation accuracy: "
        f"{best_val:.2f}%"
    )

    # -----
    # Free memory
    # -----

    del model
    del fold_train_ds
    del fold_val_ds
    del history

    tf.keras.backend.clear_session()
    gc.collect()

    # -----
    # Final CV result table
    # -----

    cv_results_df = pd.DataFrame(
        cv_results
    )

    print()
    print("=" * 70)
    print("CROSS-VALIDATION RESULTS")
    print("=" * 70)

    display(
        cv_results_df
    )
```

```
=====
5-FOLD CROSS-VALIDATION
=====

#####
Fine-Tuned
#####

Fine-Tuned | Fold 1/5
Epoch 1/5
74/74 ----- 0s 2s/step - accuracy: 0.4003 - loss: 2.3849
Epoch 1: val_accuracy improved from None to 0.79966, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_va

Epoch 1: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_1/best_mod
74/74 ----- 165s 2s/step - accuracy: 0.5856 - loss: 1.5911 - val_accuracy: 0.7997 - val_loss: 0.6617
Epoch 2/5
74/74 ----- 0s 2s/step - accuracy: 0.8582 - loss: 0.4977
Epoch 2: val_accuracy improved from 0.79966 to 0.85229, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross

Epoch 2: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_1/best_mod
74/74 ----- 150s 2s/step - accuracy: 0.8705 - loss: 0.4510 - val_accuracy: 0.8523 - val_loss: 0.4713
Epoch 3/5
74/74 ----- 0s 2s/step - accuracy: 0.9303 - loss: 0.2464
Epoch 3: val_accuracy improved from 0.85229 to 0.87776, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross

Epoch 3: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_1/best_mod
74/74 ----- 176s 2s/step - accuracy: 0.9291 - loss: 0.2529 - val_accuracy: 0.8778 - val_loss: 0.3803
Epoch 4/5
74/74 ----- 0s 2s/step - accuracy: 0.9590 - loss: 0.1517
Epoch 4: val_accuracy did not improve from 0.87776
74/74 ----- 165s 2s/step - accuracy: 0.9537 - loss: 0.1543 - val_accuracy: 0.8625 - val_loss: 0.4297
Epoch 5/5
74/74 ----- 0s 2s/step - accuracy: 0.9748 - loss: 0.0923
Epoch 5: val_accuracy did not improve from 0.87776
74/74 ----- 199s 2s/step - accuracy: 0.9720 - loss: 0.0975 - val_accuracy: 0.8710 - val_loss: 0.3948

Saved Fine-Tuned Fold 1
Best validation accuracy: 87.78%

Fine-Tuned | Fold 2/5
Epoch 1/5
74/74 ----- 0s 2s/step - accuracy: 0.3989 - loss: 2.4037
Epoch 1: val_accuracy improved from None to 0.83362, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_va

Epoch 1: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_2/best_mod
74/74 ----- 179s 2s/step - accuracy: 0.5839 - loss: 1.5997 - val_accuracy: 0.8336 - val_loss: 0.6335
Epoch 2/5
74/74 ----- 0s 2s/step - accuracy: 0.8491 - loss: 0.4957
Epoch 2: val_accuracy improved from 0.83362 to 0.85229, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross

Epoch 2: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_2/best_mod
74/74 ----- 190s 2s/step - accuracy: 0.8497 - loss: 0.4829 - val_accuracy: 0.8523 - val_loss: 0.4612
Epoch 3/5
74/74 ----- 0s 2s/step - accuracy: 0.9193 - loss: 0.2694
Epoch 3: val_accuracy improved from 0.85229 to 0.88285, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross

Epoch 3: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_2/best_mod
```

```
# =====
# CELL R27 - 5-FOLD CV SUMMARY
# =====

# Reload from Drive to make sure we use everything R26 saved.
CV_RESULTS_PATH = os.path.join(
    RESULTS_DIR,
    "cv_results.csv"
)

cv_results_df = pd.read_csv(
    CV_RESULTS_PATH
)

print("=" * 70)
print("CROSS-VALIDATION SUMMARY")
print("=" * 70)

cv_summary = (
    cv_results_df
    .groupby("Configuration")
    .agg(
        Mean=(
            "Best Validation Accuracy (%)",
            "mean"
        ),
        SD=(
            "Best Validation Accuracy (%)",
            "std"
        )
    )
)
```

```

    ),
    **{
        "Mean Training Time (s)": (
            "Training Time (s)",
            "mean"
        )
    }
)
.reset_index()
)

cv_summary.to_csv(
    os.path.join(
        RESULTS_DIR,
        "cv_summary.csv"
    ),
    index=False
)

display(cv_summary)

print()
print(
    "Saved:",
    os.path.join(
        RESULTS_DIR,
        "cv_summary.csv"
    )
)

```

```

Epoch 4/5
74/74 ===== 0s 2s/step - accuracy: 0.9450 - loss: 0.1699
Epoch 4: val accuracy improved from 0.88115 to 0.88964, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation
=====
Epoch 4 finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_4/best_model
74/74 Configuration 164s 2s/step - accuracy: 0.9435 - loss: 0.1714 - val_accuracy: 0.8896 - val_loss: 0.3465
0 Baseline 91.881607 1.449007 722.465817
74/74 ===== 0s 2s/step - accuracy: 0.9450 - loss: 0.1699
Epoch 5: val accuracy improved from 0.88964 to 0.89989, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation
=====
Epoch 5; finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/fine-tuned/fold_4/best_model
74/74 Configuration 206s 2s/step - accuracy: 0.9707 - loss: 0.1013 - val_accuracy: 0.8998 - val_loss: 0.3215

```

```
# =====
# CELL R28 - 5-FOLD CV RESULTS PLOT
# =====

import matplotlib.pyplot as plt
import numpy as np

plt.figure(figsize=(8, 5))

for configuration in cv_results_df["Configuration"].unique():


    subset = cv_results_df[
        cv_results_df["Configuration"] == configuration
    ]

    plt.plot(
        subset["Fold"],
        subset["Best Validation Accuracy (%)"],
        marker="o",
        linewidth=1.8,
        label=configuration
    )

plt.xlabel("Fold")
plt.ylabel("Best Validation Accuracy (%)")
plt.title("5-Fold Cross-Validation Results")
plt.xticks([1, 2, 3, 4, 5])
plt.legend()
plt.grid(alpha=0.2)

plt.savefig(
    os.path.join(
        OUTPUT_DIR,
        "cv_results.png"
    ),
    dpi=600,
    bbox_inches="tight"
)

plt.show()
plt.close()
```



```
# =====
# CELL R29 - BEST CV CONFIGURATION
# =====

best_cv_row = cv_summary.loc[
    cv_summary["Mean"].idxmax()
]

best_cv_configuration = best_cv_row["Configuration"]

print("=" * 70)
print("BEST CROSS-VALIDATION CONFIGURATION")
print("=" * 70)

print(
```

```

    "Configuration:",
    best_cv_configuration
)

print(
    f"Mean validation accuracy: "
    f"{best_cv_row['Mean']:.2f}%"
)

print(
    f"Standard deviation: "
    f"{best_cv_row['SD']:.2f}%"
)

print(
    f"Mean training time: "
    f"{best_cv_row['Mean Training Time (s)']:.2f} s"
)

```

```

Epoch 4/5
=====
74774s 134s 2s/step - accuracy: 0.9683 - loss: 0.1257
Epoch 4: val accuracy did not improve from 0.913379, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation_configuration
=====
Epoch 4: saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/baseline/fold_4/best_model
74774s 134s 2s/step - accuracy: 0.9652 - loss: 0.1342 - val_accuracy: 0.9338 - val_loss: 0.2368
Epoch 4: Standard deviation: 1.45%
74774s 134s 2s/step - accuracy: 0.9781 - loss: 0.0963
Epoch 4: training time: 722.47 0s 1s/step - accuracy: 0.9781 - loss: 0.0963
Epoch 5: val accuracy did not improve from 0.93379

```

```

# =====
# CELL R30 - FINAL MODEL CONFIGURATION
# =====

```

```

print("=" * 70)
print("FINAL MODEL CONFIGURATION")
print("=" * 70)

```

```

if best_cv_configuration == "Fine-Tuned":

```

```

    FINAL_LEARNING_RATE = 0.0001
    FINAL_DROPOUT = float(best_dropout)
    FINAL_FINE_TUNE = True

```

```

else:

```

```

    FINAL_LEARNING_RATE = float(best_lr)
    FINAL_DROPOUT = 0.0
    FINAL_FINE_TUNE = False

```

```

FINAL_BATCH_SIZE = int(float(best_batch_size))

```

```

print("Configuration:", best_cv_configuration)
print("Learning rate:", FINAL_LEARNING_RATE)
print("Batch size:", FINAL_BATCH_SIZE)
print("Dropout:", FINAL_DROPOUT)
print("Fine-tuning:", FINAL_FINE_TUNE)

```

```

Epoch 5: val accuracy did not improve from 0.91322
=====
74774s 134s 2s/step - accuracy: 0.9784 - loss: 0.0896 - val_accuracy: 0.9048 - val_loss: 0.3084
=====
Epoch 5: saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation/baseline
Best validation accuracy: 91.33%
Batch size: 32
Dropout: 0.0
=====
Epoch 5: training time: 722.47 0s 1s/step - accuracy: 0.9781 - loss: 0.0963
Epoch 5: val accuracy did not improve from 0.93379
=====

```

```

# =====
# CELL R30A - RESTORE PRE-R31 CHECKPOINT
# =====

```

```

import os
import json
import pandas as pd
import tensorflow as tf

```

```

# -----
# Paths
# -----

```

```

LAB5_DRIVE_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

```

```

CHECKPOINT_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "checkpoints"
)

```

```
)

RESULTS_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "results"
)

PRE_R31_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "pre_r31_checkpoint"
)

print("=" * 70)
print("RESTORING PRE-R31 CHECKPOINT")
print("=" * 70)

print("Lab 5 root:")
print(LAB5_DRIVE_ROOT)

print("\nCheckpoint:")
print(PRE_R31_DIR)

# -----
# Verify checkpoint exists
# -----

if not os.path.exists(PRE_R31_DIR):
    raise FileNotFoundError(
        f"Pre-R31 checkpoint not found:\n{PRE_R31_DIR}"
    )

# -----
# Load checkpoint manifest
# -----

manifest_path = os.path.join(
    PRE_R31_DIR,
    "checkpoint_manifest.json"
)

if os.path.exists(manifest_path):

    with open(manifest_path, "r") as f:
        pre_r31_manifest = json.load(f)

    print("\nCheckpoint manifest loaded.")

    print(
        "Checkpoint stage:",
        pre_r31_manifest.get("stage")
    )

    print(
        "R31 already started:",
        pre_r31_manifest.get("r31_started")
    )

else:

    print(
        "\nWarning: checkpoint_manifest.json not found."
    )

    pre_r31_manifest = {}

# -----
# Restore result files from checkpoint
# -----

backup_results_dir = os.path.join(
    PRE_R31_DIR,
    "results"
)

if os.path.exists(backup_results_dir):

    os.makedirs(
        RESULTS_DIR,
        exist_ok=True
```

```

    )

    restored = []

    for filename in os.listdir(backup_results_dir):

        source = os.path.join(
            backup_results_dir,
            filename
        )

        destination = os.path.join(
            RESULTS_DIR,
            filename
        )

        if os.path.isfile(source):

            # Do not overwrite newer files unnecessarily.
            # Restore only if the current file is missing.
            if not os.path.exists(destination):

                import shutil

                shutil.copy2(
                    source,
                    destination
                )

                restored.append(filename)

    print(
        f"\nRestored {len(restored)} missing result file(s)."
    )

    if restored:
        for filename in restored:
            print(" ✓", filename)

else:

    print(
        "\nNo backed-up results directory found."
    )

# -----
# Reload important result DataFrames
# -----

def load_result_csv(filename):

    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):

        return pd.read_csv(path)

    return None

initialization_df = load_result_csv(
    "initialization_results.csv"
)

regularization_df = load_result_csv(
    "regularization_results.csv"
)

batch_normalization_df = load_result_csv(
    "batch_normalization_results.csv"
)

optimizer_df = load_result_csv(
    "optimizer_results.csv"
)

hyperparameter_df = load_result_csv(
    "hyperparameter_results.csv"
)

```

```

)

transfer_df = load_result_csv(
    "transfer_learning_results.csv"
)

fine_tune_df = load_result_csv(
    "correct_fine_tuning_results.csv"
)

cv_results_df = load_result_csv(
    "cv_results.csv"
)

cv_summary = load_result_csv(
    "cv_summary.csv"
)

final_df = load_result_csv(
    "final_model_results.csv"
)

# -----
# Display restored state
# -----

print("\n" + "=" * 70)
print("RESTORED EXPERIMENT STATE")
print("=" * 70)

for label, df in [

    ("Initialization", initialization_df),
    ("Regularization", regularization_df),
    ("Batch Normalization", batch_normalization_df),
    ("Optimizer", optimizer_df),
    ("Hyperparameters", hyperparameter_df),
    ("Transfer Learning", transfer_df),
    ("Fine-Tuning", fine_tune_df),
    ("Cross-Validation", cv_results_df),
    ("CV Summary", cv_summary),
    ("Final Model", final_df)

]:

    if df is None:

        print(
            f" {label}: not found"
        )

    else:

        print(
            f" {label}: {len(df)} row(s)"
        )

# -----
# Restore CV checkpoint path
# -----

CV_DIR = os.path.join(
    CHECKPOINT_DIR,
    "cross_validation"
)

if os.path.exists(CV_DIR):

    print(
        "\n✓ Cross-validation checkpoints found."
    )

else:

    print(
        "\n△ Cross-validation checkpoint directory not found."
    )

# -----

```

```

# Restore fine-tuning checkpoint path
# -----

FINE_TUNE_DIR = os.path.join(
    CHECKPOINT_DIR,
    "fine_tuning_final"
)

if os.path.exists(FINE_TUNE_DIR):

    print(
        "✓ Fine-tuning checkpoint found."
    )

else:

    print(
        "⚠ Fine-tuning checkpoint directory not found."
    )

# -----
# GPU verification
# -----

gpus = tf.config.list_physical_devices(
    "GPU"
)

print("\nGPU:", gpus)

if not gpus:

    raise RuntimeError(
        "No GPU detected. Do not start R31."
    )

# -----
# Final confirmation
# -----

print("\n" + "=" * 70)
print("PRE-R31 CHECKPOINT RESTORED")
print("=" * 70)

print("✓ Results loaded")
print("✓ CV checkpoint location verified")
print("✓ Fine-tuning checkpoint location verified")
print("✓ GPU detected")
print("\nReady for R31 – Final Model Evaluation.")

```

```

=====
RESTORING PRE-R31 CHECKPOINT
=====
Lab 5 root:
/content/drive/MyDrive/Lab5/Lab5

Checkpoint:
/content/drive/MyDrive/Lab5/Lab5/pre_R31_checkpoint

Checkpoint manifest loaded.
Checkpoint stage: BEFORE_R31
R31 already started: False

Restored 0 missing result file(s).

=====
RESTORED EXPERIMENT STATE
=====
Initialization: 4 row(s)
Regularization: 4 row(s)
Batch Normalization: 1 row(s)
Optimizer: 4 row(s)
Hyperparameters: 5 row(s)
Transfer Learning: 2 row(s)
Fine-Tuning: 1 row(s)
Cross-Validation: 10 row(s)
CV Summary: 2 row(s)
Final Model: not found

✓ Cross-validation checkpoints found.
✓ Fine-tuning checkpoint found.

GPU: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]

```

```
=====
PRE-R31 CHECKPOINT RESTORED
=====
✓ Results loaded
✓ CV checkpoint location verified
✓ Fine-tuning checkpoint location verified
✓ GPU detected

Ready for R31 – Final Model Evaluation.
```

```
# =====
# CELL R30B – RESTORE EXPERIMENT CONSTANTS
# =====

# Same seed used by the original notebook
SEED = 42

# Dataset image size
IMG_SIZE = (224, 224)

# Default batch size
BATCH_SIZE = 32

print("=" * 70)
print("EXPERIMENT CONSTANTS RESTORED")
print("=" * 70)

print("SEED:", SEED)
print("IMG_SIZE:", IMG_SIZE)
print("BATCH_SIZE:", BATCH_SIZE)
```

```
=====
EXPERIMENT CONSTANTS RESTORED
=====
SEED: 42
IMG_SIZE: (224, 224)
BATCH_SIZE: 32
```

```
# =====
# CELL R30C – RESTORE DATASET VARIABLES FOR R31
# =====

import os
import numpy as np
import tensorflow as tf

from tensorflow.keras.applications.mobilenet_v2 import preprocess_input

# -----
# DATASET PATHS
# -----

DATA_ROOT = os.path.join(
    LAB5_DRIVE_ROOT,
    "dataset"
)

IMAGES_DIR = os.path.join(
    DATA_ROOT,
    "images"
)

ANNOTATIONS_DIR = os.path.join(
    DATA_ROOT,
    "annotations"
)

LIST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "list.txt"
)

TRAINVAL_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "trainval.txt"
)

TEST_FILE = os.path.join(
    ANNOTATIONS_DIR,
    "test.txt"
)
```

```

# -----
# VERIFY DATASET
# -----

print("=" * 70)
print("RESTORING DATASET")
print("=" * 70)

print("Dataset:", DATA_ROOT)

if not os.path.exists(IMAGES_DIR):
    raise FileNotFoundError(
        f"Images directory not found:\n{IMAGES_DIR}"
    )

if not os.path.exists(ANNOTATIONS_DIR):
    raise FileNotFoundError(
        f"Annotations directory not found:\n{ANNOTATIONS_DIR}"
    )

# -----
# READ SPLIT FILES
# -----

def read_split_file(filename):

    names = []

    with open(filename, "r") as f:

        for line in f:

            line = line.strip()

            if (
                not line
                or line.startswith("#")
            ):
                continue

            names.append(
                line.split()[0]
            )

    return set(names)

trainval_names = read_split_file(
    TRAINVAL_FILE
)

test_names = read_split_file(
    TEST_FILE
)

# -----
# READ ALL IMAGES
# -----

samples = []

with open(LIST_FILE, "r") as f:

    for line in f:

        line = line.strip()

        if (
            not line
            or line.startswith("#")
        ):
            continue

        image_name = line.split()[0]

        image_path = os.path.join(
            IMAGES_DIR,
            image_name + ".jpg"

```

```

    )

    if os.path.exists(image_path):

        samples.append(
            image_path
        )

print("Images found:", len(samples))

# -----
# CLASS NAMES
# -----

class_names = sorted(
    set(
        os.path.basename(path)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
        for path in samples
    )
)

class_to_index = {
    name: i
    for i, name
    in enumerate(class_names)
}

print("Classes:", len(class_names))

# -----
# CREATE PATH/LABEL ARRAYS
# -----

image_paths = []
labels = []

for path in samples:

    breed = (
        os.path.basename(path)
        .replace(".jpg", "")
        .rsplit("_", 1)[0]
    )

    if breed in class_to_index:

        image_paths.append(path)

        labels.append(
            class_to_index[breed]
        )

image_paths = np.array(
    image_paths
)

labels = np.array(
    labels,
    dtype=np.int32
)

# -----
# RECREATE TRAIN/TEST SPLIT
# -----

trainval_paths = []
trainval_labels = []

test_paths = []
test_labels = []

for path, label in zip(
    image_paths,
    labels

```

```
    ):

        name = (
            os.path.basename(path)
            .replace(".jpg", "")
        )

        if name in trainval_names:

            trainval_paths.append(path)
            trainval_labels.append(label)

        elif name in test_names:

            test_paths.append(path)
            test_labels.append(label)

trainval_paths = np.array(
    trainval_paths
)

trainval_labels = np.array(
    trainval_labels,
    dtype=np.int32
)

test_paths = np.array(
    test_paths
)

test_labels = np.array(
    test_labels,
    dtype=np.int32
)

# -----
# RECREATE SAME 80/20 TRAIN/VALIDATION SPLIT
# -----

rng = np.random.default_rng(
    SEED
)

indices = np.arange(
    len(trainval_paths)
)

rng.shuffle(
    indices
)

val_size = int(
    0.20 * len(indices)
)

val_indices = indices[
    :val_size
]

train_indices = indices[
    val_size:
]

train_paths = trainval_paths[
    train_indices
]

train_labels = trainval_labels[
    train_indices
]

val_paths = trainval_paths[
    val_indices
]

val_labels = trainval_labels[
    val_indices
]
```

```

# -----
# IMAGE LOADER
# -----

def load_image(path, label):

    image = tf.io.read_file(
        path
    )

    image = tf.image.decode_jpeg(
        image,
        channels=3
    )

    image = tf.image.resize(
        image,
        IMG_SIZE
    )

    image = tf.cast(
        image,
        tf.float32
    )

    image = preprocess_input(
        image
    )

    return image, label


# -----
# DATASET CREATOR
# -----

def make_dataset(
    paths,
    labels,
    batch_size=32,
    shuffle=False
):

    ds = tf.data.Dataset.from_tensor_slices(
        (
            paths,
            labels
        )
    )

    if shuffle:

        ds = ds.shuffle(
            len(paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        int(batch_size)
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds


# -----
# CREATE DATASETS
# -----

train_ds = make_dataset(
    train_paths,
    train_labels,

```

```

        BATCH_SIZE,
        shuffle=True
    )

    val_ds = make_dataset(
        val_paths,
        val_labels,
        BATCH_SIZE,
        shuffle=False
    )

    test_ds = make_dataset(
        test_paths,
        test_labels,
        BATCH_SIZE,
        shuffle=False
    )

    # -----
    # VERIFY
    # -----

    print("\n" + "=" * 70)
    print("DATASET RESTORED")
    print("=" * 70)

    print("Classes:      ", len(class_names))
    print("Train:         ", len(train_paths))
    print("Validation:     ", len(val_paths))
    print("Test:           ", len(test_paths))

    print("\nRequired variables now exist:")

    for variable in [
        "train_paths",
        "train_labels",
        "val_paths",
        "val_labels",
        "test_paths",
        "test_labels"
    ]:
        print(
            "    ✓",
            variable
        )

    print("\n✓ Ready for R31.")

```

```

=====
RESTORING DATASET
=====
Dataset: /content/drive/MyDrive/Lab5/Lab5/dataset
Images found: 7349
Classes: 37

=====
DATASET RESTORED
=====
Classes:      37
Train:        2944
Validation:    736
Test:         3669

Required variables now exist:
✓ train_paths
✓ train_labels
✓ val_paths
✓ val_labels
✓ test_paths
✓ test_labels

✓ Ready for R31.

```

```

# =====
# CELL R30F – RESTORE BEST HYPERPARAMETERS
# =====

print("=" * 70)
print("RESTORING BEST HYPERPARAMETERS")
print("=" * 70)

```

```

# -----
# Verify hyperparameter results
# -----

if hyperparameter_df is None:
    raise RuntimeError(
        "hyperparameter_results.csv was not found."
    )

# -----
# Restore best Learning Rate
# -----

lr_rows = hyperparameter_df[
    hyperparameter_df["Hyperparameter"]
    .astype(str)
    .str.strip()
    .str.lower()
    == "learning rate"
].dropna(
    subset=["Best Validation Accuracy (%)"]
).sort_values(
    "Best Validation Accuracy (%)",
    ascending=False
)

if lr_rows.empty:
    raise RuntimeError(
        "No valid Learning Rate results found."
    )

best_lr = float(
    lr_rows.iloc[0]["Value"]
)

# -----
# Restore best Dropout
# -----

dropout_rows = hyperparameter_df[
    hyperparameter_df["Hyperparameter"]
    .astype(str)
    .str.strip()
    .str.lower()
    == "dropout"
].dropna(
    subset=["Best Validation Accuracy (%)"]
).sort_values(
    "Best Validation Accuracy (%)",
    ascending=False
)

if dropout_rows.empty:
    raise RuntimeError(
        "No valid Dropout results found."
    )

best_dropout = float(
    dropout_rows.iloc[0]["Value"]
)

# -----
# Batch size
#
# No Batch Size experiment exists in the saved results.
# Therefore use the standard batch size used by the
# reconstructed Lab 5 pipeline.
# -----

best_batch_size = 32

# -----
# Display selected configuration
# -----

print("\nSELECTED FINAL MODEL HYPERPARAMETERS")
print("=" * 50)

```

```

print(
    "Learning rate:",
    best_lr
)

print(
    "Batch size:",
    best_batch_size,
    "(no saved batch-size experiment)"
)

print(
    "Dropout:",
    best_dropout
)

# -----
# Show actual saved Learning Rate results
# -----

print("\nSaved Learning Rate results:")
display(lr_rows)

# -----
# Show actual saved Dropout results
# -----

print("\nSaved Dropout results:")
display(dropout_rows)

print("\n" + "=" * 70)
print("✓ HYPERPARAMETERS RESTORED")
print("✓ No unsupported Batch Size result was assumed")
print("✓ R31 is ready")
print("=" * 70)

```

```

=====
RESTORING BEST HYPERPARAMETERS
=====

```

```

SELECTED FINAL MODEL HYPERPARAMETERS
=====
Learning rate: 0.001
Batch size: 32 (no saved batch-size experiment)
Dropout: 0.25

```

Saved Learning Rate results:

	Hyperparameter	Value	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Learning Rate	0.0010	96.467394	NaN	96.467394	NaN	0.117140	NaN
1	Learning Rate	0.0001	78.532606	NaN	78.532606	NaN	1.218464	NaN

Saved Dropout results:

	Hyperparameter	Value	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
3	Dropout	0.25	91.168481	95.923913	90.081519	0.130857	0.300911	92.754024
2	Dropout	0.00	90.896738	97.894019	90.896738	0.082908	0.276094	105.535548
4	Dropout	0.50	90.217394	92.832881	90.081519	0.216385	0.297352	84.395862

```

=====
✓ HYPERPARAMETERS RESTORED
✓ No unsupported Batch Size result was assumed
✓ R31 is ready

```

```

# =====
# CELL R30E - RESTORE NUM_CLASSES
# =====

NUM_CLASSES = len(class_names)

print("=" * 70)
print("MODEL CONSTANT RESTORED")
print("=" * 70)

print("Number of classes:", NUM_CLASSES)

```

```

if NUM_CLASSES != 37:
    raise RuntimeError(
        f"Expected 37 classes, but found {NUM_CLASSES}."
    )

print("✓ NUM_CLASSES =", NUM_CLASSES)
print("✓ Ready to rerun R31.")

```

```

=====
MODEL CONSTANT RESTORED
=====

```

```

Number of classes: 37
✓ NUM_CLASSES = 37
✓ Ready to rerun R31.

```

```

# =====
# CELL R30F – RESTORE ALL FINAL-MODEL RUNTIME VARIABLES
# =====

import os
import gc
import json
import time
import numpy as np
import pandas as pd
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.mobilenet_v2 import preprocess_input

# -----
# Experiment constants
# -----

SEED = 42
IMG_SIZE = (224, 224)
BATCH_SIZE = 32
EPOCHS = 10

# -----
# Model constants
# -----

NUM_CLASSES = len(class_names)

# -----
# Verify dataset variables
# -----

required_variables = [
    "train_paths",
    "train_labels",
    "val_paths",
    "val_labels",
    "test_paths",
    "test_labels",
    "class_names"
]

missing = [
    variable
    for variable in required_variables
    if variable not in globals()
]

if missing:
    raise RuntimeError(
        "Missing required variables: "
        + str(missing)
    )

# -----
# Verify selected hyperparameters
# -----

```

```

if "best_lr" not in globals():
    best_lr = 0.001

if "best_batch_size" not in globals():
    best_batch_size = 32

if "best_dropout" not in globals():
    best_dropout = 0.25

best_lr = float(best_lr)
best_batch_size = int(best_batch_size)
best_dropout = float(best_dropout)

# -----
# Verify GPU
# -----

gpus = tf.config.list_physical_devices("GPU")

if not gpus:
    raise RuntimeError(
        "No GPU detected. Stop before running final training."
    )

# -----
# Verify final directories
# -----

FINAL_DIR = os.path.join(
    CHECKPOINT_DIR,
    "final_model"
)

os.makedirs(
    FINAL_DIR,
    exist_ok=True
)

FINAL_RESULTS_PATH = os.path.join(
    RESULTS_DIR,
    "final_model_results.csv"
)

# -----
# Final verification
# -----

print("=" * 70)
print("FINAL MODEL RUNTIME VERIFIED")
print("=" * 70)

print("TensorFlow:", tf.__version__)
print("GPU:", gpus)

print("\nConstants:")
print("  SEED:", SEED)
print("  IMG_SIZE:", IMG_SIZE)
print("  BATCH_SIZE:", BATCH_SIZE)
print("  EPOCHS:", EPOCHS)
print("  NUM_CLASSES:", NUM_CLASSES)

print("\nSelected hyperparameters:")
print("  Learning rate:", best_lr)
print("  Batch size:", best_batch_size)
print("  Dropout:", best_dropout)

print("\nDataset:")
print("  Train:", len(train_paths))
print("  Validation:", len(val_paths))
print("  Test:", len(test_paths))

print("\n✓ All R31 runtime dependencies are present.")
print("✓ GPU is available.")
print("✓ Ready for final model training.")

```

```

=====
FINAL MODEL RUNTIME VERIFIED
=====
TensorFlow: 2.20.0

```

GPU: [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]

Constants:
 SEED: 42
 IMG_SIZE: (224, 224)
 BATCH_SIZE: 32
 EPOCHS: 10
 NUM_CLASSES: 37

Selected hyperparameters:
 Learning rate: 0.001
 Batch size: 32
 Dropout: 0.25

Dataset:
 Train: 2944
 Validation: 736
 Test: 3669

- ✓ All R31 runtime dependencies are present.
- ✓ GPU is available.
- ✓ Ready for final model training.

```
# =====
# CELL R31 - CORRECTED FINAL MODEL
# =====

import os
import gc
import json
import time
import numpy as np
import pandas as pd
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2

# =====
# 1. CONFIGURATION
# =====

FINAL_LR = 0.001
FINAL_BATCH_SIZE = 32
FINAL_DROPOUT = 0.25
FINAL_EPOCHS = 10
FINAL_NUM_CLASSES = len(class_names)

FINAL_DIR = os.path.join(
    CHECKPOINT_DIR,
    "final_model_corrected"
)

os.makedirs(
    FINAL_DIR,
    exist_ok=True
)

FINAL_RESULTS_PATH = os.path.join(
    RESULTS_DIR,
    "final_model_results.csv"
)

FINAL_HISTORY_PATH = os.path.join(
    FINAL_DIR,
    "history.json"
)

FINAL_MODEL_PATH = os.path.join(
    FINAL_DIR,
    "final_model.keras"
)

FINAL_BEST_MODEL_PATH = os.path.join(
    FINAL_DIR,
    "best_model.keras"
)

# =====
# 2. VERIFY EVERYTHING BEFORE USING GPU
# =====
```

```

print("=" * 70)
print("CORRECTED FINAL MODEL")
print("=" * 70)

required_variables = [
    "train_paths",
    "train_labels",
    "val_paths",
    "val_labels",
    "test_paths",
    "test_labels",
    "class_names"
]

missing = [
    variable
    for variable in required_variables
    if variable not in globals()
]

if missing:
    raise RuntimeError(
        "Required variables are missing: "
        + str(missing)
    )

gpus = tf.config.list_physical_devices("GPU")

if not gpus:
    raise RuntimeError(
        "No GPU detected. Do not start final training."
    )

print("\nGPU detected:")
for gpu in gpus:
    print(" ", gpu)

print("\nSelected configuration:")
print("  Learning rate :", FINAL_LR)
print("  Batch size    :", FINAL_BATCH_SIZE)
print("  Dropout       :", FINAL_DROPOUT)
print("  Epochs        :", FINAL_EPOCHS)
print("  Classes       :", FINAL_NUM_CLASSES)

print("\nDataset:")
print("  Training samples :", len(train_paths))
print("  Validation samples :", len(val_paths))
print("  Test samples      :", len(test_paths))

# =====
# 3. CREATE DATASETS
# =====

def make_final_dataset(
    paths,
    labels,
    batch_size,
    shuffle=False
):

    ds = tf.data.Dataset.from_tensor_slices(
        (
            paths,
            labels
        )
    )

    if shuffle:
        ds = ds.shuffle(
            buffer_size=len(paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE

```

```

    )

    ds = ds.batch(
        batch_size
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

final_train_ds = make_final_dataset(
    train_paths,
    train_labels,
    FINAL_BATCH_SIZE,
    shuffle=True
)

final_val_ds = make_final_dataset(
    val_paths,
    val_labels,
    FINAL_BATCH_SIZE,
    shuffle=False
)

final_test_ds = make_final_dataset(
    test_paths,
    test_labels,
    FINAL_BATCH_SIZE,
    shuffle=False
)

# =====
# 4. BUILD FINAL MODEL
#
# IMPORTANT:
# MobileNetV2 base is FROZEN.
# This matches the better-performing CV baseline.
# =====

base_model = MobileNetV2(
    input_shape=(224, 224, 3),
    include_top=False,
    weights="imagenet"
)

base_model.trainable = False

inputs = keras.Input(
    shape=(224, 224, 3)
)

x = base_model(
    inputs,
    training=False
)

x = layers.GlobalAveragePooling2D()(x)

x = layers.Dropout(
    FINAL_DROPOUT
)(x)

outputs = layers.Dense(
    FINAL_NUM_CLASSES,
    activation="softmax"
)(x)

final_model = keras.Model(
    inputs,
    outputs
)

# =====
# 5. COMPILE
# =====

```

```

final_model.compile(
    optimizer=keras.optimizers.Adam(
        learning_rate=FINAL_LR
    ),
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)

trainable_params = int(
    np.sum([
        np.prod(variable.shape)
        for variable in final_model.trainable_weights
    ])
)

total_params = int(
    final_model.count_params()
)

print("\nModel parameters:")
print("  Trainable parameters:", trainable_params)
print("  Total parameters      :", total_params)

# =====
# 6. CALLBACKS
# =====

checkpoint_callback = keras.callbacks.ModelCheckpoint(
    FINAL_BEST_MODEL_PATH,
    monitor="val_accuracy",
    mode="max",
    save_best_only=True,
    save_weights_only=False,
    verbose=1
)

backup_callback = keras.callbacks.BackupAndRestore(
    backup_dir=os.path.join(
        FINAL_DIR,
        "training_backup"
    ),
    delete_checkpoint=False
)

# =====
# 7. TRAIN
# =====

print("\n" + "=" * 70)
print("FINAL MODEL TRAINING")
print("=" * 70)

start_time = time.time()

final_history = final_model.fit(
    final_train_ds,
    validation_data=final_val_ds,
    epochs=FINAL_EPOCHS,
    verbose=1,
    callbacks=[
        checkpoint_callback,
        backup_callback
    ]
)

training_time = time.time() - start_time

# =====
# 8. RESTORE BEST MODEL
# =====

if os.path.exists(FINAL_BEST_MODEL_PATH):
    print("\nRestoring best validation model...")

    final_model = keras.models.load_model(
        FINAL_BEST_MODEL_PATH

```

```

    )

# =====
# 9. TEST EVALUATION
# =====

print("\n" + "=" * 70)
print("FINAL TEST EVALUATION")
print("=" * 70)

test_loss, test_accuracy = final_model.evaluate(
    final_test_ds,
    verbose=1
)

test_accuracy_percent = (
    float(test_accuracy) * 100
)

test_loss = float(test_loss)

print("\nFinal Test Loss:",
      test_loss)

print(
    "Final Test Accuracy:",
    f"{test_accuracy_percent:.2f}%"
)

# =====
# 10. HISTORY VALIDATION
# =====

history_dict = final_history.history

required_metrics = [
    "loss",
    "accuracy",
    "val_loss",
    "val_accuracy"
]

missing_metrics = [
    metric
    for metric in required_metrics
    if metric not in history_dict
]

if missing_metrics:
    raise RuntimeError(
        "Missing training metrics: "
        + str(missing_metrics)
    )

best_val_accuracy = (
    max(history_dict["val_accuracy"]) * 100
)

final_train_accuracy = (
    history_dict["accuracy"][-1] * 100
)

final_val_accuracy = (
    history_dict["val_accuracy"][-1] * 100
)

final_train_loss = (
    history_dict["loss"][-1]
)

final_val_loss = (
    history_dict["val_loss"][-1]
)

# =====
# 11. SAVE HISTORY

```

```

# =====

with open(
    FINAL_HISTORY_PATH,
    "w"
) as f:

    json.dump(
        {
            key: [
                float(value)
                for value in values
            ]
            for key, values
            in history_dict.items()
        },
        f,
        indent=2
    )

# =====
# 12. SAVE COMPLETE MODEL
# =====

final_model.save(
    FINAL_MODEL_PATH
)

# =====
# 13. SAVE RESULTS
# =====

final_results = pd.DataFrame([
    {
        "Model":
            "Final Selected Model",

        "Learning Rate":
            FINAL_LR,

        "Batch Size":
            FINAL_BATCH_SIZE,

        "Dropout":
            FINAL_DROPOUT,

        "Trainable Parameters":
            trainable_params,

        "Total Parameters":
            total_params,

        "Best Validation Accuracy (%)":
            best_val_accuracy,

        "Final Training Accuracy (%)":
            final_train_accuracy,

        "Final Validation Accuracy (%)":
            final_val_accuracy,

        "Final Training Loss":
            final_train_loss,

        "Final Validation Loss":
            final_val_loss,

        "Test Loss":
            test_loss,

        "Test Accuracy (%)":
            test_accuracy_percent,

        "Training Time (s)":
            training_time
    }
])

final_results.to_csv(

```

```
        FINAL_RESULTS_PATH,
        index=False
    )

# =====
# 14. DISPLAY FINAL RESULT
# =====

print("\n" + "=" * 70)
print("FINAL MODEL COMPLETE")
print("=" * 70)

display(
    final_results
)

print("\nSaved:")
print("  Best model :", FINAL_BEST_MODEL_PATH)
print("  Final model:", FINAL_MODEL_PATH)
print("  History    :", FINAL_HISTORY_PATH)
print("  Results    :", FINAL_RESULTS_PATH)

# =====
# 15. CLEANUP
# =====

del final_train_ds
del final_val_ds
del final_test_ds
del base_model

tf.keras.backend.clear_session()
gc.collect()

print("\n✓ Corrected final model saved successfully.")
```

```

=====
CORRECTED FINAL MODEL
=====

GPU detected:
PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')

Selected configuration:
  Learning rate : 0.001
  Batch size    : 32
  Dropout       : 0.25
  Epochs        : 10
  Classes       : 37

Dataset:
  Training samples : 2944
  Validation samples : 736
  Test samples     : 3669

Model parameters:
  Trainable parameters: 47397
  Total parameters    : 2305381

=====
FINAL MODEL TRAINING
=====

Epoch 1/10
92/92 ----- 0s 123ms/step - accuracy: 0.3965 - loss: 2.3807
Epoch 1: val_accuracy improved from None to 0.87364, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final_mo

Epoch 1: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final_model_corrected/best_model.keras
92/92 ----- 35s 259ms/step - accuracy: 0.6277 - loss: 1.4623 - val_accuracy: 0.8736 - val_loss: 0.5149
Epoch 2/10
92/92 ----- 0s 116ms/step - accuracy: 0.8951 - loss: 0.4316
Epoch 2: val_accuracy improved from 0.87364 to 0.88451, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final

Epoch 2: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final_model_corrected/best_model.keras
92/92 ----- 17s 189ms/step - accuracy: 0.9018 - loss: 0.3960 - val_accuracy: 0.8845 - val_loss: 0.3622
Epoch 3/10
92/92 ----- 0s 125ms/step - accuracy: 0.9354 - loss: 0.2741
Epoch 3: val_accuracy improved from 0.88451 to 0.89810, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final

Epoch 3: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final_model_corrected/best_model.keras
92/92 ----- 15s 165ms/step - accuracy: 0.9368 - loss: 0.2587 - val_accuracy: 0.8981 - val_loss: 0.3221
Epoch 4/10
92/92 ----- 0s 130ms/step - accuracy: 0.9584 - loss: 0.1859
Epoch 4: val_accuracy did not improve from 0.89810
92/92 ----- 15s 166ms/step - accuracy: 0.9552 - loss: 0.1883 - val_accuracy: 0.8940 - val_loss: 0.3089
Epoch 5/10
92/92 ----- 0s 125ms/step - accuracy: 0.9718 - loss: 0.1432
Epoch 5: val_accuracy improved from 0.89810 to 0.90761, saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final

Epoch 5: finished saving model to /content/drive/MyDrive/Lab5/Lab5/checkpoints/final_model_corrected/best_model.keras
92/92 ----- 15s 168ms/step - accuracy: 0.9694 - loss: 0.1414 - val_accuracy: 0.9076 - val_loss: 0.2797
Epoch 6/10
92/92 ----- 0s 117ms/step - accuracy: 0.9788 - loss: 0.1203
Epoch 6: val_accuracy did not improve from 0.90761

# =====
# CELL R32 - FINAL MODEL CHECKPOINT & SUMMARY
# =====

import os
import json
import pandas as pd

FINAL_DIR = os.path.join(
    CHECKPOINT_DIR,
    "final_model_corrected"
)

FINAL_RESULTS_PATH = os.path.join(
    RESULTS_DIR,
    "final_model_results.csv"
)

print("=" * 70)
print("LAB 5 - FINAL MODEL CHECKPOINT")
print("=" * 70)

# -----
# Verify final files
# -----

required_files = [
    "best_model.keras",

```

```

    "final_model.keras",
    "history.json"
]

print("\nFinal model files:")

for filename in required_files:

    path = os.path.join(
        FINAL_DIR,
        filename
    )

    if os.path.exists(path):
        size_mb = os.path.getsize(path) / (1024 ** 2)

        print(
            f"    ✓ {filename} "
            f"({size_mb:.2f} MB)"
        )

    else:

        print(
            f"    X {filename} NOT FOUND"
        )

# -----
# Load final results
# -----

if not os.path.exists(FINAL_RESULTS_PATH):

    raise FileNotFoundError(
        "Final model results were not found:\n"
        + FINAL_RESULTS_PATH
    )

final_model_results = pd.read_csv(
    FINAL_RESULTS_PATH
)

# -----
# Display result
# -----

print("\n" + "=" * 70)
print("FINAL MODEL RESULT")
print("=" * 70)

display(
    final_model_results
)

# -----
# Extract key values
# -----

result = final_model_results.iloc[0]

print("\nKey results:")
print(
    f"    Best validation accuracy : "
    f"{result['Best Validation Accuracy (%)']:.2f}%"
)

print(
    f"    Test accuracy           : "
    f"{result['Test Accuracy (%)']:.2f}%"
)

print(
    f"    Test loss                : "
    f"{result['Test Loss']:.4f}"
)

print(
    f"    Training time            : "

```

```

    f"{result['Training Time (s)'] / 60:.2f} minutes"
)

# -----
# Save final checkpoint manifest
# -----

manifest = {

    "stage": "final_model_completed",

    "final_model": {
        "learning_rate": float(
            result["Learning Rate"]
        ),
        "batch_size": int(
            result["Batch Size"]
        ),
        "dropout": float(
            result["Dropout"]
        ),
        "trainable_parameters": int(
            result["Trainable Parameters"]
        ),
        "total_parameters": int(
            result["Total Parameters"]
        )
    },

    "performance": {
        "best_validation_accuracy_percent":
            float(
                result[
                    "Best Validation Accuracy (%)"
                ]
            ),
        "test_accuracy_percent":
            float(
                result[
                    "Test Accuracy (%)"
                ]
            ),
        "test_loss":
            float(
                result["Test Loss"]
            )
    },

    "files": {
        "best_model":
            os.path.join(
                FINAL_DIR,
                "best_model.keras"
            ),
        "final_model":
            os.path.join(
                FINAL_DIR,
                "final_model.keras"
            ),
        "history":
            os.path.join(
                FINAL_DIR,
                "history.json"
            ),
        "results":
            FINAL_RESULTS_PATH
    }
}

manifest_path = os.path.join(
    FINAL_DIR,
    "final_model_manifest.json"
)

with open(
    manifest_path,
    "w"
) as f:

```

```

    json.dump(
        manifest,
        f,
        indent=2
    )

print("\n✓ Final-model checkpoint verified.")
print("✓ Results saved.")
print("✓ Manifest saved:")
print(" ", manifest_path)

print("\n" + "=" * 70)
print("R32 COMPLETE")
print("=" * 70)

```

```

=====
LAB 5 – FINAL MODEL CHECKPOINT
=====

```

```

Final model files:
✓ best_model.keras (9.72 MB)
✓ final_model.keras (9.71 MB)
✓ history.json (0.00 MB)

```

```

=====
FINAL MODEL RESULT
=====

```

	Model	Learning Rate	Batch Size	Dropout	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Test Loss
0	Final Selected Model	0.001	32	0.25	47397	2305381	91.168481	99.252719	90.625	0.059887	0.267998	0.31620

```

Key results:
Best validation accuracy : 91.17%
Test accuracy           : 89.72%
Test loss                : 0.3162
Training time            : 3.04 minutes

```

```

✓ Final-model checkpoint verified.
✓ Results saved.
✓ Manifest saved:
/content/drive/MyDrive/Lab5/Lab5/checkpoints/final_model_corrected/final_model_manifest.json

```

```

=====
R32 COMPLETE

```

```

# =====
# CELL R33 – FINAL MODEL CLASSIFICATION EVALUATION
# =====

```

```

import os
import gc
import numpy as np
import pandas as pd
import tensorflow as tf
import matplotlib.pyplot as plt

```

```

from sklearn.metrics import (
    confusion_matrix,
    classification_report,
    accuracy_score
)

```

```

# -----
# Paths
# -----

```

```

FINAL_DIR = os.path.join(
    CHECKPOINT_DIR,
    "final_model_corrected"
)

```

```

BEST_MODEL_PATH = os.path.join(
    FINAL_DIR,
    "best_model.keras"
)

```

```

CM_PATH = os.path.join(
    FINAL_DIR,
    "confusion_matrix.eps"
)

REPORT_PATH = os.path.join(
    FINAL_DIR,
    "classification_report.csv"
)

# -----
# Load best final model
# -----

if not os.path.exists(BEST_MODEL_PATH):
    raise FileNotFoundError(
        "Best final model not found:\n"
        + BEST_MODEL_PATH
    )

print("=" * 70)
print("FINAL MODEL - CLASSIFICATION EVALUATION")
print("=" * 70)

evaluation_model = keras.models.load_model(
    BEST_MODEL_PATH
)

print("✓ Best final model loaded.")

# -----
# Recreate test dataset
# -----

if "test_paths" not in globals():
    raise RuntimeError(
        "test_paths is not available."
    )

if "test_labels" not in globals():
    raise RuntimeError(
        "test_labels is not available."
    )

final_test_ds = make_final_dataset(
    test_paths,
    test_labels,
    FINAL_BATCH_SIZE,
    shuffle=False
)

# -----
# Generate predictions
# -----

print("\nGenerating predictions...")

predictions = evaluation_model.predict(
    final_test_ds,
    verbose=1
)

y_pred = np.argmax(
    predictions,
    axis=1
)

y_true = np.asarray(
    test_labels
).astype(int)

# -----
# Accuracy verification
# -----

accuracy = accuracy_score(
    y_true,
    y_pred
)

```

```

)

print("\n" + "=" * 70)
print("ACCURACY VERIFICATION")
print("=" * 70)

print(
    f"Test accuracy from predictions: "
    f"{accuracy * 100:.2f}%"
)

# -----
# Classification report
# -----

report = classification_report(
    y_true,
    y_pred,
    target_names=class_names,
    output_dict=True,
    zero_division=0
)

classification_df = pd.DataFrame(
    report
).transpose()

display(
    classification_df
)

# -----
# Save classification report
# -----

classification_df.to_csv(
    REPORT_PATH
)

print(
    "\n✓ Classification report saved:"
)
print(
    REPORT_PATH
)

# -----
# Confusion matrix
# -----

cm = confusion_matrix(
    y_true,
    y_pred
)

plt.figure(
    figsize=(16, 14)
)

plt.imshow(cm)

plt.title(
    "Final Model Confusion Matrix"
)

plt.xlabel(
    "Predicted Class"
)

plt.ylabel(
    "True Class"
)

plt.xticks(
    np.arange(len(class_names)),
    class_names,
    rotation=90
)

```

```
plt.yticks(
    np.arange(len(class_names)),
    class_names
)

plt.colorbar()

plt.tight_layout()

plt.savefig(
    CM_PATH,
    format="eps",
    dpi=600,
    bbox_inches="tight"
)

plt.show()

print(
    "\n✓ Confusion matrix saved:"
)
print(
    CM_PATH
)

# -----
# Cleanup
# -----

del evaluation_model
del final_test_ds
del predictions

tf.keras.backend.clear_session()
gc.collect()

print("\n" + "=" * 70)
print("R33 COMPLETE")
print("=" * 70)
```

```
=====
FINAL MODEL - CLASSIFICATION EVALUATION
=====
```

✓ Best final model loaded.

Generating predictions...

115/115  27s 170ms/step

```
=====
ACCURACY VERIFICATION
=====
```

Test accuracy from predictions: 89.72%

	precision	recall	f1-score	support
Abyssinian	0.793103	0.938776	0.859813	98.000000
Bengal	0.771930	0.880000	0.822430	100.000000
Birman	0.750000	0.810000	0.778846	100.000000
Bombay	0.849462	0.897727	0.872928	88.000000
British_Shorthair	0.851064	0.800000	0.824742	100.000000
Egyptian_Mau	0.939759	0.804124	0.866667	97.000000
Maine_Coon	0.870968	0.810000	0.839378	100.000000
Persian	0.875000	0.910000	0.892157	100.000000
Ragdoll	0.693069	0.700000	0.696517	100.000000
Russian_Blue	0.855556	0.770000	0.810526	100.000000
Siamese	0.964286	0.810000	0.880435	100.000000
Sphynx	0.988764	0.880000	0.931217	100.000000
american_bulldog	0.700787	0.890000	0.784141	100.000000
american_pit_bull_terrier	0.776316	0.590000	0.670455	100.000000
basset_hound	0.959184	0.940000	0.949495	100.000000
beagle	0.913462	0.950000	0.931373	100.000000
boxer	0.854369	0.888889	0.871287	99.000000
chihuahua	0.900990	0.910000	0.905473	100.000000
english_cocker_spaniel	0.949495	0.940000	0.944724	100.000000
english_setter	0.960000	0.960000	0.960000	100.000000
german_shorthaired	0.970297	0.980000	0.975124	100.000000
great_pyrenees	0.980392	1.000000	0.990099	100.000000
havanese	0.950000	0.950000	0.950000	100.000000
japanese_chin	0.990000	0.990000	0.990000	100.000000

```
# =====
# CELL R34 - FINAL EXPERIMENT SUMMARY
# =====
```

```
import os
import pandas as pd
```

```
print("=" * 70)
print("LAB 5 - COMPLETE EXPERIMENT SUMMARY")
print("=" * 70)
```

```
# -----
# Helper
# -----
```

```
def load_csv_if_exists(filename):
    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):
        return pd.read_csv(path)

    return None
```

```
# -----
```

```

# Load all completed experiment results
# -----

initialization_df = load_csv_if_exists(
    "initialization_results.csv"
)

regularization_df = load_csv_if_exists(
    "regularization_results.csv"
)

batch_normalization_df = load_csv_if_exists(
    "batch_normalization_results.csv"
)

optimizer_df = load_csv_if_exists(
    "optimizer_results.csv"
)

hyperparameter_df = load_csv_if_exists(
    "hyperparameter_results.csv"
)

transfer_df = load_csv_if_exists(
    "transfer_learning_results.csv"
)

fine_tune_df = load_csv_if_exists(
    "correct_fine_tuning_results.csv"
)

cv_results_df = load_csv_if_exists(
    "cv_results.csv"
)

cv_summary_df = load_csv_if_exists(
    "cv_summary.csv"
)

final_df = load_csv_if_exists(
    "final_model_results.csv"
)

# -----
# Display completion status
# -----

experiments = {
    "Initialization": initialization_df,
    "Regularization": regularization_df,
    "Batch Normalization": batch_normalization_df,
    "Optimizer": optimizer_df,
    "Hyperparameters": hyperparameter_df,
    "Transfer Learning": transfer_df,
    "Fine-Tuning": fine_tune_df,
    "Cross-Validation": cv_results_df,
    "CV Summary": cv_summary_df,
    "Final Model": final_df
}

print("\nCompleted experiments:")

for name, df in experiments.items():

    if df is None:
        print(f"  X {name}")
    else:
        print(
            f"  ✓ {name} "
            f"({len(df)} row(s))"
        )

# -----
# Final model result
# -----

if final_df is not None:

    print("\n" + "=" * 70)
    print("FINAL MODEL")

```

```

print("=" * 70)

display(final_df)

# -----
# Cross-validation result
# -----

if cv_summary_df is not None:

    print("\n" + "=" * 70)
    print("CROSS-VALIDATION")
    print("=" * 70)

    display(cv_summary_df)

# -----
# Save consolidated summary
# -----

summary_rows = []

if final_df is not None:

    row = final_df.iloc[0]

    summary_rows.append({

        "Final Model":
            row["Model"],

        "Learning Rate":
            row["Learning Rate"],

        "Batch Size":
            row["Batch Size"],

        "Dropout":
            row["Dropout"],

        "Best Validation Accuracy (%)":
            row["Best Validation Accuracy (%)"],

        "Test Accuracy (%)":
            row["Test Accuracy (%)"],

        "Test Loss":
            row["Test Loss"],

        "Training Time (s)":
            row["Training Time (s)"]

    })

summary_df = pd.DataFrame(
    summary_rows
)

SUMMARY_PATH = os.path.join(
    RESULTS_DIR,
    "lab5_final_summary.csv"
)

summary_df.to_csv(
    SUMMARY_PATH,
    index=False
)

print("\n" + "=" * 70)
print("R34 COMPLETE")
print("=" * 70)

print(
    "\n✓ Consolidated summary saved:"
)

print(

```

```

SUMMARY_PATH
)

```

LAB 5 – COMPLETE EXPERIMENT SUMMARY

Completed experiments:

- ✓ Initialization (4 row(s))
- ✓ Regularization (4 row(s))
- ✓ Batch Normalization (1 row(s))
- ✓ Optimizer (4 row(s))
- ✓ Hyperparameters (5 row(s))
- ✓ Transfer Learning (2 row(s))
- ✓ Fine-Tuning (1 row(s))
- ✓ Cross-Validation (10 row(s))
- ✓ CV Summary (2 row(s))
- ✓ Final Model (1 row(s))

FINAL MODEL

	Model	Learning Rate	Batch Size	Dropout	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Test Loss
0	Final Selected Model	0.001	32	0.25	47397	2305381	91.168481	99.252719	90.625	0.059887	0.267998	0.31620

CROSS-VALIDATION

	Configuration	Mean	SD	Mean Training Time (s)
0	Baseline	91.881607	1.449007	722.465817
1	Fine-Tuned	89.198631	1.019683	926.829785

R34 COMPLETE

✓ Consolidated summary saved:

```

# =====
# CELL R35 – EXPERIMENT COMPARISON TABLE
# =====

import os
import pandas as pd
import numpy as np

print("=" * 70)
print("LAB 5 – EXPERIMENT COMPARISON")
print("=" * 70)

# -----
# Load completed experiment results
# -----

def load_result(filename):
    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):
        return pd.read_csv(path)

    return None

initialization_df = load_result(
    "initialization_results.csv"
)

regularization_df = load_result(
    "regularization_results.csv"
)

```

```

batch_norm_df = load_result(
    "batch_normalization_results.csv"
)

optimizer_df = load_result(
    "optimizer_results.csv"
)

hyperparameter_df = load_result(
    "hyperparameter_results.csv"
)

transfer_df = load_result(
    "transfer_learning_results.csv"
)

fine_tuning_df = load_result(
    "correct_fine_tuning_results.csv"
)

cv_summary_df = load_result(
    "cv_summary.csv"
)

final_df = load_result(
    "final_model_results.csv"
)

# -----
# Build comparison table
# -----

comparison_rows = []

# Initialization
if initialization_df is not None:
    for _, row in initialization_df.iterrows():
        comparison_rows.append({
            "Experiment": "Initialization",
            "Configuration":
                row.get("Initialization", "N/A"),
            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),
            "Final Validation Accuracy (%)":
                row.get(
                    "Final Validation Accuracy (%)",
                    np.nan
                )
        })

# Regularization
if regularization_df is not None:
    for _, row in regularization_df.iterrows():
        comparison_rows.append({
            "Experiment": "Regularization",
            "Configuration":
                row.get("Method", "N/A"),
            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),
            "Final Validation Accuracy (%)":

```

```

        row.get(
            "Final Validation Accuracy (%)",
            np.nan
        )

    })

# Batch Normalization
if batch_norm_df is not None:

    for _, row in batch_norm_df.iterrows():

        comparison_rows.append({

            "Experiment": "Batch Normalization",

            "Configuration":
                row.get("Method", "N/A"),

            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),

            "Final Validation Accuracy (%)":
                row.get(
                    "Final Validation Accuracy (%)",
                    np.nan
                )

        })

# Optimizer
if optimizer_df is not None:

    for _, row in optimizer_df.iterrows():

        comparison_rows.append({

            "Experiment": "Optimizer",

            "Configuration":
                row.get("Optimizer", "N/A"),

            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),

            "Final Validation Accuracy (%)":
                row.get(
                    "Final Validation Accuracy (%)",
                    np.nan
                )

        })

# Hyperparameters
if hyperparameter_df is not None:

    for _, row in hyperparameter_df.iterrows():

        comparison_rows.append({

            "Experiment": "Hyperparameter",

            "Configuration":
                f"{row.get('Hyperparameter', 'N/A')} = "
                f"{row.get('Value', 'N/A')}",

            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),

            "Final Validation Accuracy (%)":

```

```

        row.get(
            "Final Validation Accuracy (%)",
            np.nan
        )
    })

# Transfer learning
if transfer_df is not None:
    for _, row in transfer_df.iterrows():
        comparison_rows.append({
            "Experiment": "Transfer Learning",
            "Configuration":
                row.get("Method", "N/A"),
            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),
            "Final Validation Accuracy (%)":
                row.get(
                    "Final Validation Accuracy (%)",
                    np.nan
                )
        })

# Fine tuning
if fine_tuning_df is not None:
    for _, row in fine_tuning_df.iterrows():
        comparison_rows.append({
            "Experiment": "Fine-Tuning",
            "Configuration":
                row.get("Method", "Fine-Tuning"),
            "Best Validation Accuracy (%)":
                row.get(
                    "Best Validation Accuracy (%)",
                    np.nan
                ),
            "Final Validation Accuracy (%)":
                row.get(
                    "Final Validation Accuracy (%)",
                    np.nan
                )
        })

# -----
# Add CV summary
# -----

if cv_summary_df is not None:
    for _, row in cv_summary_df.iterrows():
        comparison_rows.append({
            "Experiment": "Cross-Validation",
            "Configuration":
                row.get("Configuration", "N/A"),
            "Best Validation Accuracy (%)":
                row.get("Mean", np.nan),
            "Final Validation Accuracy (%)":
                np.nan
        })

```

```

    })

# -----
# Add final model
# -----

if final_df is not None:

    row = final_df.iloc[0]

    comparison_rows.append({

        "Experiment": "Final Model",

        "Configuration":
            "Selected Configuration",

        "Best Validation Accuracy (%)":
            row.get(
                "Best Validation Accuracy (%)",
                np.nan
            ),

        "Final Validation Accuracy (%)":
            row.get(
                "Final Validation Accuracy (%)",
                np.nan
            )

    })

# -----
# Create DataFrame
# -----

comparison_df = pd.DataFrame(
    comparison_rows
)

if comparison_df.empty:

    raise RuntimeError(
        "No experiment results were found."
    )

# -----
# Round numerical values
# -----

numeric_columns = [
    "Best Validation Accuracy (%)",
    "Final Validation Accuracy (%)"
]

for column in numeric_columns:

    if column in comparison_df.columns:

        comparison_df[column] = (
            pd.to_numeric(
                comparison_df[column],
                errors="coerce"
            ).round(2)
        )

# -----
# Display
# -----

display(
    comparison_df
)

# -----
# Save

```

```
# -----

COMPARISON_PATH = os.path.join(
    RESULTS_DIR,
    "experiment_comparison.csv"
)

comparison_df.to_csv(
    COMPARISON_PATH,
    index=False
)

print("\n✓ Comparison table saved:")
print(COMPARISON_PATH)

print("\n" + "=" * 70)
print("R35 COMPLETE")
print("\n" * 70)
```

```
=====
LAB 5 – EXPERIMENT COMPARISON
=====
```

	Experiment	Configuration	Best Validation Accuracy (%)	Final Validation Accuracy (%)
0	Initialization	Zero	93.21	93.21
1	Initialization	Random	91.98	91.98
2	Initialization	Xavier	91.85	91.85
3	Initialization	He	91.30	90.76
4	Regularization	N/A	91.85	91.85
5	Regularization	N/A	92.26	92.26
6	Regularization	N/A	92.53	92.26
7	Regularization	N/A	91.03	90.90
8	Batch Normalization	N/A	91.17	91.17
9	Optimizer	SGD	36.14	36.14
10	Optimizer	Momentum	88.45	88.45
11	Optimizer	RMSProp	91.58	91.58
12	Optimizer	Adam	91.58	90.62
13	Hyperparameter	Learning Rate = 0.001	96.47	96.47
14	Hyperparameter	Learning Rate = 0.0001	78.53	78.53
15	Hyperparameter	Dropout = 0.0	90.90	90.90
16	Hyperparameter	Dropout = 0.25	91.17	90.08
17	Hyperparameter	Dropout = 0.5	90.22	90.08
18	Transfer Learning	N/A	90.76	90.76
19	Transfer Learning	N/A	89.27	89.27
20	Fine-Tuning	Fine-Tuning	89.27	89.27
21	Cross-Validation	Baseline	91.88	NaN
22	Cross-Validation	Fine-Tuned	89.20	NaN
23	Final Model	Selected Configuration	91.17	90.62

```
✓ Comparison table saved:
/content/drive/MyDrive/Lab5/Lab5/results/experiment_comparison.csv
```

```
=====
R35 COMPLETE
=====
```

```
# =====
# CELL R36 – FINAL LAB 5 CHECKPOINT
# =====

import os
import json
import shutil
import time
import pandas as pd
```

```

FINAL_CHECKPOINT_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "final_checkpoint"
)

os.makedirs(
    FINAL_CHECKPOINT_DIR,
    exist_ok=True
)

print("=" * 70)
print("LAB 5 – CREATING FINAL CHECKPOINT")
print("=" * 70)

print("\nCheckpoint:")
print(FINAL_CHECKPOINT_DIR)

# -----
# Files/directories to preserve
# -----

items_to_copy = [

    # Results
    (
        os.path.join(RESULTS_DIR, "initialization_results.csv"),
        "initialization_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "regularization_results.csv"),
        "regularization_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "batch_normalization_results.csv"),
        "batch_normalization_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "optimizer_results.csv"),
        "optimizer_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "hyperparameter_results.csv"),
        "hyperparameter_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "transfer_learning_results.csv"),
        "transfer_learning_results.csv"
    ),
    (
        os.path.join(
            RESULTS_DIR,
            "correct_fine_tuning_results.csv"
        ),
        "correct_fine_tuning_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "cv_results.csv"),
        "cv_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "cv_summary.csv"),
        "cv_summary.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "final_model_results.csv"),
        "final_model_results.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "experiment_comparison.csv"),
        "experiment_comparison.csv"
    ),
    (
        os.path.join(RESULTS_DIR, "lab5_final_summary.csv"),
        "lab5_final_summary.csv"
    )
]

# -----
# Copy files
# -----

```

```

copied_files = []
missing_files = []

for source, filename in items_to_copy:

    destination = os.path.join(
        FINAL_CHECKPOINT_DIR,
        filename
    )

    if os.path.exists(source):

        shutil.copy2(
            source,
            destination
        )

        copied_files.append(filename)

    else:

        missing_files.append(filename)

# -----
# Preserve final model
# -----

final_model_source = os.path.join(
    CHECKPOINT_DIR,
    "final_model_corrected"
)

final_model_destination = os.path.join(
    FINAL_CHECKPOINT_DIR,
    "final_model_corrected"
)

if os.path.exists(final_model_source):

    if os.path.exists(final_model_destination):
        shutil.rmtree(final_model_destination)

    shutil.copytree(
        final_model_source,
        final_model_destination
    )

    print(
        "\n✓ Final model checkpoint copied."
    )

else:

    print(
        "\n△ Final model checkpoint not found."
    )

# -----
# Preserve fine-tuning checkpoint
# -----

fine_source = os.path.join(
    CHECKPOINT_DIR,
    "fine_tuning_final"
)

fine_destination = os.path.join(
    FINAL_CHECKPOINT_DIR,
    "fine_tuning_final"
)

if os.path.exists(fine_source):

    if os.path.exists(fine_destination):
        shutil.rmtree(fine_destination)

    shutil.copytree(
        fine_source,
        fine_destination
    )

```

```
)

print(
    "✓ Fine-tuning checkpoint copied."
)

# -----
# Preserve CV checkpoints
# -----

cv_source = os.path.join(
    CHECKPOINT_DIR,
    "cross_validation"
)

cv_destination = os.path.join(
    FINAL_CHECKPOINT_DIR,
    "cross_validation"
)

if os.path.exists(cv_source):

    if os.path.exists(cv_destination):
        shutil.rmtree(cv_destination)

    shutil.copytree(
        cv_source,
        cv_destination
    )

    print(
        "✓ Cross-validation checkpoints copied."
    )

# -----
# Create manifest
# -----

manifest = {

    "checkpoint": "LAB 5 FINAL CHECKPOINT",

    "created": time.strftime(
        "%Y-%m-%d %H:%M:%S"
    ),

    "stage":
        "Final model evaluation completed",

    "gpu_required_for_completed_results":
        False,

    "final_model_completed":
        True,

    "final_model_result":
        "89.72% test accuracy",

    "copied_result_files":
        copied_files,

    "missing_result_files":
        missing_files,

    "final_model_checkpoint":
        os.path.exists(final_model_destination),

    "fine_tuning_checkpoint":
        os.path.exists(fine_destination),

    "cross_validation_checkpoint":
        os.path.exists(cv_destination)
}

manifest_path = os.path.join(
    FINAL_CHECKPOINT_DIR,
    "final_checkpoint_manifest.json"
)
```

```

with open(
    manifest_path,
    "w"
) as f:

    json.dump(
        manifest,
        f,
        indent=2
    )

# -----
# Summary
# -----

print("\n" + "=" * 70)
print("FINAL CHECKPOINT CREATED")
print("=" * 70)

print(
    f"\nResult files copied: "
    f"{len(copied_files)}"
)

if missing_files:

    print(
        "\nMissing result files:"
    )

    for filename in missing_files:
        print("  △", filename)

else:

    print(
        "✓ All expected result files found."
    )

print(
    "\nCheckpoint location:"
)

print(
    FINAL_CHECKPOINT_DIR
)

print(
    "\n✓ Manifest saved:"
)

print(
    manifest_path
)

print("\n" + "=" * 70)
print("R36 COMPLETE")
print("=" * 70)

```

```

=====
LAB 5 — CREATING FINAL CHECKPOINT
=====

```

```

Checkpoint:
/content/drive/MyDrive/Lab5/Lab5/final_checkpoint

```

```

✓ Final model checkpoint copied.
✓ Fine-tuning checkpoint copied.
✓ Cross-validation checkpoints copied.

```

```

=====
FINAL CHECKPOINT CREATED
=====

```

```

Result files copied: 12
✓ All expected result files found.

```

```

Checkpoint location:
/content/drive/MyDrive/Lab5/Lab5/final_checkpoint

```

```

✓ Manifest saved:
/content/drive/MyDrive/Lab5/Lab5/final_checkpoint/final_checkpoint_manifest.json

```

```
=====
R36 COMPLETE
=====
```

```
# =====
# CELL R37 - FINAL LAB 5 VERIFICATION
# =====

import os
import json
import pandas as pd

print("=" * 70)
print("LAB 5 - FINAL VERIFICATION")
print("=" * 70)

# -----
# Required result files
# -----

required_results = [
    "initialization_results.csv",
    "regularization_results.csv",
    "batch_normalization_results.csv",
    "optimizer_results.csv",
    "hyperparameter_results.csv",
    "transfer_learning_results.csv",
    "correct_fine_tuning_results.csv",
    "cv_results.csv",
    "cv_summary.csv",
    "final_model_results.csv",
    "experiment_comparison.csv",
    "lab5_final_summary.csv"
]

print("\nRESULT FILES")
print("-" * 70)

missing_results = []

for filename in required_results:

    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):

        print("✓", filename)

    else:

        print("X", filename)

        missing_results.append(
            filename
        )

# -----
# Final model files
# -----

final_model_dir = os.path.join(
    CHECKPOINT_DIR,
    "final_model_corrected"
)

final_model_files = [
    "best_model.keras",
    "final_model.keras",
    "history.json"
]

print("\nFINAL MODEL FILES")
print("-" * 70)

missing_model_files = []
```

```

for filename in final_model_files:

    path = os.path.join(
        final_model_dir,
        filename
    )

    if os.path.exists(path):

        print("✓", filename)

    else:

        print("X", filename)

        missing_model_files.append(
            filename
        )

# -----
# Load final result
# -----

final_results_path = os.path.join(
    RESULTS_DIR,
    "final_model_results.csv"
)

print("\nFINAL MODEL RESULT")
print("-" * 70)

if os.path.exists(final_results_path):

    final_results = pd.read_csv(
        final_results_path
    )

    display(final_results)

    if len(final_results) > 0:

        final_row = final_results.iloc[0]

        print(
            f"\nReported test accuracy: "
            f"{float(final_row['Test Accuracy (%)']):.2f}%"
        )

    else:

        final_results = None

        print(
            "X Final model results CSV not found."
        )

# -----
# Verify confusion matrix
# -----

cm_path = os.path.join(
    final_model_dir,
    "confusion_matrix.eps"
)

print("\nDIAGNOSTIC OUTPUTS")
print("-" * 70)

if os.path.exists(cm_path):

    print("✓ Confusion matrix EPS")

else:

    print("△ Confusion matrix EPS not found.")

report_path = os.path.join(
    final_model_dir,
    "classification_report.csv"
)

```

```

if os.path.exists(report_path):
    print("✓ Classification report")
else:
    print("⚠ Classification report not found.")

# -----
# Final status
# -----

print("\n" + "=" * 70)

if not missing_results and not missing_model_files:

    print("✓ LAB 5 FINAL VERIFICATION PASSED")

else:

    print("⚠ LAB 5 VERIFICATION FOUND MISSING FILES")

    if missing_results:

        print("\nMissing result files:")

        for filename in missing_results:
            print("  -", filename)

    if missing_model_files:

        print("\nMissing final-model files:")

        for filename in missing_model_files:
            print("  -", filename)

print("=" * 70)

```

```

=====
LAB 5 – FINAL VERIFICATION
=====

```

RESULT FILES

```

-----
✓ initialization_results.csv
✓ regularization_results.csv
✓ batch_normalization_results.csv
✓ optimizer_results.csv
✓ hyperparameter_results.csv
✓ transfer_learning_results.csv
✓ correct_fine_tuning_results.csv
✓ cv_results.csv
✓ cv_summary.csv
✓ final_model_results.csv
✓ experiment_comparison.csv
✓ lab5_final_summary.csv

```

FINAL MODEL FILES

```

-----
✓ best_model.keras
✓ final_model.keras
✓ history.json

```

FINAL MODEL RESULT

```

-----

```

	Model	Learning Rate	Batch Size	Dropout	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Test Loss
0	Final Selected Model	0.001	32	0.25	47397	2305381	91.168481	99.252719	90.625	0.059887	0.267998	0.31620

Reported test accuracy: 89.72%

DIAGNOSTIC OUTPUTS

```

-----
✓ Confusion matrix EPS
✓ Classification report

```

```

=====
✓ LAB 5 FINAL VERIFICATION PASSED

```

```
# =====
# CELL R38 - PRINT ALL REQUIRED LAB 5 RESULTS
# =====

import os
import json
import numpy as np
import pandas as pd

print("=" * 80)
print("LAB 5 - ALL REQUIRED RESULTS")
print("=" * 80)

# =====
# 1. PATHS
# =====

RESULTS_DIR = "/content/drive/MyDrive/Lab5/Lab5/results"

FINAL_DIR = (
    "/content/drive/MyDrive/Lab5/Lab5/"
    "checkpoints/final_model_corrected"
)

# =====
# 2. HELPER
# =====

def load_csv(filename):

    path = os.path.join(
        RESULTS_DIR,
        filename
    )

    if os.path.exists(path):

        return pd.read_csv(path)

    return None

def get_column(df, possible_names):

    if df is None:
        return np.nan

    for name in possible_names:

        if name in df.columns:
            return df[name].iloc[0]

    return np.nan

def get_best_accuracy(df):

    if df is None or len(df) == 0:
        return np.nan

    candidates = [
        "Best Validation Accuracy (%)",
        "Final Validation Accuracy (%)",
        "Validation Accuracy (%)",
        "Mean"
    ]

    for column in candidates:

        if column in df.columns:

            values = pd.to_numeric(
                df[column],
                errors="coerce"
            )

            if values.notna().any():

                return values.max()
```

```

        return np.nan

# =====
# 3. LOAD ALL EXPERIMENT RESULTS
# =====

initialization_df = load_csv(
    "initialization_results.csv"
)

regularization_df = load_csv(
    "regularization_results.csv"
)

batch_norm_df = load_csv(
    "batch_normalization_results.csv"
)

optimizer_df = load_csv(
    "optimizer_results.csv"
)

hyperparameter_df = load_csv(
    "hyperparameter_results.csv"
)

transfer_df = load_csv(
    "transfer_learning_results.csv"
)

fine_tuning_df = load_csv(
    "correct_fine_tuning_results.csv"
)

cv_results_df = load_csv(
    "cv_results.csv"
)

cv_summary_df = load_csv(
    "cv_summary.csv"
)

final_df = load_csv(
    "final_model_results.csv"
)

# =====
# 4. INITIALIZATION
# =====

print("\n" + "=" * 80)
print("1. WEIGHT INITIALIZATION")
print("=" * 80)

if initialization_df is not None:

    display(initialization_df)

else:

    print("Initialization results not found.")

# =====
# 5. REGULARIZATION
# =====

print("\n" + "=" * 80)
print("2. REGULARIZATION")
print("=" * 80)

if regularization_df is not None:

    display(regularization_df)

else:

    print("Regularization results not found.")

```

```
# =====
# 6. BATCH NORMALIZATION
# =====

print("\n" + "=" * 80)
print("3. BATCH NORMALIZATION")
print("=" * 80)

if batch_norm_df is not None:

    display(batch_norm_df)

else:

    print("Batch Normalization results not found.")

# =====
# 7. OPTIMIZERS
# =====

print("\n" + "=" * 80)
print("4. OPTIMIZATION ALGORITHMS")
print("=" * 80)

if optimizer_df is not None:

    display(optimizer_df)

else:

    print("Optimizer results not found.")

# =====
# 8. HYPERPARAMETER TUNING
# =====

print("\n" + "=" * 80)
print("5. HYPERPARAMETER TUNING")
print("=" * 80)

if hyperparameter_df is not None:

    display(hyperparameter_df)

else:

    print("Hyperparameter results not found.")

# =====
# 9. TRANSFER LEARNING
# =====

print("\n" + "=" * 80)
print("6. TRANSFER LEARNING")
print("=" * 80)

if transfer_df is not None:

    display(transfer_df)

else:

    print("Transfer-learning results not found.")

# =====
# 10. FINE-TUNING
# =====

print("\n" + "=" * 80)
print("7. FINE-TUNING")
print("=" * 80)

if fine_tuning_df is not None:

    display(fine_tuning_df)

else:
```

```

    print("Fine-tuning results not found.")

# =====
# 11. CROSS-VALIDATION – INDIVIDUAL FOLDS
# =====

print("\n" + "=" * 80)
print("8. 5-FOLD CROSS-VALIDATION – INDIVIDUAL FOLDS")
print("=" * 80)

if cv_results_df is not None:

    display(cv_results_df)

else:

    print("Cross-validation results not found.")

# =====
# 12. CROSS-VALIDATION SUMMARY
# =====

print("\n" + "=" * 80)
print("9. 5-FOLD CROSS-VALIDATION SUMMARY")
print("=" * 80)

if cv_summary_df is not None:

    display(cv_summary_df)

else:

    print("Cross-validation summary not found.")

# =====
# 13. FINAL MODEL
# =====

print("\n" + "=" * 80)
print("10. FINAL MODEL")
print("=" * 80)

if final_df is not None:

    display(final_df)

else:

    print("Final model results not found.")

# =====
# 14. REQUIRED OVERALL RESULTS TABLE
# =====

print("\n" + "=" * 80)
print("11. OVERALL RESULTS – REQUIRED BY EXPERIMENT")
print("=" * 80)

overall_rows = []

# -----
# Baseline
# -----

baseline_accuracy = np.nan

if cv_summary_df is not None:

    baseline_rows = cv_summary_df[
        cv_summary_df["Configuration"]
        .astype(str)
        .str.lower()
        .str.contains("baseline")
    ]

```

```

        if len(baseline_rows) > 0:

            baseline_accuracy = float(
                baseline_rows["Mean"].iloc[0]
            )

overall_rows.append({

    "Configuration":
        "Baseline",

    "CV Accuracy (%)":
        baseline_accuracy,

    "SD (%)":
        (
            baseline_rows["SD"].iloc[0]
            if len(baseline_rows) > 0
            else np.nan
        ),

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":
        (
            baseline_rows[
                "Mean Training Time (s)"
            ].iloc[0]
            if (
                len(baseline_rows) > 0
                and "Mean Training Time (s)"
                in baseline_rows.columns
            )
            else np.nan
        )
    })

# -----
# Best Initialization
# -----

overall_rows.append({

    "Configuration":
        "Best Initialization",

    "CV Accuracy (%)":
        np.nan,

    "SD (%)":
        np.nan,

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":
        np.nan
    })

# -----
# Best Regularization
# -----

overall_rows.append({

    "Configuration":
        "Best Regularization",

    "CV Accuracy (%)":
        np.nan,

    "SD (%)":
        np.nan,

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":

```

```

        np.nan
    })

# -----
# Best Optimizer
# -----

overall_rows.append({

    "Configuration":
        "Best Optimizer",

    "CV Accuracy (%)":
        np.nan,

    "SD (%)":
        np.nan,

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":
        np.nan
})

# -----
# Best Hyperparameters
# -----

overall_rows.append({

    "Configuration":
        "Best Hyperparameters",

    "CV Accuracy (%)":
        np.nan,

    "SD (%)":
        np.nan,

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":
        np.nan
})

# -----
# Fine-Tuned Model
# -----

fine_cv_accuracy = np.nan
fine_cv_sd = np.nan
fine_cv_time = np.nan

if cv_summary_df is not None:

    fine_rows = cv_summary_df[
        cv_summary_df["Configuration"]
            .astype(str)
            .str.lower()
            .str.contains("fine")
    ]

    if len(fine_rows) > 0:

        fine_cv_accuracy = float(
            fine_rows["Mean"].iloc[0]
        )

        fine_cv_sd = float(
            fine_rows["SD"].iloc[0]
        )

        if "Mean Training Time (s)" in fine_rows:

            fine_cv_time = float(
                fine_rows[
                    "Mean Training Time (s)"

```

```

        ].iloc[0]
    )

overall_rows.append({

    "Configuration":
        "Fine-Tuned Model",

    "CV Accuracy (%)":
        fine_cv_accuracy,

    "SD (%)":
        fine_cv_sd,

    "Test Accuracy (%)":
        np.nan,

    "Training Time (s)":
        fine_cv_time
})

overall_df = pd.DataFrame(
    overall_rows
)

display(overall_df)

# =====
# 15. REQUIRED FINAL MODEL METRICS
# =====

print("\n" + "=" * 80)
print("12. REQUIRED FINAL MODEL METRICS")
print("=" * 80)

final_metrics = {

    "Mean CV Accuracy (%)":
        np.nan,

    "CV Standard Deviation (%)":
        np.nan,

    "Test Accuracy (%)":
        np.nan,

    "Precision":
        np.nan,

    "Recall":
        np.nan,

    "F1-score":
        np.nan,

    "Training Time (s)":
        np.nan,

    "Number of Parameters":
        np.nan
}

# -----
# CV metrics
# -----

if cv_summary_df is not None:

    # Prefer the selected/final configuration.
    selected_cv = cv_summary_df[
        cv_summary_df["Configuration"]
        .astype(str)
        .str.lower()
        .str.contains("baseline")
    ]

```

```

if len(selected_cv) == 0:

    selected_cv = cv_summary_df.iloc[
        [cv_summary_df["Mean"].idxmax()]
    ]

    final_metrics[
        "Mean CV Accuracy (%)"
    ] = float(
        selected_cv["Mean"].iloc[0]
    )

    final_metrics[
        "CV Standard Deviation (%)"
    ] = float(
        selected_cv["SD"].iloc[0]
    )

# -----
# Final model CSV
# -----

if final_df is not None and len(final_df) > 0:

    row = final_df.iloc[0]

    if "Test Accuracy (%)" in final_df.columns:

        final_metrics[
            "Test Accuracy (%)"
        ] = float(
            row["Test Accuracy (%)"]
        )

    if "Training Time (s)" in final_df.columns:

        final_metrics[
            "Training Time (s)"
        ] = float(
            row["Training Time (s)"]
        )

    if "Total Parameters" in final_df.columns:

        final_metrics[
            "Number of Parameters"
        ] = int(
            row["Total Parameters"]
        )

# -----
# Classification report
# -----

report_path = os.path.join(
    FINAL_DIR,
    "classification_report.csv"
)

if os.path.exists(report_path):

    report_df = pd.read_csv(
        report_path,
        index_col=0
    )

    if "weighted avg" in report_df.index:

        final_metrics["Precision"] = float(
            report_df.loc[
                "weighted avg",
                "precision"
            ]
        )

        final_metrics["Recall"] = float(
            report_df.loc[
                "weighted avg",
                "recall"
            ]
        )

```

```
)

final_metrics["F1-score"] = float(
    report_df.loc[
        "weighted avg",
        "f1-score"
    ]
)

final_metrics_df = pd.DataFrame(
    [final_metrics]
)

display(
    final_metrics_df.round(4)
)

# =====
# 16. FINAL SELECTED CONFIGURATION
# =====

print("\n" + "=" * 80)
print("13. FINAL SELECTED CONFIGURATION")
print("=" * 80)

print("""
Learning Rate : 0.001
Batch Size    : 32
Dropout       : 0.25
Epochs       : 10
Architecture  : MobileNetV2
Classes       : 37
""")

# =====
# 17. FINAL STATUS
# =====

print("=" * 80)
print("LAB 5 REQUIRED RESULTS PRINTED")
print("=" * 80)

print("\nImportant final result:")
print("  Test Accuracy = 89.72%")

print("\nNo training was performed by this cell.")
print("All values were loaded from the saved experiment results.")
```

=====

LAB 5 – ALL REQUIRED RESULTS

=====

=====

1. WEIGHT INITIALIZATION

=====

	Initialization	Best Validation Accuracy (%)	Final Validation Accuracy (%)	Final Training Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Zero	93.206519	93.206519	98.879075	0.096306	0.241588	194.542761
1	Random	91.983694	91.983694	98.471469	0.104582	0.263262	209.935266
2	Xavier	91.847825	91.847825	98.403531	0.105380	0.256203	192.279224
3	He	91.304350	90.760869	98.505437	0.101310	0.249884	183.985518

=====

2. REGULARIZATION

=====

	Regularization	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	NaN	91.847825	98.947012	91.847825	0.099481	0.250932	196.603703
1	L2	92.255437	98.573369	92.255437	0.220735	0.367931	191.783631
2	Dropout	92.527175	93.070650	92.255437	0.230227	0.261738	189.963952
3	Batch Normalization	91.032606	99.558425	90.896738	0.048230	0.275540	185.832275

=====

3. BATCH NORMALIZATION

=====

	Configuration	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	Without BN	91.168481	98.675269	91.168481	0.103367	0.270477	261.090611

=====

4. OPTIMIZATION ALGORITHMS

=====

	Optimizer	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Training Time (s)
0	SGD	36.141303	33.899456	36.141303	2.713774	2.653474	286.185142
1	Momentum	88.451087	90.760869	88.451087	0.510050	0.554425	194.851837
2	RMSProp	91.576087	97.961956	91.576087	0.091073	0.248172	210.569297
3	Adam	91.576087	98.573369	90.625000	0.100249	0.274059	186.688971

```
# =====
# CELL R31A – RESTORE 5-FOLD INDICES FOR ADDITIONAL EXERCISE
# =====

import os
import numpy as np
from sklearn.model_selection import StratifiedKFold

print("=" * 70)
print("RESTORING 5-FOLD CROSS-VALIDATION INDICES")
print("=" * 70)

# -----
# Verify dataset variables
# -----

required = [
    "train_paths",
    "train_labels"
]

missing = [
    x for x in required
    if x not in globals()
]

if missing:
```

```

    raise RuntimeError(
        "Required dataset variables are missing: "
        + str(missing)
        + "\nRun the dataset restoration cell first."
    )

# -----
# Recreate the EXACT 5-fold split
# -----

skf = StratifiedKFold(
    n_splits=5,
    shuffle=True,
    random_state=SEED
)

folds = []

for train_idx, val_idx in skf.split(
    train_paths,
    train_labels
):

    folds.append(
        (
            train_idx,
            val_idx
        )
    )

# -----
# Verify
# -----

print(
    f"\nNumber of folds: {len(folds)}"
)

for i, (
    train_idx,
    val_idx
) in enumerate(
    folds,
    start=1
):

    print(
        f"Fold {i}: "
        f"Train = {len(train_idx)}, "
        f"Validation = {len(val_idx)}"
    )

if len(folds) != 5:

    raise RuntimeError(
        "Expected exactly 5 folds."
    )

print("\n" + "=" * 70)
print("5-FOLD INDICES RESTORED")
print("=" * 70)

print(
    "✓ Ready to run R32."
)

```

```

0          91.8816          1.449          89.7247          0.9  0.8972  0.8969  182.592  2305381
=====
RESTORING 5-FOLD CROSS-VALIDATION INDICES
=====
13. FINAL SELECTED CONFIGURATION
Number-of-folds=5=====
Fold 1: Train = 2355, Validation = 589
Fold 2: Train = 2355, Validation = 589
Fold 3: Train = 2355, Validation = 589
Fold 4: Train = 2355, Validation = 589
Fold 5: Train = 2356, Validation = 588

```

```
=====
5-FOLD INDICES RESTORED
=====
✓ Ready to run R32.
```

```
# =====
# CELL R32 – ADDITIONAL EXERCISE: TWO NEW CONFIGURATIONS
# =====

import os
import gc
import time
import json
import numpy as np
import pandas as pd
import tensorflow as tf

from tensorflow import keras
from tensorflow.keras import layers
from tensorflow.keras.applications import MobileNetV2

print("=" * 70)
print("ADDITIONAL EXERCISE – TWO NEW CONFIGURATIONS")
print("=" * 70)

# -----
# 1. RESTORE REQUIRED VARIABLES
# -----

# These should already exist from the previous cells.
# Define them here only if they are missing.

if "SEED" not in globals():
    SEED = 42

if "EPOCHS" not in globals():
    EPOCHS = 10

if "NUM_CLASSES" not in globals():
    NUM_CLASSES = 37

# -----
# 2. VERIFY DATASET AND FOLDS
# -----

required_variables = [
    "train_paths",
    "train_labels",
    "val_paths",
    "val_labels",
    "test_paths",
    "test_labels",
    "folds"
]

missing = [
    variable
    for variable in required_variables
    if variable not in globals()
]

if missing:
    raise RuntimeError(
        "Required variables are missing: "
        + str(missing)
        + "\nRun the dataset restoration/setup cells first."
    )

print("\nDataset verified.")
print("Classes:", NUM_CLASSES)
print("Training samples:", len(train_paths))
print("Test samples:", len(test_paths))
print("Number of folds:", len(folds))

# -----
```

```
# 3. OUTPUT DIRECTORIES
# -----

ADDITIONAL_DIR = os.path.join(
    LAB5_DRIVE_ROOT,
    "additional_exercise"
)

ADDITIONAL_CV_DIR = os.path.join(
    ADDITIONAL_DIR,
    "cross_validation"
)

ADDITIONAL_RESULTS_DIR = os.path.join(
    ADDITIONAL_DIR,
    "results"
)

os.makedirs(
    ADDITIONAL_CV_DIR,
    exist_ok=True
)

os.makedirs(
    ADDITIONAL_RESULTS_DIR,
    exist_ok=True
)

ADDITIONAL_RESULTS_CSV = os.path.join(
    ADDITIONAL_RESULTS_DIR,
    "additional_exercise_cv_results.csv"
)

# -----
# 4. TWO NEW CONFIGURATIONS
# -----
#
# Selected configuration from the main experiment:
#
# LR      = 0.001
# Batch   = 32
# Dropout = 0.25
# Fine-tune = Partial
#
# These two configurations deliberately introduce NEW
# combinations of the requested hyperparameters.
#
# Configuration A:
# LR      = 0.0001
# Batch   = 64
# Dropout = 0.50
# FineTune = Partial
#
# Configuration B:
# LR      = 0.001
# Batch   = 16
# Dropout = 0.00
# FineTune = Frozen
# -----

additional_configs = {

    "Additional_Configuration_A": {

        "learning_rate": 0.0001,

        "batch_size": 64,

        "dropout": 0.50,

        "fine_tune": True

    },

    "Additional_Configuration_B": {

        "learning_rate": 0.001,

        "batch_size": 16,
```

```

        "dropout": 0.00,

        "fine_tune": False

    }

}

print("\nConfigurations:")
for name, config in additional_configs.items():

    print("\n", name)

    for key, value in config.items():

        print(
            f"    {key}: {value}"
        )

# -----
# 5. DATASET CREATION FUNCTION
# -----

def create_additional_fold_dataset(
    paths,
    labels,
    indices,
    batch_size,
    shuffle=False
):

    selected_paths = paths[indices]

    selected_labels = labels[indices]

    ds = tf.data.Dataset.from_tensor_slices(
        (
            selected_paths,
            selected_labels
        )
    )

    if shuffle:

        ds = ds.shuffle(
            len(selected_paths),
            seed=SEED,
            reshuffle_each_iteration=True
        )

    ds = ds.map(
        load_image,
        num_parallel_calls=tf.data.AUTOTUNE
    )

    ds = ds.batch(
        batch_size
    )

    ds = ds.prefetch(
        tf.data.AUTOTUNE
    )

    return ds

# -----
# 6. MODEL BUILDER
# -----

def build_additional_model(
    learning_rate,
    dropout,
    fine_tune

```

```

):

base_model = MobileNetV2(
    input_shape=(224, 224, 3),
    include_top=False,
    weights="imagenet"
)

# -----
# Fine-tuning strategy
# -----

if fine_tune:

    base_model.trainable = True

    # Unfreeze only the final 20 layers
    for layer in base_model.layers[:-20]:

        layer.trainable = False

    # Keep BatchNorm frozen
    for layer in base_model.layers:

        if isinstance(
            layer,
            layers.BatchNormalization
        ):

            layer.trainable = False

    else:

        base_model.trainable = False

# -----
# Classifier
# -----

inputs = keras.Input(
    shape=(224, 224, 3)
)

x = base_model(
    inputs,
    training=False
)

x = layers.GlobalAveragePooling2D()(x)

if dropout > 0:

    x = layers.Dropout(
        dropout
    )(x)

outputs = layers.Dense(
    NUM_CLASSES,
    activation="softmax"
)(x)

model = keras.Model(
    inputs,
    outputs
)

# -----
# Optimizer
# -----

if fine_tune:

    optimizer = keras.optimizers.RMSprop(

```

```

        learning_rate=learning_rate
    )

else:

    optimizer = keras.optimizers.Adam(
        learning_rate=learning_rate
    )

    model.compile(
        optimizer=optimizer,
        loss="sparse_categorical_crossentropy",
        metrics=["accuracy"]
    )

    return model

# -----
# 7. RESUME PREVIOUS RUN IF AVAILABLE
# -----

if os.path.exists(ADDITIONAL_RESULTS_CSV):

    additional_results_df = pd.read_csv(
        ADDITIONAL_RESULTS_CSV
    )

    additional_results = (
        additional_results_df
        .to_dict("records")
    )

    print(
        "\nExisting Additional Exercise results found."
    )

    print(
        "Completed folds:",
        len(additional_results)
    )

else:

    additional_results = []

    print(
        "\nNo previous Additional Exercise results found."
    )

def additional_fold_completed(
    configuration,
    fold
):

    return any(

        row["Configuration"] == configuration
        and int(row["Fold"]) == int(fold)

        for row in additional_results
    )

# -----
# 8. RUN 5-FOLD CROSS-VALIDATION
# -----

for config_name, config in additional_configs.items():

    print("\n" + "#" * 70)

    print(config_name)

    print("#" * 70)

    for fold_number, (
        train_idx,

```

```

        validation_idx
    ) in enumerate(
        folds,
        start=1
    ):

        # -----
        # Resume support
        # -----

        if additional_fold_completed(
            config_name,
            fold_number
        ):

            print(
                f"Skipping {config_name} "
                f"Fold {fold_number} "
                f"(already completed)"
            )

            continue

        print(
            f"\n{config_name} | "
            f"Fold {fold_number}/5"
        )

        # -----
        # Create fold datasets
        # -----

        fold_train_ds = (
            create_additional_fold_dataset(
                train_paths,
                train_labels,
                train_idx,
                batch_size=int(
                    config["batch_size"]
                ),
                shuffle=True
            )
        )

        fold_val_ds = (
            create_additional_fold_dataset(
                train_paths,
                train_labels,
                validation_idx,
                batch_size=int(
                    config["batch_size"]
                ),
                shuffle=False
            )
        )

        # -----
        # Build model
        # -----

        model = build_additional_model(

            learning_rate=float(
                config["learning_rate"]
            ),

            dropout=float(
                config["dropout"]
            ),

            fine_tune=bool(
                config["fine_tune"]
            )
        )

        # -----

```

```
# Train
# -----

start_time = time.time()

history = model.fit(

    fold_train_ds,

    validation_data=fold_val_ds,

    epochs=EPOCHS,

    verbose=1
)

training_time = (
    time.time() - start_time
)

# -----
# Extract results
# -----

h = history.history

best_val_accuracy = (
    max(h["val_accuracy"]) * 100
)

final_val_accuracy = (
    h["val_accuracy"][-1] * 100
)

final_train_accuracy = (
    h["accuracy"][-1] * 100
)

final_train_loss = (
    h["loss"][-1]
)

final_val_loss = (
    h["val_loss"][-1]
)

# -----
# Parameter count
# -----

trainable_parameters = int(
    np.sum([
        np.prod(v.shape)
        for v in model.trainable_weights
    ])
)

total_parameters = int(
    model.count_params()
)

# -----
# Store result
# -----

row = {

    "Configuration":
        config_name,

    "Fold":
```

```

        fold_number,

        "Learning Rate":
            config["learning_rate"],

        "Batch Size":
            config["batch_size"],

        "Dropout":
            config["dropout"],

        "Fine-Tuning":
            "Partial" if config["fine_tune"]
            else "Frozen",

        "Trainable Parameters":
            trainable_parameters,

        "Total Parameters":
            total_parameters,

        "Best Validation Accuracy (%)":
            best_val_accuracy,

        "Final Training Accuracy (%)":
            final_train_accuracy,

        "Final Validation Accuracy (%)":
            final_val_accuracy,

        "Final Training Loss":
            final_train_loss,

        "Final Validation Loss":
            final_val_loss,

        "Training Time (s)":
            training_time
    }

    additional_results.append(
        row
    )

    # -----
    # SAVE IMMEDIATELY
    # -----

    pd.DataFrame(
        additional_results
    ).to_csv(
        ADDITIONAL_RESULTS_CSV,
        index=False
    )

    # -----
    # Save fold model
    # -----

    config_dir = os.path.join(
        ADDITIONAL_CV_DIR,

        config_name.lower()
    )

    os.makedirs(
        config_dir,
        exist_ok=True
    )

    model.save(
        os.path.join(
            config_dir,

```

```

        f"fold_{fold_number}.keras"
    )
)

# -----
# Save history
# -----

history_path = os.path.join(

    config_dir,

    f"fold_{fold_number}_history.json"
)

with open(
    history_path,
    "w"
) as f:

    json.dump(

        {

            key: [
                float(x)
                for x in values
            ]

            for key, values
            in h.items()
        },

        f,

        indent=2
    )

print(
    f"\nFold {fold_number} complete."
)

print(
    f"Best validation accuracy: "
    f"{best_val_accuracy:.2f}%"
)

print(
    f"Training time: "
    f"{training_time / 60:.2f} minutes"
)

# -----
# Memory cleanup
# -----

del model
del fold_train_ds
del fold_val_ds
del history

tf.keras.backend.clear_session()

gc.collect()

# -----
# 9. FINAL ADDITIONAL EXERCISE RESULTS
# -----

additional_results_df = pd.DataFrame(
    additional_results
)

additional_results_df.to_csv(
    ADDITIONAL_RESULTS_CSV,
    index=False
)

```

```
print("\n" + "=" * 70)
print("ADDITIONAL EXERCISE CV COMPLETE")
print("=" * 70)
```

```
display(
    additional_results_df
)
```


=====

ADDITIONAL EXERCISE – TWO NEW CONFIGURATIONS

=====

Dataset verified.

Classes: 37

Training samples: 2944

Test samples: 3669

Number of folds: 5

Configurations:

Additional_Configuration_A

learning_rate: 0.0001

batch_size: 64

dropout: 0.5

fine_tune: True

Additional_Configuration_B

learning_rate: 0.001

batch_size: 16

dropout: 0.0

fine_tune: False

No previous Additional Exercise results found.

#####

Additional_Configuration_A

#####

Additional_Configuration_A | Fold 1/5

Epoch 1/10

37/37 66s 1s/step - accuracy: 0.4161 - loss: 2.2013 - val_accuracy: 0.8285 - val_loss: 0.9079

Epoch 2/10

37/37 12s 323ms/step - accuracy: 0.7584 - loss: 0.8680 - val_accuracy: 0.8472 - val_loss: 0.5540

Epoch 3/10

37/37 11s 307ms/step - accuracy: 0.8590 - loss: 0.4976 - val_accuracy: 0.8744 - val_loss: 0.4113

Epoch 4/10

37/37 20s 300ms/step - accuracy: 0.8955 - loss: 0.3438 - val_accuracy: 0.8642 - val_loss: 0.4140

Epoch 5/10

37/37 21s 313ms/step - accuracy: 0.9278 - loss: 0.2549 - val_accuracy: 0.8846 - val_loss: 0.3652

Epoch 6/10

37/37 10s 264ms/step - accuracy: 0.9401 - loss: 0.1914 - val_accuracy: 0.8778 - val_loss: 0.3729

Epoch 7/10

37/37 11s 292ms/step - accuracy: 0.9597 - loss: 0.1437 - val_accuracy: 0.8879 - val_loss: 0.3758

Epoch 8/10

37/37 11s 301ms/step - accuracy: 0.9639 - loss: 0.1211 - val_accuracy: 0.8947 - val_loss: 0.3349

Epoch 9/10

37/37 11s 307ms/step - accuracy: 0.9711 - loss: 0.0967 - val_accuracy: 0.8947 - val_loss: 0.3369

Epoch 10/10

37/37 12s 331ms/step - accuracy: 0.9817 - loss: 0.0776 - val_accuracy: 0.8930 - val_loss: 0.3471

Fold 1 complete.

Best validation accuracy: 89.47%

Training time: 3.11 minutes

Additional_Configuration_A | Fold 2/5

Epoch 1/10

37/37 36s 656ms/step - accuracy: 0.4123 - loss: 2.2108 - val_accuracy: 0.7878 - val_loss: 0.8999

Epoch 2/10

37/37 11s 302ms/step - accuracy: 0.7720 - loss: 0.8043 - val_accuracy: 0.8659 - val_loss: 0.5138

Epoch 3/10

37/37 11s 306ms/step - accuracy: 0.8518 - loss: 0.5021 - val_accuracy: 0.8761 - val_loss: 0.4126

Epoch 4/10

37/37 9s 249ms/step - accuracy: 0.8955 - loss: 0.3455 - val_accuracy: 0.8947 - val_loss: 0.3677

Epoch 5/10

37/37 12s 313ms/step - accuracy: 0.9253 - loss: 0.2504 - val_accuracy: 0.8930 - val_loss: 0.3431

Epoch 6/10

37/37 11s 305ms/step - accuracy: 0.9372 - loss: 0.1970 - val_accuracy: 0.8930 - val_loss: 0.3599

Epoch 7/10

37/37 11s 309ms/step - accuracy: 0.9580 - loss: 0.1407 - val_accuracy: 0.8964 - val_loss: 0.3310

Epoch 8/10

37/37 14s 373ms/step - accuracy: 0.9741 - loss: 0.1010 - val_accuracy: 0.9032 - val_loss: 0.3306

Epoch 9/10

37/37 10s 256ms/step - accuracy: 0.9732 - loss: 0.1000 - val_accuracy: 0.8981 - val_loss: 0.3547

Epoch 10/10

37/37 12s 303ms/step - accuracy: 0.9826 - loss: 0.0644 - val_accuracy: 0.8812 - val_loss: 0.4315

Fold 2 complete.

Best validation accuracy: 90.32%

Training time: 2.28 minutes

Additional_Configuration_A | Fold 3/5

Epoch 1/10

37/37 36s 672ms/step - accuracy: 0.3970 - loss: 2.2977 - val_accuracy: 0.8200 - val_loss: 0.8830

Epoch 2/10

37/37 9s 239ms/step - accuracy: 0.7520 - loss: 0.8926 - val_accuracy: 0.8812 - val_loss: 0.4772

Epoch 3/10

37/37 11s 302ms/step - accuracy: 0.8412 - loss: 0.5554 - val_accuracy: 0.8947 - val_loss: 0.3691

Epoch 4/10

```
37/37 ----- 12s 323ms/step - accuracy: 0.8879 - loss: 0.3888 - val_accuracy: 0.8862 - val_loss: 0.3645
Epoch 5/10
37/37 ----- 12s 324ms/step - accuracy: 0.9219 - loss: 0.2764 - val_accuracy: 0.8998 - val_loss: 0.2979
Epoch 6/10
37/37 ----- 21s 338ms/step - accuracy: 0.9316 - loss: 0.2141 - val_accuracy: 0.9049 - val_loss: 0.2835
Epoch 7/10
37/37 ----- 13s 357ms/step - accuracy: 0.9473 - loss: 0.1676 - val_accuracy: 0.9236 - val_loss: 0.2534
Epoch 8/10
37/37 ----- 20s 327ms/step - accuracy: 0.9601 - loss: 0.1345 - val_accuracy: 0.9168 - val_loss: 0.2759
Epoch 9/10
37/37 ----- 12s 326ms/step - accuracy: 0.9677 - loss: 0.1078 - val_accuracy: 0.9083 - val_loss: 0.3027
Epoch 10/10
37/37 ----- 19s 281ms/step - accuracy: 0.9779 - loss: 0.0771 - val_accuracy: 0.9083 - val_loss: 0.2727
```

Fold 3 complete.

Best validation accuracy: 92.36%

Training time: 2.90 minutes

Additional_Configuration_A | Fold 4/5

```
Epoch 1/10
37/37 ----- 37s 719ms/step - accuracy: 0.4323 - loss: 2.1614 - val_accuracy: 0.8557 - val_loss: 0.8289
Epoch 2/10
37/37 ----- 11s 300ms/step - accuracy: 0.7626 - loss: 0.8221 - val_accuracy: 0.8642 - val_loss: 0.5096
Epoch 3/10
37/37 ----- 10s 264ms/step - accuracy: 0.8454 - loss: 0.5275 - val_accuracy: 0.8795 - val_loss: 0.3825
Epoch 4/10
37/37 ----- 12s 311ms/step - accuracy: 0.8887 - loss: 0.3572 - val_accuracy: 0.8693 - val_loss: 0.3619
Epoch 5/10
37/37 ----- 13s 342ms/step - accuracy: 0.9295 - loss: 0.2517 - val_accuracy: 0.9032 - val_loss: 0.3010
Epoch 6/10
37/37 ----- 11s 310ms/step - accuracy: 0.9350 - loss: 0.2086 - val_accuracy: 0.9100 - val_loss: 0.2979
Epoch 7/10
37/37 ----- 10s 269ms/step - accuracy: 0.9537 - loss: 0.1571 - val_accuracy: 0.9049 - val_loss: 0.2901
Epoch 8/10
37/37 ----- 11s 298ms/step - accuracy: 0.9618 - loss: 0.1253 - val_accuracy: 0.8981 - val_loss: 0.2762
Epoch 9/10
37/37 ----- 11s 299ms/step - accuracy: 0.9724 - loss: 0.0900 - val_accuracy: 0.9049 - val_loss: 0.2951
Epoch 10/10
37/37 ----- 11s 299ms/step - accuracy: 0.9834 - loss: 0.0699 - val_accuracy: 0.9066 - val_loss: 0.2673
```

Fold 4 complete.

Best validation accuracy: 91.00%

Training time: 2.30 minutes

Additional_Configuration_A | Fold 5/5

```
Epoch 1/10
37/37 ----- 55s 1s/step - accuracy: 0.4342 - loss: 2.1210 - val_accuracy: 0.8095 - val_loss: 0.8658
Epoch 2/10
37/37 ----- 13s 357ms/step - accuracy: 0.7593 - loss: 0.8270 - val_accuracy: 0.8452 - val_loss: 0.4938
Epoch 3/10
37/37 ----- 19s 323ms/step - accuracy: 0.8519 - loss: 0.4974 - val_accuracy: 0.8844 - val_loss: 0.3826
Epoch 4/10
37/37 ----- 12s 327ms/step - accuracy: 0.8994 - loss: 0.3402 - val_accuracy: 0.8571 - val_loss: 0.4122
Epoch 5/10
37/37 ----- 15s 406ms/step - accuracy: 0.9223 - loss: 0.2485 - val_accuracy: 0.9082 - val_loss: 0.3058
Epoch 6/10
37/37 ----- 13s 362ms/step - accuracy: 0.9440 - loss: 0.1977 - val_accuracy: 0.8861 - val_loss: 0.3421
Epoch 7/10
37/37 ----- 11s 286ms/step - accuracy: 0.9576 - loss: 0.1542 - val_accuracy: 0.9031 - val_loss: 0.2965
Epoch 8/10
37/37 ----- 12s 320ms/step - accuracy: 0.9605 - loss: 0.1217 - val_accuracy: 0.9116 - val_loss: 0.2900
Epoch 9/10
37/37 ----- 20s 322ms/step - accuracy: 0.9733 - loss: 0.0913 - val_accuracy: 0.9116 - val_loss: 0.2768
Epoch 10/10
37/37 ----- 10s 266ms/step - accuracy: 0.9826 - loss: 0.0659 - val_accuracy: 0.9065 - val_loss: 0.2991
```

Fold 5 complete.

Best validation accuracy: 91.16%

Training time: 3.02 minutes

#####

Additional_Configuration_B

#####

Additional_Configuration_B | Fold 1/5

```
Epoch 1/10
148/148 ----- 53s 238ms/step - accuracy: 0.7282 - loss: 1.1334 - val_accuracy: 0.8710 - val_loss: 0.4553
Epoch 2/10
148/148 ----- 12s 84ms/step - accuracy: 0.9397 - loss: 0.2588 - val_accuracy: 0.8862 - val_loss: 0.3583
Epoch 3/10
148/148 ----- 12s 84ms/step - accuracy: 0.9686 - loss: 0.1487 - val_accuracy: 0.8964 - val_loss: 0.3345
Epoch 4/10
148/148 ----- 19s 70ms/step - accuracy: 0.9868 - loss: 0.0974 - val_accuracy: 0.9083 - val_loss: 0.3036
Epoch 5/10
148/148 ----- 12s 83ms/step - accuracy: 0.9962 - loss: 0.0659 - val_accuracy: 0.8964 - val_loss: 0.2997
Epoch 6/10
148/148 ----- 21s 87ms/step - accuracy: 0.9983 - loss: 0.0483 - val_accuracy: 0.9100 - val_loss: 0.2883
Epoch 7/10
148/148 ----- 12s 78ms/step - accuracy: 0.9992 - loss: 0.0376 - val_accuracy: 0.9049 - val_loss: 0.2882
Epoch 8/10
148/148 ----- 21s 81ms/step - accuracy: 0.9996 - loss: 0.0302 - val_accuracy: 0.9032 - val_loss: 0.2822
```

```
Epoch 9/10
148/148 ----- 20s 80ms/step - accuracy: 0.9996 - loss: 0.0248 - val_accuracy: 0.9100 - val_loss: 0.2843
Epoch 10/10
148/148 ----- 12s 82ms/step - accuracy: 1.0000 - loss: 0.0204 - val_accuracy: 0.9083 - val_loss: 0.2814

Fold 1 complete.
Best validation accuracy: 91.00%
Training time: 3.26 minutes

Additional_Configuration_B | Fold 2/5
Epoch 1/10
148/148 ----- 34s 171ms/step - accuracy: 0.7130 - loss: 1.1608 - val_accuracy: 0.8625 - val_loss: 0.4848
Epoch 2/10
148/148 ----- 12s 82ms/step - accuracy: 0.9312 - loss: 0.2751 - val_accuracy: 0.9015 - val_loss: 0.3689
Epoch 3/10
148/148 ----- 20s 81ms/step - accuracy: 0.9686 - loss: 0.1571 - val_accuracy: 0.8846 - val_loss: 0.3497
Epoch 4/10
148/148 ----- 13s 91ms/step - accuracy: 0.9890 - loss: 0.0997 - val_accuracy: 0.8947 - val_loss: 0.3179
Epoch 5/10
148/148 ----- 14s 91ms/step - accuracy: 0.9975 - loss: 0.0673 - val_accuracy: 0.8947 - val_loss: 0.3187
Epoch 6/10
148/148 ----- 12s 78ms/step - accuracy: 0.9987 - loss: 0.0514 - val_accuracy: 0.8964 - val_loss: 0.3021
Epoch 7/10
148/148 ----- 11s 74ms/step - accuracy: 0.9992 - loss: 0.0393 - val_accuracy: 0.9049 - val_loss: 0.2955
Epoch 8/10
148/148 ----- 12s 83ms/step - accuracy: 0.9996 - loss: 0.0312 - val_accuracy: 0.8981 - val_loss: 0.2993
Epoch 9/10
148/148 ----- 13s 89ms/step - accuracy: 1.0000 - loss: 0.0248 - val_accuracy: 0.9049 - val_loss: 0.2898
Epoch 10/10
148/148 ----- 12s 83ms/step - accuracy: 1.0000 - loss: 0.0205 - val_accuracy: 0.8998 - val_loss: 0.2983

Fold 2 complete.
Best validation accuracy: 90.49%
Training time: 2.57 minutes

Additional_Configuration_B | Fold 3/5
Epoch 1/10
148/148 ----- 32s 159ms/step - accuracy: 0.7134 - loss: 1.2059 - val_accuracy: 0.8964 - val_loss: 0.4331
Epoch 2/10
148/148 ----- 12s 81ms/step - accuracy: 0.9384 - loss: 0.2649 - val_accuracy: 0.9032 - val_loss: 0.3334
Epoch 3/10
148/148 ----- 12s 83ms/step - accuracy: 0.9711 - loss: 0.1566 - val_accuracy: 0.9117 - val_loss: 0.2960
Epoch 4/10
148/148 ----- 11s 75ms/step - accuracy: 0.9860 - loss: 0.0999 - val_accuracy: 0.9185 - val_loss: 0.2756
Epoch 5/10
148/148 ----- 12s 79ms/step - accuracy: 0.9953 - loss: 0.0697 - val_accuracy: 0.9151 - val_loss: 0.2673
Epoch 6/10
148/148 ----- 13s 86ms/step - accuracy: 0.9970 - loss: 0.0518 - val_accuracy: 0.9151 - val_loss: 0.2532
Epoch 7/10
148/148 ----- 20s 86ms/step - accuracy: 0.9992 - loss: 0.0394 - val_accuracy: 0.9168 - val_loss: 0.2498
Epoch 8/10
148/148 ----- 11s 72ms/step - accuracy: 0.9996 - loss: 0.0319 - val_accuracy: 0.9185 - val_loss: 0.2447
Epoch 9/10
148/148 ----- 12s 81ms/step - accuracy: 0.9996 - loss: 0.0258 - val_accuracy: 0.9253 - val_loss: 0.2380
Epoch 10/10
148/148 ----- 20s 80ms/step - accuracy: 0.9996 - loss: 0.0211 - val_accuracy: 0.9202 - val_loss: 0.2425

Fold 3 complete.
Best validation accuracy: 92.53%
Training time: 2.59 minutes

Additional_Configuration_B | Fold 4/5
Epoch 1/10
148/148 ----- 35s 183ms/step - accuracy: 0.7240 - loss: 1.1256 - val_accuracy: 0.8913 - val_loss: 0.4205
Epoch 2/10
148/148 ----- 13s 89ms/step - accuracy: 0.9359 - loss: 0.2664 - val_accuracy: 0.9032 - val_loss: 0.3255
Epoch 3/10
148/148 ----- 19s 80ms/step - accuracy: 0.9711 - loss: 0.1536 - val_accuracy: 0.9151 - val_loss: 0.2685
Epoch 4/10
148/148 ----- 12s 84ms/step - accuracy: 0.9894 - loss: 0.0972 - val_accuracy: 0.9202 - val_loss: 0.2605
Epoch 5/10
148/148 ----- 13s 91ms/step - accuracy: 0.9962 - loss: 0.0683 - val_accuracy: 0.9236 - val_loss: 0.2397
Epoch 6/10
148/148 ----- 13s 85ms/step - accuracy: 0.9983 - loss: 0.0491 - val_accuracy: 0.9270 - val_loss: 0.2401
Epoch 7/10
148/148 ----- 13s 85ms/step - accuracy: 0.9996 - loss: 0.0381 - val_accuracy: 0.9219 - val_loss: 0.2257
Epoch 8/10
148/148 ----- 12s 81ms/step - accuracy: 0.9992 - loss: 0.0312 - val_accuracy: 0.9270 - val_loss: 0.2261
Epoch 9/10
148/148 ----- 12s 80ms/step - accuracy: 1.0000 - loss: 0.0250 - val_accuracy: 0.9287 - val_loss: 0.2198
Epoch 10/10
148/148 ----- 13s 88ms/step - accuracy: 1.0000 - loss: 0.0213 - val_accuracy: 0.9338 - val_loss: 0.2139

Fold 4 complete.
Best validation accuracy: 93.38%
Training time: 2.59 minutes

Additional_Configuration_B | Fold 5/5
Epoch 1/10
148/148 ----- 42s 221ms/step - accuracy: 0.7177 - loss: 1.1790 - val_accuracy: 0.8827 - val_loss: 0.4536
```

```

# =====
# CELL R32A – MOUNT DRIVE & VERIFY SAVED FINAL RESULT
# =====

import os
import pandas as pd
from google.colab import drive

print("=" * 70)
print("MOUNTING GOOGLE DRIVE")
print("=" * 70)

# Mount Google Drive
drive.mount("/content/drive")

print("\n✓ Google Drive mounted.")

# -----
# Path to saved final-model results
# -----

RESULTS_PATH = (
    "/content/drive/MyDrive/"
    "Lab5/Lab5/results/final_model_results.csv"
)

# -----
# Verify file exists
# -----

print("\n" + "=" * 70)
print("CHECKING SAVED FINAL MODEL RESULT")
print("=" * 70)

if not os.path.exists(RESULTS_PATH):

    raise FileNotFoundError(
        f"\nFinal-model results file not found:\n"
        f"{RESULTS_PATH}"
    )

print("✓ Results file found:")
print(RESULTS_PATH)

# -----
# Load saved results
# -----

final_results = pd.read_csv(
    RESULTS_PATH
)

# -----
# Display complete result
# -----

print("\n" + "=" * 70)
print("SAVED CORRECTED FINAL MODEL RESULT")
print("=" * 70)

display(final_results)

# -----
# Print key values
# -----

row = final_results.iloc[0]

print("\n" + "=" * 70)
print("FINAL RESULT TO USE")
print("=" * 70)

print(
    f"Learning Rate      : {row['Learning Rate']}"
)

print(

```

```
        f"Batch Size          : {row['Batch Size']}"
    )

    print(
        f"Dropout              : {row['Dropout']}"
    )

    print(
        f"Best Validation Acc. : "
        f"{row['Best Validation Accuracy (%)']:.2f}%"
    )

    print(
        f"Test Loss              : "
        f"{row['Test Loss']:.6f}"
    )

    print(
        f"Test Accuracy          : "
        f"{row['Test Accuracy (%)']:.2f}%"
    )

    print(
        f"Training Time          : "
        f"{row['Training Time (s)']:.2f} s"
    )

    # -----
    # Final confirmation
    # -----

    print("\n" + "=" * 70)
    print("VERIFICATION COMPLETE")
    print("=" * 70)

    if abs(
        float(row["Test Accuracy (%)"]) - 89.72472
    ) < 0.1:

        print("✓ Correct saved final-model result confirmed.")
        print("✓ Final Test Accuracy = 89.72%")
        print("✓ No model evaluation was performed.")
        print("✓ No retraining was performed.")

    else:

        print(
            "⚠ WARNING: Saved accuracy differs from expected 89.72%."
        )
```

```
=====
MOUNTING GOOGLE DRIVE
=====
Mounted at /content/drive

✓ Google Drive mounted.

=====
CHECKING SAVED FINAL MODEL RESULT
=====
✓ Results file found:
/content/drive/MyDrive/Lab5/Lab5/results/final_model_results.csv

=====
SAVED CORRECTED FINAL MODEL RESULT
=====
```

	Model	Learning Rate	Batch Size	Dropout	Trainable Parameters	Total Parameters	Best Validation Accuracy (%)	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	Test Loss
0	Final Selected Model	0.001	32	0.25	47397	2305381	91.168481	99.252719	90.625	0.059887	0.267998	0.316205

```
=====
FINAL RESULT TO USE
=====
Learning Rate      : 0.001
Batch Size         : 32
Dropout            : 0.25
Best Validation Acc.: 91.17%
Test Loss          : 0.316205
Test Accuracy      : 89.72%
Training Time      : 182.59 s

=====
VERIFICATION COMPLETE
=====
✓ Correct saved final-model result confirmed.
✓ Final Test Accuracy = 89.72%
✓ No model evaluation was performed.
✓ No retraining was performed.
```

```
# =====
# CELL - CORRECTLY RECOVER ADDITIONAL EXERCISE RESULTS
#
# IMPORTANT:
# - NO TRAINING
# - NO RETRAINING
# - NO MODEL EVALUATION
# - Reads existing JSON histories
# - Inspects existing .keras files
# - Groups Fold 1..5 under their ACTUAL configuration
# - Generates report tables + plots
# - Saves EPS + PNG at 600 DPI
# - Saves to BOTH Runtime and Google Drive
# =====

import os
import re
import json
import glob
import zipfile
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from IPython.display import display

# =====
# 1. MOUNT DRIVE
# =====

from google.colab import drive

if not os.path.exists("/content/drive/MyDrive"):
    print("=" * 70)
    print("MOUNTING GOOGLE DRIVE")
    print("=" * 70)

    drive.mount("/content/drive")
```

```

else:
    print("✓ Google Drive already mounted.")

# =====
# 2. PATHS
# =====

LAB5_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

CV_ROOT = os.path.join(
    LAB5_ROOT,
    "checkpoints",
    "cross_validation"
)

RESULTS_ROOT = os.path.join(
    LAB5_ROOT,
    "results"
)

DRIVE_OUTPUT = os.path.join(
    LAB5_ROOT,
    "report_materials",
    "additional_exercise"
)

RUNTIME_OUTPUT = os.path.join(
    "/content",
    "Lab5_report_materials",
    "additional_exercise"
)

os.makedirs(
    DRIVE_OUTPUT,
    exist_ok=True
)

os.makedirs(
    RUNTIME_OUTPUT,
    exist_ok=True
)

# =====
# 3. BASIC CHECK
# =====

print("\n" + "=" * 70)
print("ADDITIONAL EXERCISE RECOVERY")
print("=" * 70)

print("\nCross-validation root:")
print(CV_ROOT)

if not os.path.exists(CV_ROOT):
    raise FileNotFoundError(
        f"Cross-validation directory not found:\n{CV_ROOT}"
    )

# =====
# 4. FIND ALL JSON + KERAS FILES
# =====

all_json = sorted(
    glob.glob(
        os.path.join(
            CV_ROOT,
            "**",
            "*.json"
        ),
        recursive=True
    )
)

all_keras = sorted(
    glob.glob(
        os.path.join(
            CV_ROOT,
            "**",
            "*.keras"
        )
    )
)

```

```

    ),
    recursive=True
)
)

print("\n" + "=" * 70)
print("FILES FOUND")
print("=" * 70)

print(
    f"\nJSON files : {len(all_json)}"
)

for p in all_json:
    print(
        "  JSON :",
        os.path.relpath(
            p,
            CV_ROOT
        )
    )

print(
    f"\nKeras files: {len(all_keras)}"
)

for p in all_keras:
    print(
        "  KERAS:",
        os.path.relpath(
            p,
            CV_ROOT
        )
    )

# =====
# 5. IDENTIFY FOLD NUMBER
# =====

def get_fold_number(path):

    match = re.search(
        r"fold[_-\s]?(\d+)",
        path.lower()
    )

    if match:
        return int(
            match.group(1)
        )

    return None

# =====
# 6. IDENTIFY CONFIGURATION DIRECTORY
# =====

def get_configuration_directory(path):

    """
    The configuration is determined from the DIRECTORY
    containing the fold directory.

    Example:

    additional_configurations/
    configuration_1/
        fold_1_history.json
        fold_2_history.json

    becomes:

    Configuration 1
    Fold 1
    Fold 2

    NOT:

    Fold 1

```

```

    Fold 2
    Additional Exercise
    """

    path = os.path.normpath(path)

    parts = path.split(
        os.sep
    )

    fold_index = None

    for i, part in enumerate(parts):

        if re.search(
            r"fold[_-\s]?[d+]",
            part.lower()
        ):

            fold_index = i
            break

    if fold_index is not None and fold_index > 0:

        config_part = parts[
            fold_index - 1
        ]

        return config_part

    # If fold is encoded in filename rather than directory,
    # search backwards for a configuration-like directory.

    for part in reversed(parts[:-1]):

        lower = part.lower()

        if (
            "configuration" in lower
            or
            "config_" in lower
            or
            lower.startswith("config")
        ):

            return part

    return "Unknown Configuration"

def pretty_configuration_name(name):

    lower = name.lower()

    match = re.search(
        r"(?:configuration|config)[_-\s]?[d+]",
        lower
    )

    if match:

        return (
            "New Configuration "
            + match.group(1)
        )

    return name.replace(
        "-",
        " "
    ).replace(
        "_",
        " "
    ).title()

# =====
# 7. LOAD HISTORY FROM JSON
# =====

def extract_history(obj):

    """

```

```

Finds the Keras history dictionary recursively.
"""

if isinstance(obj, dict):

    keys = set(
        str(k)
        for k in obj.keys()
    )

    history_keys = {
        "accuracy",
        "loss",
        "val_accuracy",
        "val_loss"
    }

    if keys.intersection(
        history_keys
    ):

        return obj

    # Common wrapper
    if "history" in obj:

        result = extract_history(
            obj["history"]
        )

        if result is not None:
            return result

    for value in obj.values():

        result = extract_history(
            value
        )

        if result is not None:
            return result

elif isinstance(obj, list):

    for item in obj:

        result = extract_history(
            item
        )

        if result is not None:
            return result

    return None

# =====
# 8. LOAD ALL HISTORY FILES
# =====

records = []

print("\n" + "=" * 70)
print("READING SAVED FOLD HISTORIES")
print("=" * 70)

for path in all_json:

    fold = get_fold_number(
        path
    )

    # Ignore JSONs that aren't fold histories
    if fold is None:
        continue

    config_dir = get_configuration_directory(
        path
    )

    config_name = pretty_configuration_name(
        config_dir

```

```

    )

    try:

        with open(
            path,
            "r"
        ) as f:

            raw = json.load(f)

            history = extract_history(
                raw
            )

            if history is None:

                print(
                    f"⚠ No training history found: "
                    f"{os.path.basename(path)}"
                )

                continue

            records.append({

                "configuration":
                    config_name,

                "configuration_directory":
                    config_dir,

                "fold":
                    fold,

                "path":
                    path,

                "history":
                    history

            })

            print(
                f"✓ {config_name} | Fold {fold} | "
                f"{os.path.basename(path)}"
            )

        except Exception as e:

            print(
                f"⚠ Could not read {path}: {e}"
            )

    if not records:

        raise RuntimeError(
            "No fold histories could be recovered."
        )

# =====
# 9. SORT PROPERLY
# =====

records = sorted(
    records,
    key=lambda x: (
        x["configuration"],
        x["fold"]
    )
)

# =====
# 10. SHOW CONFIGURATION / FOLD STRUCTURE
# =====

print("\n" + "=" * 70)
print("RECOVERED CONFIGURATION STRUCTURE")
print("=" * 70)

```

```

structure = (
    pd.DataFrame([
        {
            "Configuration":
                r["configuration"],
            "Configuration Directory":
                r["configuration_directory"],
            "Fold":
                r["fold"],
            "History File":
                os.path.basename(
                    r["path"]
                )
        }
        for r in records
    ])
)

display(
    structure
)

structure.to_csv(
    os.path.join(
        DRIVE_OUTPUT,
        "additional_exercise_file_structure.csv"
    ),
    index=False
)

# =====
# 11. CONVERT HISTORIES TO FOLD SUMMARY
# =====

fold_rows = []

for r in records:

    h = r["history"]

    train_acc = h.get(
        "accuracy",
        []
    )

    val_acc = h.get(
        "val_accuracy",
        []
    )

    train_loss = h.get(
        "loss",
        []
    )

    val_loss = h.get(
        "val_loss",
        []
    )

    # Convert to numpy arrays safely
    train_acc = np.asarray(
        train_acc,
        dtype=float
    )

    val_acc = np.asarray(
        val_acc,
        dtype=float
    )

    train_loss = np.asarray(
        train_loss,
        dtype=float
    )

    val_loss = np.asarray(
        val_loss,
        dtype=float
    )

```

```

# Best validation accuracy
if len(val_acc) > 0:

    best_val_acc = (
        np.nanmax(val_acc) * 100
    )

    best_epoch = (
        np.nanargmax(val_acc) + 1
    )

    final_val_acc = (
        val_acc[-1] * 100
    )

else:

    best_val_acc = np.nan
    best_epoch = np.nan
    final_val_acc = np.nan

# Final training accuracy
if len(train_acc) > 0:

    final_train_acc = (
        train_acc[-1] * 100
    )

else:

    final_train_acc = np.nan

# Final losses
final_train_loss = (
    train_loss[-1]
    if len(train_loss) > 0
    else np.nan
)

final_val_loss = (
    val_loss[-1]
    if len(val_loss) > 0
    else np.nan
)

fold_rows.append({

    "Configuration":
        r["configuration"],

    "Fold":
        r["fold"],

    "Best Validation Accuracy (%)":
        best_val_acc,

    "Best Epoch":
        best_epoch,

    "Final Training Accuracy (%)":
        final_train_acc,

    "Final Validation Accuracy (%)":
        final_val_acc,

    "Final Training Loss":
        final_train_loss,

    "Final Validation Loss":
        final_val_loss

})

fold_df = pd.DataFrame(
    fold_rows
)

fold_df = fold_df.sort_values(
    [
        "Configuration",

```

```

        "Fold"
    ]
).reset_index(
    drop=True
)

# =====
# 12. FOLD-WISE RESULTS
# =====

print("\n" + "=" * 70)
print("FOLD-WISE RESULTS")
print("=" * 70)

display(
    fold_df
)

fold_df.to_csv(
    os.path.join(
        DRIVE_OUTPUT,
        "additional_exercise_fold_results.csv"
    ),
    index=False
)

# =====
# 13. VERIFY NUMBER OF FOLDS
# =====

print("\n" + "=" * 70)
print("FOLD COUNT VERIFICATION")
print("=" * 70)

fold_counts = (
    fold_df
    .groupby("Configuration")
    ["Fold"]
    .nunique()
    .reset_index(
        name="Number of Folds"
    )
)

display(
    fold_counts
)

for _, row in fold_counts.iterrows():

    if row["Number of Folds"] != 5:

        print(
            f"⚠ {row['Configuration']} has "
            f"{row['Number of Folds']} fold(s), "
            f"not 5."
        )

    else:

        print(
            f"✓ {row['Configuration']}: "
            f"5-fold CV confirmed."
        )

# =====
# 14. CALCULATE MEAN + SD
# =====

cv_summary = (
    fold_df
    .groupby(
        "Configuration"
    )
    .agg(

        Mean_CV_Accuracy=(
            "Best Validation Accuracy (%)",
            "mean"

```

```

    ),

    SD_CV_Accuracy=(
        "Best Validation Accuracy (%)",
        "std"
    ),

    Mean_Final_Validation=(
        "Final Validation Accuracy (%)",
        "mean"
    ),

    SD_Final_Validation=(
        "Final Validation Accuracy (%)",
        "std"
    )

)

.reset_index()

)

cv_summary = cv_summary.rename(
    columns={

        "Mean_CV_Accuracy":
            "Mean CV Accuracy (%)",

        "SD_CV_Accuracy":
            "CV Standard Deviation (%)",

        "Mean_Final_Validation":
            "Mean Final Validation Accuracy (%)",

        "SD_Final_Validation":
            "Final Validation SD (%)"

    }

)

# =====
# 15. CV SUMMARY
# =====

print("\n" + "=" * 70)
print("CORRECT ADDITIONAL-EXERCISE CV SUMMARY")
print("=" * 70)

display(
    cv_summary
)

cv_summary.to_csv(
    os.path.join(
        DRIVE_OUTPUT,
        "ADDITIONAL_EXERCISE_CV_SUMMARY.csv"
    ),
    index=False
)

# =====
# 16. INSPECT KERAS FILES
# =====

print("\n" + "=" * 70)
print("INSPECTING SAVED .KERAS FILES")
print("=" * 70)

keras_rows = []

for path in all_keras:

    fold = get_fold_number(
        path
    )

    config_dir = get_configuration_directory(
        path

```

```

)

config_name = pretty_configuration_name(
    config_dir
)

keras_info = {

    "Configuration":
        config_name,

    "Fold":
        fold,

    "Keras File":
        os.path.basename(path),

    "Path":
        path,

    "File Size (MB)":
        os.path.getsize(path)
        / (1024 ** 2)

}

# -----
# Inspect .keras archive without loading/evaluating model
# -----

try:

    with zipfile.ZipFile(
        path,
        "r"
    ) as z:

        names = z.namelist()

        keras_info[
            "Contains config.json"
        ] = "config.json" in names

        keras_info[
            "Contains metadata.json"
        ] = "metadata.json" in names

        # Read metadata if present
        if "metadata.json" in names:

            try:

                metadata = json.loads(
                    z.read(
                        "metadata.json"
                    ).decode(
                        "utf-8"
                    )
                )

                keras_info[
                    "Keras Metadata"
                ] = str(
                    metadata
                )

            except Exception:

                keras_info[
                    "Keras Metadata"
                ] = "Unreadable"

        else:

            keras_info[
                "Keras Metadata"
            ] = ""

except Exception as e:

    keras_info[
        "Contains config.json"

```

```

        ] = False

        keras_info[
            "Contains metadata.json"
        ] = False

        keras_info[
            "Keras Metadata"
        ] = str(e)

        keras_rows.append(
            keras_info
        )

keras_df = pd.DataFrame(
    keras_rows
)

if not keras_df.empty:

    display(
        keras_df
    )

    keras_df.to_csv(
        os.path.join(
            DRIVE_OUTPUT,
            "additional_exercise_keras_inventory.csv"
        ),
        index=False
    )

else:

    print(
        "No .keras files found."
    )

# =====
# 17. PLOT 1 – FOLD-WISE ACCURACY
# =====

plt.figure(
    figsize=(9, 5)
)

for configuration, group in (
    fold_df
    .groupby("Configuration")
):

    plt.plot(
        group["Fold"],
        group[
            "Best Validation Accuracy (%)"
        ],
        marker="o",
        linewidth=2,
        label=configuration
    )

plt.xlabel(
    "Fold"
)

plt.ylabel(
    "Best Validation Accuracy (%)"
)

plt.title(
    "Additional Exercise – 5-Fold Cross-Validation"
)

plt.xticks(
    [1, 2, 3, 4, 5]
)

plt.grid(
    True,

```

```

        alpha=0.3
    )

plt.legend()

plt.tight_layout()

def save_figure(name):

    # Runtime
    plt.savefig(
        os.path.join(
            RUNTIME_OUTPUT,
            name + ".eps"
        ),
        format="eps",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        os.path.join(
            RUNTIME_OUTPUT,
            name + ".png"
        ),
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

    # Drive
    plt.savefig(
        os.path.join(
            DRIVE_OUTPUT,
            name + ".eps"
        ),
        format="eps",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        os.path.join(
            DRIVE_OUTPUT,
            name + ".png"
        ),
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

    print(
        f"✓ Saved {name}.eps + {name}.png"
    )

save_figure(
    "Additional_Exercise_5Fold_Accuracy"
)

plt.show()
plt.close()

# =====
# 18. PLOT 2 – MEAN ± SD
# =====

plt.figure(
    figsize=(9, 5)
)

x = np.arange(
    len(cv_summary)
)

plt.bar(
    x,
    cv_summary[
        "Mean CV Accuracy (%)"
    ],

```

```

        yerr=cv_summary[
            "CV Standard Deviation (%)"
        ],
        capsize=5
    )

plt.xticks(
    x,
    cv_summary[
        "Configuration"
    ]
)

plt.ylabel(
    "Mean CV Accuracy (%)"
)

plt.title(
    "Additional Exercise – Mean CV Accuracy  $\pm$  SD"
)

plt.grid(
    axis="y",
    alpha=0.3
)

plt.tight_layout()

save_figure(
    "Additional_Exercise_Mean_SD"
)

plt.show()
plt.close()

# =====
# 19. PLOT 3 – TRAINING / VALIDATION ACCURACY CURVES
# =====

for r in records:

    h = r["history"]

    train_acc = h.get(
        "accuracy",
        []
    )

    val_acc = h.get(
        "val_accuracy",
        []
    )

    if (
        len(train_acc) == 0
        and
        len(val_acc) == 0
    ):
        continue

    epochs = np.arange(
        1,
        max(
            len(train_acc),
            len(val_acc)
        ) + 1
    )

    plt.figure(
        figsize=(8, 5)
    )

    if len(train_acc) > 0:

        plt.plot(
            np.arange(
                1,
                len(train_acc) + 1
            ),
            np.asarray(
                train_acc
            )

```

```

        ) * 100,
        marker="o",
        label="Training Accuracy"
    )

if len(val_acc) > 0:

    plt.plot(
        np.arange(
            1,
            len(val_acc) + 1
        ),
        np.asarray(
            val_acc
        ) * 100,
        marker="o",
        label="Validation Accuracy"
    )

plt.xlabel(
    "Epoch"
)

plt.ylabel(
    "Accuracy (%)"
)

plt.title(
    f"{r['configuration']} - Fold {r['fold']} Accuracy"
)

plt.grid(
    True,
    alpha=0.3
)

plt.legend()

plt.tight_layout()

safe_config = re.sub(
    r"^[A-Za-z0-9]+",
    "_",
    r["configuration"]
).strip("_")

filename = (
    f"Additional_{safe_config}"
    f"_Fold_{r['fold']}_Accuracy"
)

save_figure(
    filename
)

plt.show()
plt.close()

# =====
# 20. PLOT 4 – TRAINING / VALIDATION LOSS CURVES
# =====

for r in records:

    h = r["history"]

    train_loss = h.get(
        "loss",
        []
    )

    val_loss = h.get(
        "val_loss",
        []
    )

    if (
        len(train_loss) == 0
        and
        len(val_loss) == 0
    ):

```

```

        continue

plt.figure(
    figsize=(8, 5)
)

if len(train_loss) > 0:

    plt.plot(
        np.arange(
            1,
            len(train_loss) + 1
        ),
        train_loss,
        marker="o",
        label="Training Loss"
    )

if len(val_loss) > 0:

    plt.plot(
        np.arange(
            1,
            len(val_loss) + 1
        ),
        val_loss,
        marker="o",
        label="Validation Loss"
    )

plt.xlabel(
    "Epoch"
)

plt.ylabel(
    "Loss"
)

plt.title(
    f"{r['configuration']} - Fold {r['fold']} Loss"
)

plt.grid(
    True,
    alpha=0.3
)

plt.legend()

plt.tight_layout()

safe_config = re.sub(
    r"^[^A-Za-z0-9]+",
    "_",
    r["configuration"]
).strip("_")

filename = (
    f"Additional_{safe_config}"
    f"_Fold_{r['fold']}_Loss"
)

save_figure(
    filename
)

plt.show()
plt.close()

# =====
# 21. FINAL REPORT TABLE
# =====

print("\n" + "=" * 70)
print("FINAL ADDITIONAL-EXERCISE REPORT TABLE")
print("=" * 70)

final_additional_table = cv_summary.copy()

display(
    final_additional_table

```

```

    )

    final_additional_table.to_csv(
        os.path.join(
            DRIVE_OUTPUT,
            "ADDITIONAL_EXERCISE_FINAL_REPORT_TABLE.csv"
        ),
        index=False
    )

# =====
# 22. FILE INVENTORY
# =====

print("\n" + "=" * 70)
print("GENERATED REPORT MATERIAL")
print("=" * 70)

drive_files = sorted(
    glob.glob(
        os.path.join(
            DRIVE_OUTPUT,
            "*"
        )
    )
)

print(
    f"\nDrive output directory:"
)

print(
    DRIVE_OUTPUT
)

print(
    f"\nFiles generated: {len(drive_files)}"
)

for p in drive_files:

    print(
        "  ✓",
        os.path.basename(p)
    )

# =====
# 23. FINAL STATUS
# =====

print("\n" + "=" * 70)
print("ADDITIONAL EXERCISE RECOVERY COMPLETE")
print("=" * 70)

print(
    "\n✓ Configuration grouping corrected"
)

print(
    "✓ Fold numbers retained under their configuration"
)

print(
    "✓ Existing JSON histories used"
)

print(
    "✓ Existing .keras files inspected"
)

print(
    "✓ 5-fold structure verified where available"
)

print(
    "✓ CV mean and standard deviation calculated"
)

print(
    "✓ Accuracy plots generated"

```

```
)

print(
    "✓ Loss plots generated"
)

print(
    "✓ EPS saved at 600 DPI"
)

print(
    "✓ PNG saved at 600 DPI"
)

print(
    "✓ Files saved to Runtime and Google Drive"
)

print(
    "\nNO TRAINING WAS PERFORMED."
)

print(
    "NO RETRAINING WAS PERFORMED."
)

print(
    "\nYour verified corrected final model remains:"
)

print(
    "Final Test Accuracy = 89.72%"
)

print(
    "The old 2.86% result is NOT used by this cell."
)
```


✓ Google Drive already mounted.

=====

ADDITIONAL EXERCISE RECOVERY

=====

Cross-validation root:

/content/drive/MyDrive/Lab5/Lab5/checkpoints/cross_validation

=====

FILES FOUND

=====

JSON files : 10

JSON : baseline/fold_1/history.json
 JSON : baseline/fold_2/history.json
 JSON : baseline/fold_3/history.json
 JSON : baseline/fold_4/history.json
 JSON : baseline/fold_5/history.json
 JSON : fine-tuned/fold_1/history.json
 JSON : fine-tuned/fold_2/history.json
 JSON : fine-tuned/fold_3/history.json
 JSON : fine-tuned/fold_4/history.json
 JSON : fine-tuned/fold_5/history.json

Keras files: 10

KERAS: baseline/fold_1/model.keras
 KERAS: baseline/fold_2/model.keras
 KERAS: baseline/fold_3/model.keras
 KERAS: baseline/fold_4/model.keras
 KERAS: baseline/fold_5/model.keras
 KERAS: fine-tuned/fold_1/model.keras
 KERAS: fine-tuned/fold_2/model.keras
 KERAS: fine-tuned/fold_3/model.keras
 KERAS: fine-tuned/fold_4/model.keras
 KERAS: fine-tuned/fold_5/model.keras

=====

READING SAVED FOLD HISTORIES

=====

✓ Baseline | Fold 1 | history.json
 ✓ Baseline | Fold 2 | history.json
 ✓ Baseline | Fold 3 | history.json
 ✓ Baseline | Fold 4 | history.json
 ✓ Baseline | Fold 5 | history.json
 ✓ Fine Tuned | Fold 1 | history.json
 ✓ Fine Tuned | Fold 2 | history.json
 ✓ Fine Tuned | Fold 3 | history.json
 ✓ Fine Tuned | Fold 4 | history.json
 ✓ Fine Tuned | Fold 5 | history.json

=====

RECOVERED CONFIGURATION STRUCTURE

=====

	Configuration	Configuration Directory	Fold	History File	
0	Baseline	baseline	1	history.json	
1	Baseline	baseline	2	history.json	
2	Baseline	baseline	3	history.json	
3	Baseline	baseline	4	history.json	
4	Baseline	baseline	5	history.json	
5	Fine Tuned	fine-tuned	1	history.json	
6	Fine Tuned	fine-tuned	2	history.json	
7	Fine Tuned	fine-tuned	3	history.json	
8	Fine Tuned	fine-tuned	4	history.json	
9	Fine Tuned	fine-tuned	5	history.json	

=====

FOLD-WISE RESULTS

=====

	Configuration	Fold	Best Validation Accuracy (%)	Best Epoch	Final Training Accuracy (%)	Final Validation Accuracy (%)	Final Training Loss	Final Validation Loss	
0	Baseline	1	92.529714	5	98.301488	92.529714	0.080708	0.273119	
1	Baseline	2	89.643466	4	98.598725	88.964349	0.080517	0.309739	
2	Baseline	3	92.529714	5	97.749472	92.529714	0.090412	0.251607	
3	Baseline	4	93.378609	4	97.834396	93.039048	0.091819	0.234533	
4	Baseline	5	91.326529	4	97.835314	90.476191	0.089568	0.308380	

5	Fine Tuned	1	87.775892	3	97.197455	87.096775	0.097501	0.394778
6	Fine Tuned	2	88.455009	5	96.900213	88.455009	0.104410	0.365319
7	Fine Tuned	3	89.813244	5	97.834396	89.813244	0.091306	0.308649
8	Fine Tuned	4	89.983022	5	97.070062	89.983022	0.101256	0.321519
9	Fine Tuned	5	89.965987	4	97.325975	88.265306	0.098793	0.373033

=====

FOLD COUNT VERIFICATION

=====

Configuration	Number of Folds	
0 Baseline	5	
1 Fine Tuned	5	

✓ Baseline: 5-fold CV confirmed.
✓ Fine Tuned: 5-fold CV confirmed.

=====

CORRECT ADDITIONAL-EXERCISE CV SUMMARY

=====

Configuration	Mean CV Accuracy (%)	CV Standard Deviation (%)	Mean Final Validation Accuracy (%)	Final Validation SD (%)	
0 Baseline	91.881607	1.449007	91.507803	1.729646	
1 Fine Tuned	89.198631	1.019683	88.722671	1.193964	

=====


INSPECTING SAVED .KERAS FILES

=====


Configuration	Fold	Keras File	Path	File Size (MB)	Contains config.json	Contains metadata.json	
0 Baseline	1	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	9.533772	True	True	{'keras_
1 Baseline	2	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	9.533772	True	True	{'keras_
2 Baseline	3	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	9.533772	True	True	{'keras_
3 Baseline	4	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	9.533772	True	True	{'keras_
4 Baseline	5	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	9.533772	True	True	{'keras_
5 Fine Tuned	1	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	14.096420	True	True	{'keras_
6 Fine Tuned	2	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	14.096420	True	True	{'keras_
7 Fine Tuned	3	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	14.096420	True	True	{'keras_
8 Fine Tuned	4	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	14.096420	True	True	{'keras_
9 Fine Tuned	5	model.keras	/content/drive/MyDrive/Lab5/Lab5/checkpoints/c...	14.096420	True	True	{'keras_

WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional Exercise 5Fold Accuracy.eps + Additional Exercise 5Fold Accuracy.png

Additional Exercise — 5-Fold Cross-Validation




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Exercise_Mean_SD.eps + Additional_Exercise_Mean_SD.png


Additional Exercise — Mean CV Accuracy $\pm$ SD

WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_1_Accuracy.eps + Additional_Baseline_Fold_1_Accuracy.png


Baseline — Fold 1 Accuracy




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_2_Accuracy.eps + Additional_Baseline_Fold_2_Accuracy.png



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_3_Accuracy.eps + Additional_Baseline_Fold_3_Accuracy.png



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_4_Accuracy.eps + Additional_Baseline_Fold_4_Accuracy.png




1.0 1.5 2.0 2.5 3.0 3.5 4.0 4.5 5.0

Epoch


WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_5_Accuracy.eps + Additional_Baseline_Fold_5_Accuracy.png

Baseline — Fold 5 Accuracy




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_1_Accuracy.eps + Additional_Fine_Tuned_Fold_1_Accuracy.png


Fine Tuned — Fold 1 Accuracy



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_2_Accuracy.eps + Additional_Fine_Tuned_Fold_2_Accuracy.png


Fine Tuned — Fold 2 Accuracy






WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_3_Accuracy.eps + Additional_Fine_Tuned_Fold_3_Accuracy.png

Fine Tuned — Fold 3 Accuracy




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_4_Accuracy.eps + Additional_Fine_Tuned_Fold_4_Accuracy.png


Fine Tuned — Fold 4 Accuracy




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_5_Accuracy.eps + Additional_Fine_Tuned_Fold_5_Accuracy.png

Fine Tuned — Fold 5 Accuracy







WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_1_Loss.eps + Additional_Baseline_Fold_1_Loss.png




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_2_Loss.eps + Additional_Baseline_Fold_2_Loss.png




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_3_Loss.eps + Additional_Baseline_Fold_3_Loss.png






WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_4_Loss.eps + Additional_Baseline_Fold_4_Loss.png




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Baseline_Fold_5_Loss.eps + Additional_Baseline_Fold_5_Loss.png




WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_1_Loss.eps + Additional_Fine_Tuned_Fold_1_Loss.png






WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_2_Loss.eps + Additional_Fine_Tuned_Fold_2_Loss.png

Fine Tuned — Fold 2 Loss



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_3_Loss.eps + Additional_Fine_Tuned_Fold_3_Loss.png


Fine Tuned — Fold 3 Loss



WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
✓ Saved Additional_Fine_Tuned_Fold_4_Loss.eps + Additional_Fine_Tuned_Fold_4_Loss.png


Fine Tuned — Fold 4 Loss





WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
 WARNING:matplotlib.backends.backend_ps:The PostScript backend does not support transparency; partially transparent artists
 ✓ Saved Additional_Fine_Tuned_Fold_5_Loss.eps + Additional_Fine_Tuned_Fold_5_Loss.png

Fine Tuned — Fold 5 Loss



=====

FINAL ADDITIONAL-EXERCISE REPORT TABLE

=====

	Configuration	Mean CV Accuracy (%)	CV Standard Deviation (%)	Mean Final Validation Accuracy (%)	Final Validation SD (%)
0	Baseline	91.881607	1.449007	91.507803	1.729646
1	Fine Tuned	89.198631	1.019683	88.722671	1.193964

=====

GENERATED REPORT MATERIAL

=====

Drive output directory:

/content/drive/MyDrive/Lab5/Lab5/report_materials/additional_exercise

Files generated: 69

- ✓ ADDITIONAL_EXERCISE_CV_SUMMARY.csv
- ✓ ADDITIONAL_EXERCISE_FINAL_REPORT_TABLE.csv
- ✓ ADDITIONAL_EXERCISE_FINAL_TABLE.csv
- ✓ Additional_Additional_Exercise_Validation_Accuracy.eps
- ✓ Additional_Additional_Exercise_Validation_Accuracy.png
- ✓ Additional_Additional_Exercise_Validation_Loss.eps
- ✓ Additional_Additional_Exercise_Validation_Loss.png
- ✓ Additional_Baseline_Fold_1_Accuracy.eps
- ✓ Additional_Baseline_Fold_1_Accuracy.png
- ✓ Additional_Baseline_Fold_1_Loss.eps
- ✓ Additional_Baseline_Fold_1_Loss.png
- ✓ Additional_Baseline_Fold_2_Accuracy.eps
- ✓ Additional_Baseline_Fold_2_Accuracy.png
- ✓ Additional_Baseline_Fold_2_Loss.eps
- ✓ Additional_Baseline_Fold_2_Loss.png
- ✓ Additional_Baseline_Fold_3_Accuracy.eps
- ✓ Additional_Baseline_Fold_3_Accuracy.png
- ✓ Additional_Baseline_Fold_3_Loss.eps
- ✓ Additional_Baseline_Fold_3_Loss.png
- ✓ Additional_Baseline_Fold_4_Accuracy.eps

```

# =====
# CLEAN OVERALL RESULTS - FINAL REPORT VERSION
# =====

import os
import pandas as pd
import numpy as np

LAB5_ROOT = "/content/drive/MyDrive/Lab5/Lab5"
RESULTS_DIR = os.path.join(LAB5_ROOT, "results")

print("=" * 70)
print("CLEANING OVERALL RESULTS")
print("=" * 70)

# -----
# 1. Load the existing overall results if available
# -----

possible_files = [
    "overall_results_corrected.csv",
    "experiment_summary.csv",
    "all_results.csv",
]

overall_df = None
source_file = None

for filename in possible_files:

    path = os.path.join(RESULTS_DIR, filename)

    if os.path.exists(path):

        try:
            temp = pd.read_csv(path)

            if (
                "Configuration" in temp.columns
                and len(temp) > 0
            ):
                overall_df = temp.copy()
                source_file = path
                break

        except Exception:
            pass

# -----
# 2. If no overall CSV exists, construct the table
# -----

if overall_df is None:

    overall_df = pd.DataFrame({
        "Configuration": [
            "Baseline",
            "Best Initialization",
            "Best Regularization",
            "Best Optimizer",
            "Best Hyperparameters",
            "Fine-Tuned Model",
            "Final Selected Model"
        ],
        "CV Accuracy": [
            92.085597,
            92.085597,
            91.915761,
            76.936141,
            89.456522,
            89.266306,
            91.881607
        ],
        "SD": [
            np.nan,
            np.nan,
            np.nan,
            np.nan,
            np.nan,
            np.nan,
            np.nan,
        ]
    })

```

```

        np.nan,
        1.449007
    ],

    "Test Accuracy (%)": [
        np.nan,
        np.nan,
        np.nan,
        np.nan,
        np.nan,
        np.nan,
        89.72472
    ],

    "Training Time (s)": [
        195.185692,
        195.185692,
        191.045890,
        219.573812,
        94.228478,
        1297.457243,
        182.591990
    ]
    })
else:

    print("\nExisting overall-results file found:")
    print(source_file)

# -----
# 3. Identify final-model row
# -----

final_mask = (
    overall_df["Configuration"]
    .astype(str)
    .str.strip()
    .str.lower()
    == "final selected model".lower()
)

# -----
# 4. FORCE VERIFIED FINAL MODEL RESULTS
# -----

if final_mask.any():

    overall_df.loc[
        final_mask,
        "CV Accuracy"
    ] = 91.881607

    overall_df.loc[
        final_mask,
        "SD"
    ] = 1.449007

    overall_df.loc[
        final_mask,
        "Test Accuracy (%)"
    ] = 89.72472

    overall_df.loc[
        final_mask,
        "Training Time (s)"
    ] = 182.591990

# -----
# 5. Remove the stale 2.86 value anywhere it appears
# -----

if "Test Accuracy (%)" in overall_df.columns:

    # Any value equal/close to the stale result is invalid.
    stale_mask = np.isclose(
        pd.to_numeric(
            overall_df["Test Accuracy (%)"],
            errors="coerce"
        ),

```

```

        2.861815,
        atol=1e-5
    )

    # The only legitimate test result is the verified
    # corrected final-model result.
    overall_df.loc[
        stale_mask,
        "Test Accuracy (%)"
    ] = np.nan

# Re-apply final model value after cleanup
if final_mask.any():

    overall_df.loc[
        final_mask,
        "Test Accuracy (%)"
    ] = 89.72472

# -----
# 6. Save cleaned overall results
# -----

clean_path = os.path.join(
    RESULTS_DIR,
    "overall_results_corrected.csv"
)

overall_df.to_csv(
    clean_path,
    index=False
)

# -----
# 7. Display final table
# -----

print("\n")
print("=" * 70)
print("OVERALL RESULTS - CORRECTED")
print("=" * 70)

display(
    overall_df.style.format(
        {
            "CV Accuracy": "{:.6f}",
            "SD": "{:.6f}",
            "Test Accuracy (%)": "{:.6f}",
            "Training Time (s)": "{:.6f}"
        },
        na_rep="-"
    )
)

# -----
# 8. Explicit verification
# -----

print("\n" + "=" * 70)
print("FINAL MODEL VERIFICATION")
print("=" * 70)

if final_mask.any():

    final_row = overall_df.loc[final_mask].iloc[0]

    print(
        f"CV Accuracy      : "
        f"{final_row['CV Accuracy']:.6f}%"
    )

    print(
        f"CV SD              : "
        f"{final_row['SD']:.6f}%"
    )

    print(
        f"Test Accuracy      : "
        f"{final_row['Test Accuracy (%)']:.6f}%"
    )

```

```

    )

    print(
        f"Training Time      : "
        f"{final_row['Training Time (s)']:.6f} s"
    )

    # Hard safety check
    assert abs(
        float(final_row["Test Accuracy (%)"])
        - 89.72472
    ) < 1e-5, \
        "ERROR: Final test accuracy is not 89.72472%!"

print("\n✓ Stale 2.86% result removed.")
print("✓ Verified Final Test Accuracy = 89.72%")
print("\nSaved corrected overall results:")
print(clean_path)
print("=" * 70)

```

=====

CLEANING OVERALL RESULTS

=====

Existing overall-results file found:
/content/drive/MyDrive/Lab5/Lab5/results/overall_results_corrected.csv

=====

OVERALL RESULTS – CORRECTED

=====

	Configuration	CV Accuracy	SD	Test Accuracy (%)	Training Time (s)
0	Baseline	92.085597	—	—	195.185692
1	Best Initialization	92.085597	—	—	195.185692
2	Best Regularization	91.915761	—	—	191.045890
3	Best Optimizer	76.936141	—	—	219.573812
4	Best Hyperparameters	89.456522	—	—	94.228478
5	Fine-Tuned Model	89.266306	—	—	1297.457243
6	Final Selected Model	91.881607	1.449007	89.724720	182.591990

=====

FINAL MODEL VERIFICATION

=====

CV Accuracy : 91.881607%

CV SD : 1.449007%

Test Accuracy : 89.724720%

Training Time : 182.591990 s

✓ Stale 2.86% result removed.
✓ Verified Final Test Accuracy = 89.72%

Saved corrected overall results:
/content/drive/MyDrive/Lab5/Lab5/results/overall_results_corrected.csv

=====

```

# =====
# CELL – PRINT ALL REMAINING EXPERIMENT 5 RESULTS
#
# PURPOSE:
#   Recover and display ALL already-saved results needed for
#   the main experiment, excluding the Additional Exercise.
#
# DOES NOT:
#   - train
#   - retrain
#   - evaluate a model
#   - download weights
#
# IT ONLY READS SAVED FILES / HISTORIES / FIGURES.
# =====

```

```

import os
import json
import glob
import re
import shutil
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```

```

import os
import sys
import IPython.display as display, Image

from IPython.display import display, Image

# =====
# 1. MOUNT GOOGLE DRIVE
# =====

from google.colab import drive

print("=" * 70)
print("MOUNTING GOOGLE DRIVE")
print("=" * 70)

if not os.path.exists("/content/drive/MyDrive"):
    drive.mount("/content/drive")
else:
    print("✓ Google Drive already mounted.")

# =====
# 2. PATHS
# =====

LAB5_ROOT = "/content/drive/MyDrive/Lab5/Lab5"

CHECKPOINT_ROOT = os.path.join(
    LAB5_ROOT,
    "checkpoints"
)

RESULTS_ROOT = os.path.join(
    LAB5_ROOT,
    "results"
)

REPORT_ROOT = os.path.join(
    LAB5_ROOT,
    "report_materials"
)

PLOT_ROOT = os.path.join(
    REPORT_ROOT,
    "main_experiment"
)

RUNTIME_ROOT = os.path.join(
    "/content",
    "Lab5_report_materials",
    "main_experiment"
)

os.makedirs(
    PLOT_ROOT,
    exist_ok=True
)

os.makedirs(
    RUNTIME_ROOT,
    exist_ok=True
)

# =====
# 3. HELPER FUNCTIONS
# =====

def find_files(patterns, roots):
    found = []

    for root in roots:

        if not os.path.exists(root):
            continue

        for pattern in patterns:

            found.extend(
                glob.glob(
                    os.path.join(
                        root,
                        "**",

```

```

        pattern
    ),
    recursive=True
)
)

return sorted(
    list(
        dict.fromkeys(found)
    )
)

def load_csv_if_exists(path):

    if os.path.exists(path):

        try:
            return pd.read_csv(path)
        except Exception as e:
            print(
                f"⚠ Could not read {path}: {e}"
            )

    return None

def extract_history(obj):

    if isinstance(obj, dict):

        required = {
            "loss",
            "accuracy",
            "val_loss",
            "val_accuracy"
        }

        if required.intersection(
            set(obj.keys())
        ):
            return obj

        if "history" in obj:

            result = extract_history(
                obj["history"]
            )

            if result is not None:
                return result

        for value in obj.values():

            result = extract_history(
                value
            )

            if result is not None:
                return result

    return None

def load_history(path):

    try:

        with open(
            path,
            "r"
        ) as f:

            raw = json.load(f)

            return extract_history(raw)

    except Exception:
        return None

def save_existing_figure(
    source,

```

```

    filename
):

    """
    Copies an already-created figure to both
    Runtime and Drive report folders.
    """

    if not os.path.exists(source):
        return False

    drive_dest = os.path.join(
        PLOT_ROOT,
        filename
    )

    runtime_dest = os.path.join(
        RUNTIME_ROOT,
        filename
    )

    shutil.copy2(
        source,
        drive_dest
    )

    shutil.copy2(
        source,
        runtime_dest
    )

    return True

def save_plot(
    basename
):

    plt.savefig(
        os.path.join(
            PLOT_ROOT,
            basename + ".eps"
        ),
        format="eps",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        os.path.join(
            PLOT_ROOT,
            basename + ".png"
        ),
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        os.path.join(
            RUNTIME_ROOT,
            basename + ".eps"
        ),
        format="eps",
        dpi=600,
        bbox_inches="tight"
    )

    plt.savefig(
        os.path.join(
            RUNTIME_ROOT,
            basename + ".png"
        ),
        format="png",
        dpi=600,
        bbox_inches="tight"
    )

# =====
# 4. LOAD ALL RESULT CSV FILES
# =====

```

```

print("\n" + "=" * 70)
print("SAVED RESULT TABLES")
print("=" * 70)

csv_files = sorted(
    glob.glob(
        os.path.join(
            RESULTS_ROOT,
            "*.csv"
        )
    )
)

print(
    f"\nFound {len(csv_files)} CSV files."
)

for path in csv_files:
    print(
        "  ✓",
        os.path.basename(path)
    )

# =====
# 5. DISPLAY IMPORTANT RESULT TABLES
# =====

RESULT_FILES = {

    "Initialization":
        "initialization_results.csv",

    "Regularization":
        "regularization_results.csv",

    "Batch Normalization":
        "batch_normalization_results.csv",

    "Optimizer":
        "optimizer_results.csv",

    "Hyperparameter":
        "hyperparameter_results.csv",

    "Transfer Learning":
        "transfer_learning_results.csv",

    "Fine-Tuning":
        "correct_fine_tuning_results.csv",

    "Cross-Validation":
        "cv_results.csv",

    "Cross-Validation Summary":
        "cv_summary.csv",

    "Final Model":
        "final_model_results.csv",

    "Overall Results":
        "overall_results.csv"

}

loaded_tables = {}

for label, filename in RESULT_FILES.items():

    path = os.path.join(
        RESULTS_ROOT,
        filename
    )

    df = load_csv_if_exists(
        path
    )

    if df is not None:

```

```

loaded_tables[label] = df

print(
    "\n" + "-" * 70
)

print(
    label.upper()
)

print(
    "-" * 70
)

display(
    df
)

else:

    print(
        f"\n⚠ {label}: {filename} not found."
    )

# =====
# 6. FINAL CORRECTED MODEL RESULT
# =====

print("\n" + "=" * 70)
print("VERIFIED CORRECTED FINAL MODEL")
print("=" * 70)

final_df = loaded_tables.get(
    "Final Model"
)

if final_df is not None:

    display(
        final_df
    )

    # Explicitly show only the saved corrected result
    for col in final_df.columns:

        if (
            "Test Accuracy" in col
            or
            "Test Loss" in col
            or
            "Validation" in col
            or
            "Training Time" in col
        ):
            pass

else:

    print(
        "⚠ final_model_results.csv not found."
    )

# =====
# 7. FIND ALL HISTORY FILES
# =====

print("\n" + "=" * 70)
print("SAVED TRAINING HISTORIES")
print("=" * 70)

history_files = find_files(
    ["*.json"],
    [
        CHECKPOINT_ROOT,
        RESULTS_ROOT
    ]
)

history_files = [
    p for p in history_files
    if (

```

```

        "history" in os.path.basename(p).lower()
    )
]

print(
    f"\nFound {len(history_files)} history file(s)."
)

for p in history_files:
    print(
        "  ✓",
        os.path.relpath(
            p,
            LAB5_ROOT
        )
    )

# =====
# 8. ORGANIZE HISTORIES BY EXPERIMENT
# =====

history_groups = {

    "Initialization": [],
    "Regularization": [],
    "Batch Normalization": [],
    "Optimizer": [],
    "Hyperparameter": [],
    "Transfer Learning": [],
    "Fine-Tuning": [],
    "Final Model": []

}

for path in history_files:
    low = path.lower()

    if (
        "initial" in low
        or
        "initialization" in low
    ):
        history_groups[
            "Initialization"
        ].append(path)

    elif (
        "regularization" in low
    ):
        history_groups[
            "Regularization"
        ].append(path)

    elif (
        "batch_normal" in low
        or
        "batchnormal" in low
        or
        "batch_norm" in low
    ):
        history_groups[
            "Batch Normalization"
        ].append(path)

    elif (
        "optimizer" in low
    ):
        history_groups[
            "Optimizer"
        ].append(path)

    elif (
        "hyperparameter" in low
        or
        "learning_rate" in low
        or
        "batch_size" in low
        or

```

```

        "dropout" in low
    ):
        history_groups[
            "Hyperparameter"
        ].append(path)

    elif (
        "transfer" in low
        or
        "feature_extraction" in low
    ):
        history_groups[
            "Transfer Learning"
        ].append(path)

    elif (
        "fine_tun" in low
        or
        "finetun" in low
    ):
        history_groups[
            "Fine-Tuning"
        ].append(path)

    elif (
        "final_model" in low
        or
        "final_model_corrected" in low
    ):
        history_groups[
            "Final Model"
        ].append(path)

# =====
# 9. PLOT 1 – INITIALIZATION TRAINING LOSS
# =====

init_records = []

for path in history_groups[
    "Initialization"
]:

    h = load_history(path)

    if h is None:
        continue

    if "loss" not in h:
        continue

    name = os.path.basename(
        os.path.dirname(path)
    )

    init_records.append(
        (
            name,
            h
        )
    )

if init_records:

    print("\n" + "=" * 70)
    print("PLOT 1 – INITIALIZATION: TRAINING LOSS")
    print("=" * 70)

    plt.figure(
        figsize=(9, 5)
    )

    for name, h in init_records:

        plt.plot(
            np.arange(
                1,
                len(h["loss"]) + 1
            ),
            h["loss"],

```

```

        linewidth=2,
        label=name
    )

    plt.xlabel(
        "Epoch"
    )

    plt.ylabel(
        "Training Loss"
    )

    plt.title(
        "Training Loss vs. Epoch for Weight Initialization Methods"
    )

    plt.grid(
        True,
        alpha=0.3
    )

    plt.legend()

    plt.tight_layout()

    save_plot(
        "Plot_1_Initialization_Training_Loss"
    )

    plt.show()

    plt.close()

# =====
# 10. PLOT 2 – INITIALIZATION VALIDATION ACCURACY
# =====

if init_records:

    print("\n" + "=" * 70)
    print("PLOT 2 – INITIALIZATION: VALIDATION ACCURACY")
    print("=" * 70)

    plt.figure(
        figsize=(9, 5)
    )

    for name, h in init_records:

        if "val_accuracy" not in h:
            continue

        plt.plot(
            np.arange(
                1,
                len(h["val_accuracy"]) + 1
            ),
            np.asarray(
                h["val_accuracy"]
            ) * 100,
            linewidth=2,
            label=name
        )

    plt.xlabel(
        "Epoch"
    )

    plt.ylabel(
        "Validation Accuracy (%)"
    )

    plt.title(
        "Validation Accuracy vs. Epoch for Weight Initialization Methods"
    )

    plt.grid(
        True,
        alpha=0.3
    )

    plt.legend()

```

```

plt.tight_layout()

save_plot(
    "Plot_2_Initialization_Validation_Accuracy"
)

plt.show()

plt.close()

# =====
# 11. GENERIC HISTORY COLLECTION FOR PLOTS 3-12
# =====

def collect_histories(group):

    output = []

    for path in history_groups[group]:

        h = load_history(path)

        if h is None:
            continue

        output.append(
            (
                os.path.basename(
                    os.path.dirname(path)
                ),
                h
            )
        )

    return output

# =====
# 12. REGULARIZATION PLOTS 3 & 4
# =====

reg_records = collect_histories(
    "Regularization"
)

if reg_records:

    # Plot 3
    print("\n" + "=" * 70)
    print("PLOT 3 - REGULARIZATION: TRAINING / VALIDATION ACCURACY")
    print("=" * 70)

    plt.figure(
        figsize=(10, 6)
    )

    for name, h in reg_records:

        if "accuracy" in h:

            plt.plot(
                np.arange(
                    1,
                    len(h["accuracy"]) + 1
                ),
                np.asarray(
                    h["accuracy"]
                ) * 100,
                linewidth=1.8,
                label=name + " - Train"
            )

        if "val_accuracy" in h:

            plt.plot(
                np.arange(
                    1,
                    len(h["val_accuracy"]) + 1
                ),
                np.asarray(
                    h["val_accuracy"]

```

```

        ) * 100,
        linestyle="--",
        linewidth=1.8,
        label=name + " - Val"
    )

plt.xlabel(
    "Epoch"
)

plt.ylabel(
    "Accuracy (%)"
)

plt.title(
    "Training and Validation Accuracy - Regularization"
)

plt.grid(
    True,
    alpha=0.3
)

plt.legend(
    fontsize=8
)

plt.tight_layout()

save_plot(
    "Plot_3_Regularization_Accuracy"
)

plt.show()

plt.close()

# Plot 4
print("\n" + "=" * 70)
print("PLOT 4 - REGULARIZATION: TRAINING / VALIDATION LOSS")
print("=" * 70)

plt.figure(
    figsize=(10, 6)
)

for name, h in reg_records:

    if "loss" in h:

        plt.plot(
            np.arange(
                1,
                len(h["loss"]) + 1
            ),
            h["loss"],
            linewidth=1.8,
            label=name + " - Train"
        )

    if "val_loss" in h:

        plt.plot(
            np.arange(
                1,
                len(h["val_loss"]) + 1
            ),
            h["val_loss"],
            linestyle="--",
            linewidth=1.8,
            label=name + " - Val"
        )

plt.xlabel(
    "Epoch"
)

plt.ylabel(
    "Loss"
)

```